

RADIANT Theory Manual

Physics Reference for the RADIANT EO Sensor Performance Model

RADIANT Project

v0.1.0 (v0.1.0-96-g95871c6f-dirty)
Built 2026-09-20

Contents

Notation and Symbols	4
0.1 Canonical units	4
0.2 Coordinate, attitude, and indexing conventions	5
0.3 Radiometric symbols	5
0.4 Geometry and sampling symbols	6
0.5 Spatial symbols — PSF and MTF	7
0.6 Detector, readout, and noise symbols	7
0.7 Calibration symbols	8
0.8 Physical constants	8
0.9 Operators, accents, and reused symbols	9
0.10 Acronyms	9
1 Introduction	11
1.1 The chain at a glance	11
1.2 The three radiometric regimes	13
1.3 What the model predicts	14
1.4 Model scope — what is deliberately absent	14
1.5 How to read this manual	15
2 Geometry & Sampling	16
2.1 The viewing triangle	16
2.2 Slant range	17
2.3 Circular-orbit kinematics	17
2.4 GSD — nadir and off-nadir	17
2.5 Swath and access	18
2.6 Smear kinematics and TDI line-rate matching	18
2.7 Solar geometry	18
2.8 Attitude: Euler ZYX	19
2.9 Sampling on the ground	19
3 Radiometric Signal Chain	20
3.1 Notation	20
3.2 The chain at a glance	20
3.3 Foundations	21
3.4 How the foundations feed the chain	32
3.5 Assumptions and Limitations	33
4 Atmosphere Models	34
4.1 Two families, one contract	34
4.2 The simple parametric model	35
4.3 The library-backed models	47

4.4	What the models do not represent	50
4.5	References	51
5	Spatial Model	52
5.1	The two spatial paths	52
5.2	Diffraction-limited PSF	53
5.3	Optical MTF from the pupil autocorrelation	53
5.4	Wavefront error, Zernike modes, and Strehl	54
5.5	Detector-plane kernels: pixel aperture, diffusion, IPC	54
5.6	Platform kernels: jitter and smear	55
5.7	Atmospheric turbulence	55
5.8	TDI mis-registration — the one MTF-only term	56
5.9	Ensquared energy (EE_box)	56
5.10	Sampling: Nyquist, Q, and folded MTF	57
5.11	RER and the edge response	57
5.12	Scope and deferred effects	58
6	Noise Model	59
6.1	Overview — RSS composition	59
6.2	Noise Taxonomy	59
6.3	Acquisition scaling	61
6.4	Noise regime selection	62
6.5	When each term dominates	62
6.6	Parameter cross-reference	62
7	Calibration Error Model	63
7.1	Overview — a terms-only stage after readout	63
7.2	Cal points — where the correction is anchored	64
7.3	Post-NUC residual — the correctability parabola	65
7.4	Cal-source spatial non-uniformity	66
7.5	Drift between calibration events	67
7.6	The internal-shutter cal path	68
7.7	Composition — how calibration enters the budget	68
7.8	Calibration-limited NEDT	69
7.9	The bias budget — accuracy, not precision	70
7.10	Assumptions, validity, and what is deferred	71
7.11	Parameter cross-reference	71
8	Performance Metrics	72
8.1	SNR family	72
8.2	NEDT	72
8.3	NEI, NEP, and D*	73
8.4	NIIRS via GIQE-5 (and the IIRS placeholder)	73
8.5	Johnson criteria and minimum resolvable	74
8.6	Detection range	74
8.7	Saturation, dynamic range, BLIP	75
8.8	Radiometric inversions and sensitivity relatives	75
8.9	Metric selection	76

A	Radiometric Model — Mixed Refractive and Reflective Optical Train	77
A.1	Element types and their radiometric character	77
A.2	Per-element radiometric quantities	78
A.3	Signal path through the full optical train	79
A.4	Thermal background path from each element	79
A.5	Combined focal plane irradiance	81
A.6	Kirchhoff self-consistency check	81
A.7	Implementation recipe	81
A.8	Key physical distinctions summary	83
	References	84

Notation and Symbols

This front-matter chapter is the manual’s single notation home. Every symbol, canonical unit, and sign convention used in the chapters that follow is defined here once; the chapters use the symbols without redefining them. Where a chapter needs a symbol that is local to one derivation, it introduces it at the point of use and the local meaning wins for that section only.

Two rules govern everything below:

- **One canonical unit per quantity.** RADIANT computes internally in the units of the canonical-units table below — wavelength in μm , angles in radians, length in meters, time in seconds, radiance in $\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$, noise in e- RMS. No module works in any other unit.
- **Conversion happens exactly once, at a boundary.** User input converts on entry; external files convert in their reader. A unit conversion inside a physics equation is a defect, not a convenience.

0.1 Canonical units

Quantity	Symbol	Canonical unit	Notes
Wavelength	λ	μm	The primary spectral variable; arrays ascend in λ
Wavelength in SI	λ_m	m	$\lambda_m = \lambda \cdot 10^{-6}$; used wherever h , c , or k_B appears
Wavenumber	ν_{cm}	cm^{-1}	Derived only ($\nu_{cm} = 10^4/\lambda$), never stored as the primary axis
Spectral radiance	$L(\lambda)$	$\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$	MODTRAN’s $\text{W}/\text{cm}^2/\text{sr}/\mu\text{m}$ converts by $\times 10^4$ in the reader
In-band radiance	L	$\text{W}/\text{m}^2/\text{sr}$	$\int L(\lambda) d\lambda$ over the bandpass
Spectral irradiance	$E(\lambda)$	$\text{W}/\text{m}^2/\mu\text{m}$	
In-band irradiance	E	W/m^2	
Spectral intensity	$I(\lambda)$	$\text{W}/\text{sr}/\mu\text{m}$	Point sources
Spectral photon radiance	$L_q(\lambda)$	$\text{photon}/\text{s}/\text{m}^2/\text{sr}/\mu\text{m}$	Derived: $L_q = L\lambda_m/(hc)$
Temperature	T	K	Users may enter K, $^{\circ}\text{C}$, or $^{\circ}\text{F}$; the affine conversion is applied once at entry
Angle (all internal)	$\theta, \eta, \zeta, \dots$	rad	Users enter large angles in degrees, small angles in μrad
Length	D, f, p, R_s	m	
Time	t_{int}, T_{frame}	s	Displays may use ms or μs ; storage is seconds
Signal	S	e-	Photoelectrons at the sense node, after integration
Noise	σ	e- RMS	Every noise term, before any DN conversion

Quantity	Symbol	Canonical unit	Notes
Digital number	DN	DN	$DN = S/g$ with g in e-/DN
Spatial frequency (focal plane)	ν	cy/m	Quoted as cy/mm in worked numbers; always labeled
Spatial frequency (angular)	ν_{ang}	cy/mrad	νf with f in m is cy/rad, so $\nu_{ang} = \nu f / 1000$; always labeled
Transfer function	MTF(ν)	dimensionless	$0 \leq \text{MTF} \leq 1$
Dimensionless fractions	$\varepsilon, \rho, \tau, \text{EE}$	dimensionless	Stated on the interval $[0, 1]$ where physical

Reading the tables that follow. Units in square brackets are the canonical units above. A quantity listed as dimensionless still carries a physical range, given where it matters.

0.2 Coordinate, attitude, and indexing conventions

Property	Convention
Handedness	Right-handed
+Z	Toward the target, along the boresight / optical axis
+X	Cross-track (perpendicular to the flight direction)
+Y	Along-track (flight direction projected into the image plane)
Euler sequence	3-2-1 (ZYX): yaw about +Z, then pitch about the once-rotated +Y, then roll about the twice-rotated +X
Pixel indexing	$[\text{row}, \text{col}] = [y, x] = [\text{along-track}, \text{cross-track}]$, 0-indexed

For a nadir-pointing space sensor +Z is nadir; for a ground observer +Z is toward the target in the sky. The convention holds for every viewing geometry the model supports without redefinition.

Viewing geometry is referenced to the **target**, not the sensor: the primary angle is the line-of-sight zenith angle at the target, θ_o , from which the look angle η at the sensor, the Earth central angle Λ , and the slant range R_s follow.

0.3 Radiometric symbols

Symbol	Meaning	Units
$B(\lambda, T)$	Planck blackbody spectral radiance	W/m ² /sr/μm
$B_q(\lambda, T)$	Planck photon spectral radiance, B/E_{ph}	photon/s/m ² /sr/μm
$\partial B / \partial T$	temperature derivative of Planck radiance (the NEDT kernel)	W/m ² /sr/μm/K
$\bar{L}$	band-averaged radiance	W/m ² /sr/μm
L_{band}	band-integrated radiance	W/m ² /sr

Symbol	Meaning	Units
$L_{\text{target}}, L_{\text{bg}}$	target and background spectral radiance at the source	W/m ² /sr/μm
L_{ap}	at-aperture spectral radiance	W/m ² /sr/μm
L_{path}	atmospheric path radiance (scattered plus emitted into the line of sight)	W/m ² /sr/μm
L_{sky}	hemispheric downwelling sky radiance	W/m ² /sr/μm
$E_{\text{sun}}(\lambda)$	top-of-atmosphere solar spectral irradiance	W/m ² /μm
E_{ph}	photon energy, hc/λ_m	J
$\varepsilon(\lambda)$	emissivity	dimensionless, 0–1
$\rho(\lambda)$	hemispherical reflectance (albedo)	dimensionless, 0–1
f_r	bidirectional reflectance distribution function (BRDF)	sr ⁻¹
$\tau_{\text{atm}}(\lambda)$	atmospheric transmittance along the slant path	dimensionless, 0–1
$\tau_{\text{opt}}(\lambda)$	net optical-train transmittance	dimensionless, 0–1
$R(\lambda), T(\lambda)$	surface reflectance and transmittance of an optical element	dimensionless, 0–1
θ_{sun}	solar zenith angle at the surface	rad
x	dimensionless Planck argument, $x = hc/(\lambda_m k_B T)$	dimensionless
ff	sub-pixel fill fraction, target area over pixel footprint	dimensionless, 0–1
$QE(\lambda)$	quantum efficiency	e-/photon

0.4 Geometry and sampling symbols

Symbol	Meaning	Units
R	mean Earth radius, 6371.0 km	m
h	sensor altitude above the surface	m
R_s	slant range, sensor to target	m
η	look (off-nadir) angle at the sensor	rad
θ_o	line-of-sight zenith angle at the target (equals incidence angle for a surface target)	rad
ε_{el}	grazing/elevation angle at the target, $\pi/2 - \theta_o$	rad
Λ	Earth central angle between sub-sensor point and target	rad
ζ	path zenith angle used by the atmosphere models	rad
A_t	target projected area	m ²
A_{ap}	collecting aperture area	m ²
Ω_{pix}	pixel solid angle, $p_x p_y / f^2$	sr
Ω_t	target solid angle, A_t / R_s^2	sr
D	aperture diameter	m
f	focal length	m
$F_{\#}$	f-number, f / D	dimensionless
p, p_x, p_y	pixel pitch (cross-track, along-track)	m
IFOV	instantaneous field of view of one pixel, p / f	rad
GSD	ground sample distance	m
v_{img}	image-plane motion rate during integration	m/s
θ	target angular extent, $\sqrt{A_t} / R_s$	rad

0.5 Spatial symbols — PSF and MTF

Symbol	Meaning	Units
P	complex pupil function, $P = A e^{i2\pi W/\lambda}$	dimensionless
W	wavefront error across the pupil	m (quoted in waves at a stated λ)
σ_{WFE}	RMS wavefront error	waves at the stated reference wavelength
Z_n^m	Zernike coefficient (Noll-ordered)	waves RMS
PSF	point spread function (the degraded, chain-assembled one)	dimensionless, normalized to unit volume
OTF, MTF(ν)	optical transfer function and its modulus	dimensionless
MTF _{sys} (ν)	system MTF, the product of all contributor MTFs	dimensionless
ϵ	central-obscuration ratio, D_{sec}/D	dimensionless, 0–1
ν_c	optical cutoff frequency, $1/(\lambda F_\#)$	cy/m
ν_{Nyq}	detector Nyquist frequency, $1/(2p)$	cy/m
Q	sampling parameter, $\lambda F_\#/p$	dimensionless
S_{tr}	Strehl ratio (degraded PSF peak over diffraction-limited peak)	dimensionless, 0–1
EE _{box}	ensquared energy inside one pixel, from the degraded PSF	dimensionless, 0–1
RER	relative edge response	dimensionless
σ_j	RMS angular jitter	rad (entered in μ rad)
L_{sm}	smear length during integration	m at the focal plane
r_0	Fried parameter	m
$C_n^2(z)$	refractive-index structure constant profile	$m^{-2/3}$
L_d	minority-carrier diffusion length	m
α_{IPC}	interpixel-capacitance coupling coefficient	dimensionless

0.6 Detector, readout, and noise symbols

Symbol	Meaning	Units
S, S_{bg}	target and background signal	e-
N_{fw}	full-well capacity	e-
g	conversion gain, $N_{fw}/2^n$ for an n -bit converter	e-/DN
t_{int}	integration time	s
T_{frame}	frame period	s
J_{dark}	dark-current rate	e-/s
E_a	dark-current activation energy	eV
σ_i	one noise term, input-referred to electrons	e- RMS
σ_{tot}	total noise, the RSS of all contributing terms	e- RMS
σ_{read}	read noise	e- RMS
σ_{kTC}	reset (kTC) noise, $\sqrt{k_B T C_{node}}/q$	e- RMS

Symbol	Meaning	Units
σ_{ADC}	quantization noise, $g/\sqrt{12}$	e- RMS
k	PRNU fraction (multiplicative gain dispersion) — the <i>residual</i> after NUC in the noise chapter's $\sigma_{PRNU} = kS$, the <i>pre-correction</i> value in the calibration chapter's one-point residual, which is what an offset-only correction leaves intact	dimensionless
σ_{DSNU}	dark-signal non-uniformity	e- RMS
$\sigma_{clutter}$	scene-induced spatial variance term	e- RMS
$N_{TDI}, N_{coadd}, M_{bin}$	TDI stages, co-added frames, binned pixels	dimensionless counts
dS/dT	signal derivative with respect to scene temperature	e-/K
$\oplus$	root-sum-square combination, $a \oplus b = \sqrt{a^2 + b^2}$	—

0.7 Calibration symbols

Symbol	Meaning	Units
S_1, S_2, S_3	calibration-point signals (low, mid/high, high)	e-
$T_{cal,j}$	calibration-source temperature at point j	K
f_j	calibration-point signal declared as a fraction of the scene signal	dimensionless
S_{ref}	full-scale reference for the nonlinearity coefficient	e-
β	1 σ per-pixel response-nonlinearity fraction at full scale	dimensionless
D_j	calibration-point signal derivative, dS/dT at $T_{cal,j}$	e-/K
ΔT_{unif}	1 σ spatial non-uniformity of the calibration source	K
r_g	gain-drift rate since calibration	s ⁻¹
r_o	offset-drift rate since calibration	e-/s
t_{cal}	elapsed time since the calibration event	s
σ_{cal}	RSS of the calibration residual noise terms	e- RMS
b_i, b_{tot}	per-source and total fractional radiometric bias (1 σ)	dimensionless
$g(T)$	band-shift log-derivative of the photon-weighted band integral	μm^{-1}
$\delta\lambda$	band-center (spectral-calibration) uncertainty	μm
f_{fore}	photon-weighted fraction of near-field emission from elements ahead of an internal cal shutter	dimensionless, 0–1

0.8 Physical constants

Every constant is defined once in the code and never re-entered in an equation. The first six rows are the CODATA 2018 values — the first five exact by the SI definitions, the second radiation constant derived from them. The last row is not a CODATA quantity: it is the mean Earth radius, a geodetic convention (see the geometry chapter).

Constant	Symbol	Value	Units
Speed of light	c	2.99792458×10^8	m/s
Planck constant	h	$6.62607015 \times 10^{-34}$	J·s
Boltzmann constant	k_B	1.380649×10^{-23}	J/K
Stefan-Boltzmann constant	σ_{SB}	$5.670374419 \times 10^{-8}$	W/m ² /K ⁴
Elementary charge	q	$1.602176634 \times 10^{-19}$	C
Second radiation constant	$c_2 = hc/k_B$	1.4387769×10^{-2}	m·K
Mean Earth radius	R	6.3710×10^6	m

0.9 Operators, accents, and reused symbols

Form	Meaning
$\bar{X}$	band average of X over the bandpass
X_m	X expressed in SI meters (used only for λ_m)
$a \oplus b$	root-sum-square, $\sqrt{a^2 + b^2}$
$\star$	correlation (used for the pupil autocorrelation)
$\mathcal{F}\{\cdot\}$	Fourier transform
$\prod_i, \sum_i$	product and sum over contributor index i

Four symbols carry more than one meaning across the classical literature this manual follows. They are disambiguated by chapter, and never used in two senses in one equation:

- η is the **look (off-nadir) angle** everywhere except the performance chapter's NEI/NEP/D* section, where the detector literature's η (quantum efficiency) and η_{sys} (end-to-end photon-to-electron efficiency) are kept; that section defines both at the point of use.
- ε is **emissivity** in the radiometric, atmosphere, and calibration chapters. The geometry chapter's grazing/elevation angle is written ε_{el} here to keep the two apart.
- ϵ is the **central-obscuratio ratio** in the spatial chapter only.
- ν is **spatial frequency** [cy/m] throughout the spatial and performance chapters; spectral wavenumber, which the atmospheric literature also writes ν , appears only as ν_{cm} [cm⁻¹] at MODTRAN interfaces. Angular spatial frequency is written f_a [cy/rad] in the spatial chapter's turbulence section, where the classical results are stated that way.

Two chapters keep a legacy spelling that predates this table and is retained rather than churned: the radiometric-chain chapter writes the slant range as R (here and in the geometry chapter it is R_s , with R reserved for the mean Earth radius), and it writes spectral power as $\Phi(\lambda)$ [W/μm].

0.10 Acronyms

Acronym	Expansion
ADC	analog-to-digital converter
BLIP	background-limited infrared performance
BRDF	bidirectional reflectance distribution function
CDS	correlated double sampling
CSNR	contrast signal-to-noise ratio
DN	digital number
DROIC	digital readout integrated circuit (photon-counting)
DSNU	dark-signal non-uniformity
EE	ensquared (or encircled) energy
ERF	edge response function
FPA	focal plane array
FPN	fixed-pattern noise
FWHM	full width at half maximum
GIQE	General Image Quality Equation
GSD	ground sample distance
IFOV	instantaneous field of view
IPC	interpixel capacitance
LOS	line of sight
LSF	line spread function
MRT	minimum resolvable temperature difference
MTF	modulation transfer function
NEDL	noise-equivalent differential radiance
NEDT	noise-equivalent differential temperature
NEI	noise-equivalent irradiance
NEP	noise-equivalent power
NIIRS	National Imagery Interpretability Rating Scale
NUC	non-uniformity correction
OTF	optical transfer function
PRNU	photo-response non-uniformity
PSF	point spread function
QE	quantum efficiency
RER	relative edge response
ROIC	readout integrated circuit
RSS	root-sum-square
SCNR	signal-to-clutter-plus-noise ratio
SNR	signal-to-noise ratio
TDI	time delay and integration
TOA	top of atmosphere
WFE	wavefront error

Introduction

RADIANT is a first-principles performance model for electro-optical sensors, space-based and airborne, from the ultraviolet through the long-wave infrared. It answers questions of the form *what will this sensor, looking at that scene through this atmosphere, actually measure* — and it answers them by carrying a physical quantity, in canonical units, from the emitting surface to a photoelectron count and then to a performance metric, with no unexplained factor anywhere in between.

This chapter states what the chain is, what selects the radiometric regime a target falls into, and where the model’s boundaries lie. Every equation named here is developed in a later chapter; the purpose of this one is that a reader always knows which chapter to open.

1.1 The chain at a glance

RADIANT models the end-to-end signal chain as ten sequential stages. Each stage is a pure function of an immutable chain state: it consumes what earlier stages published, adds its own radiometric frames, noise terms, bias terms, and transfer-function contributions, and returns a new state. No stage reads a file, calls another stage, or edits what an earlier stage published.

#	Stage	What it establishes	Detailed in
0	Geometry	Slant range R_s [m], target-referenced zenith angle θ_o [rad], solar geometry, ground sampling, motion rates	Ch. 2
1	Source	Target and background spectral radiance $L_{\text{target}}(\lambda)$, $L_{\text{bg}}(\lambda)$ [W/m ² /sr/μm]; tentative regime	Ch. 3
2	Atmosphere	Transmittance $\tau_{\text{atm}}(\lambda)$, path radiance $L_{\text{path}}(\lambda)$, downwelling $L_{\text{sky}}(\lambda)$, and the at-aperture radiance they produce	Ch. 3, Ch. 4
3	Optics	Aperture area A_{ap} [m ²], pixel solid angle Ω_{pix} [sr], throughput $\tau_{\text{opt}}(\lambda)$, warm-optics self-emission, the diffraction-plus-aberration PSF and optical MTF, and the final regime	Ch. 3, Ch. 5
4	Platform	Jitter and smear degradation of the PSF and the MTF budget; EE_{box} from the fully degraded PSF	Ch. 5
5	Spectral integration	The single spectral-to-scalar collapse: spectra in, photoelectrons out	Ch. 3
6	Detector	Quantum efficiency, dark current, the noise budget’s detector terms, detector-plane MTF kernels	Ch. 5, Ch. 6
7	Readout	TDI, binning, co-adding, gain, ADC — signal scaling and the readout noise terms	Ch. 6
8	Calibration	Post-NUC residual fixed-pattern noise, drift, and the radiometric bias budget	Ch. 7
9	Performance	SNR, NEDT, NIIRS, system MTF, detection range, and the rest of the metric layer	Ch. 8

1.1.1 Geometry first

Geometry runs before radiometry, not alongside it. The reason is that almost every radiometric quantity downstream is conditioned on a geometric one: the atmospheric slant column depends on the path zenith angle, the solar reflected term depends on the solar zenith angle, the sub-pixel fill fraction depends on the target's angular extent at the slant range, and the smear kernel depends on the image-plane motion rate. Deriving those inside the stages that consume them would mean deriving several of them more than once, which is how two parts of a model come to disagree about the same angle. The geometry stage solves the viewing triangle once, publishes R_s , θ_o , η , and the sampling and motion quantities, and every later stage reads them.

1.1.2 The radiometric spine

Stages 1 through 5 form the radiometric spine. The source stage builds spectral radiance from Planck emission, graybody emissivity, reflected solar irradiance through a BRDF, or a user-supplied spectrum. The atmosphere stage attenuates it and adds what the atmosphere itself contributes along the line of sight:

$$L_{\text{ap}}(\lambda) = L_{\text{target}}(\lambda) \tau_{\text{atm}}(\lambda) + L_{\text{path}}(\lambda)$$

The optics stage supplies the collecting area, the solid angle the pixel subtends, the train's transmittance, and — in the thermal infrared — the train's own emission, whose emissivity is never an independent input but is always derived from the element's reflectance and transmittance by Kirchhoff's law. The spectral-integration stage then performs the one collapse the whole model is allowed:

$$S = \int_{\lambda_1}^{\lambda_2} L_{\text{ap}}(\lambda) \tau_{\text{opt}}(\lambda) A_{\text{ap}} \Omega_{\text{pix}} \text{QE}(\lambda) \frac{\lambda_m}{hc} t_{\text{int}} d\lambda \quad [\text{e-}]$$

Before that integral every radiometric quantity is a spectrum; after it every quantity is a per-pixel scalar in electrons. There is no second collapse and no stage that partially collapses one.

1.1.3 The spatial spine

In parallel with the radiometric spine, the model carries a spatial description of the same system along two paths that must agree. The **PSF path** starts from the complex pupil, forms the point spread function, and convolves onto it every spatial degradation in turn — detector aperture, charge diffusion, interpixel capacitance, jitter, smear, turbulence. Ensquared energy, relative edge response, full width at half maximum, Strehl ratio, and the line and edge responses are computed from that one object. The **MTF product path** computes the optical transfer function directly from the autocorrelation of the same pupil and multiplies it by an analytic or kernel-derived MTF for each downstream contributor:

$$\text{MTF}_{\text{sys}}(\nu) = \prod_i \text{MTF}_i(\nu)$$

MTF budgets, MTF at Nyquist, folded MTF, and the image-quality equation consume this path. Because both paths originate in the same pupil, the Fourier transform of the convolved PSF must reproduce the MTF

product, and the model checks that on every run that computes the spatial path. A failure means a degradation was added to one path and not the other. The two paths and the check are the subject of Chapter 5.

1.1.4 Terms, not transformations

Two stages — platform and calibration — deliberately transform nothing. They contribute *terms*: a kernel and an MTF factor in the platform’s case, noise terms and bias terms in the calibration’s. The calibration stage’s position is itself load-bearing physics. Readout applies the $\sqrt{N}$ averaging of TDI and co-adding to the temporal noise terms; calibration runs after, so its residuals are structurally exempt from that averaging, which is exactly right because correlated errors do not average down.

1.2 The three radiometric regimes

How much of a target’s flux reaches one pixel depends on whether the target fills the pixel, underfills it, or is small enough that the point spread function — not the target — sets the energy distribution. RADIANT classifies every scene into one of three regimes, and the classification changes the signal equation, not merely a label.

Regime	Physical meaning	EE_{box} applied?
Extended	Target fills the pixel footprint completely	No — the pixel sees a uniform radiance field
Sub-pixel	Target smaller than the pixel but larger than the diffraction blur	Yes, to the target term only
Point source	Target smaller than the diffraction limit	Yes, to the target term

What selects the regime. The source stage computes the target’s angular extent $\theta = \sqrt{A_t}/R_s$ [rad] and compares it with the pixel instantaneous field of view $\text{IFOV} = p/f$ [rad]. A declared fill fraction below 1.0 forces the sub-pixel regime; otherwise an explicit override is honored if present; otherwise $\theta \geq 2 \text{ IFOV}$ gives extended, $\theta \leq 0.25 \text{ IFOV}$ gives point source, and anything between is sub-pixel. That classification is *tentative*. The optics stage, which is the first stage that knows the actual PSF size, finalizes it. Every downstream stage reads the final classification and no stage re-derives it.

Where the ensquared energy enters. EE_{box} — the fraction of the degraded PSF’s energy falling inside one pixel — is computed in the platform stage from the *fully* degraded PSF, so jitter, smear, and turbulence are already in it, and it is applied in the spectral-integration stage and nowhere else. It multiplies the target term in the sub-pixel and point-source regimes, never the background term, and never anything in the extended regime, where the pixel is filled by construction and the energy that leaves it is replaced by energy from the neighboring scene.

Getting the regime wrong is the most common cause of an SNR that looks inexplicable: an extended scene treated as a point source loses the factor $EE_{\text{box}} < 1$ for no physical reason, and a point source treated as extended assumes a uniformly filled pixel that does not exist.

1.3 What the model predicts

The metric layer turns the final chain state into the quantities an analyst reports. They fall into five families, each of which can be enabled or disabled per run:

- **Radiometric** — SNR, contrast SNR, signal-to-clutter-plus-noise ratio, NEDT [K], radiometric accuracy [% and K], detection range [m], minimum resolvable temperature at Nyquist [K]. The classical sensitivity relatives — noise-equivalent irradiance, power, specific detectivity, noise-equivalent radiance and reflectance, and the temperature retrieval with its Jacobians — are developed in Chapter 8 as computed relatives of the same chain state.
- **Spatial and MTF** — system MTF at Nyquist, folded MTF, aliased fraction, Strehl ratio, relative edge response, ensquared energy in 1×1 and 3×3 pixel boxes, PSF full width at half maximum [m], per-contributor MTF budgets.
- **Sampling** — ground sample distance [m] cross-track, along-track, and geometric mean; target-plane sample distance [m]; the sampling parameter Q ; swath width [m]; ground range [m]; access rate [m^2/s]; diffraction limit in μrad and in meters.
- **Saturation and dynamic range** — well margin [dB], ADC margin [dB], dynamic range [dB], maximum integration time before saturation [s].
- **Interpretability** — NIIRS via the General Image Quality Equation, with an explicit flag when an input falls outside the equation's calibration envelope.

Every metric is reported with its units, and a metric whose inputs do not support it is reported as a named failure rather than as a silent NaN.

1.4 Model scope — what is deliberately absent

A performance model earns trust by being explicit about its edges. The following are not modeled, and a result that depends on one of them is outside the model's validity rather than merely imprecise.

Scene and source. No molecular band spectroscopy or plume chemistry; no fluorescence; no scene-level interreflection or ray tracing; no temporal scene variability within or between frames; no lunar illumination and no sun-glint geometry. Target emission is Lambertian, and the reflectance structure available is the Lambertian and normalized-Phong BRDF pair rather than a microfacet model or measured BRDF tables.

Atmosphere. No refraction, and therefore no bending of the geometric path; no scintillation; no adjacency effect; no polarization; no cloud or precipitation attenuation. The simple parametric model and the library-backed models are two families with one contract, and their limits are the subject of Chapter 4.

Optics. No ghost images, no bidirectional scatter distribution function beyond the scalar-roughness halo, no chromatic aberration, no field-dependent vignetting map. Aberrations enter as wavefront error in the pupil — either a bulk RMS value or a Zernike set — and their interaction with diffraction is handled in the pupil, never as a separate multiplicative MTF factor.

Detector and readout. No optical crosstalk between pixels, no charge-trapping model, no ADC integral or differential nonlinearity, no electronics-level coupling or supply noise, no cosmic-ray transient or bad-pixel statistics. Blooming, charge-transfer efficiency in CCDs, and radiation-damage accumulation are outside the current model.

Spatial. No geometric distortion, no frame-to-frame registration error, no whiskbroom scan-mechanism transfer function, and no short-exposure turbulence branch — turbulence enters as the long-exposure Kolmogorov result, and is omitted entirely when no Fried parameter is available.

1.5 How to read this manual

Chapter 2 establishes the viewing geometry and sampling on which everything else is conditioned. Chapter 3 is the radiometric core: the governing equations from Planck’s law to the photoelectron count, each with its validity limits, its classic implementation pitfalls, and a pinned numeric anchor. Chapter 4 covers the two atmosphere model families. Chapter 5 is the spatial model, PSF and MTF together. Chapter 6 is the noise taxonomy and its acquisition scaling. Chapter 7 is the calibration error model — what survives a non-uniformity correction, and how it sets a floor that integration time cannot lower. Chapter 8 is the metric layer. The appendix treats the mixed refractive-reflective optical train in full, and the front matter holds the notation on which all of it depends.

Throughout, an equation is stated in the form the model actually evaluates. Where the implementation makes an approximation, the approximation is named where the equation is given, not deferred to a limitations list.

Every quantitative claim in the source documents behind this manual is anchored to the implementation function that evaluates it and to a committed test that pins its value; the repository sources carry those anchors as the project’s traceability record. The typeset manual omits the pointers and keeps the physics.

Geometry & Sampling

Spherical-Earth viewing geometry, orbit kinematics, ground sampling, smear kinematics, solar geometry, and attitude conventions as implemented in RADIANT, which resolves the scene geometry before any radiometry. Numeric anchors were re-derived from the literature independently of the code and then checked against it. RADIANT uses the IUGG **mean Earth radius** $R = 6371.0$ km, so the anchors below are quoted for that radius: an ~8 km radius difference — the gap to the WGS-84 equatorial value — moves a 500 km/30° slant range by ~9 m and the orbit period by ~9 s, so always match radii before comparing numbers.

Symbols and the code’s naming. Classical texts parameterize the viewing triangle by the look angle at the satellite; RADIANT is θ_o -referenced:

Classical symbol	Meaning	RADIANT name
η	look/off-nadir angle at the sensor	<code>eta(viewing_triangle.eta_from_theta_o)</code>
θ_o ($= \theta_i$ for a surface target)	LOS zenith angle at the target	<code>theta_o, path_zenith_rad</code>
$\varepsilon = \pi/2 - \theta_o$	grazing/elevation angle	<code>elevation_angle_rad</code>
Λ	Earth central angle	via ground range $R\Lambda$
R_s	slant range	<code>slant_range_m</code>
h	sensor altitude	<code>geometry.sensor_altitude_m</code>

2.1 The viewing triangle

Equations. Triangle Earth-center–sensor–target with sides R , $R + h$, R_s . Angle closure $\Lambda = \theta_o - \eta$ and the law of sines give the load-bearing identity

$$\sin \theta_o = \frac{R + h}{R} \sin \eta$$

— note the **amplifying** factor $(R + h)/R$: incidence at the target always *exceeds* the look angle on a sphere ($R/(R + h)$ belongs only in the $\varepsilon \rightarrow \eta$ direction). Closed-form solution sets exist from any one of $\{\eta, \theta_o, \Lambda\}$ plus h ; the Λ -parameterized form uses the branch-safe $\eta = \text{atan2}(R \sin \Lambda, (R + h) - R \cos \Lambda)$.

Assumptions & validity. Spherical Earth, surface (or specified-altitude) target, no refraction (refraction lifts apparent elevation up to ~0.5° at $\varepsilon \approx 0$).

Pitfalls. The flat-Earth intuition $\theta_o = \eta$ (wrong by 2.6° at $h = 500$ km, $\eta = 30^\circ$ and unbounded toward the limb); $\Lambda = \eta - \theta_o$ sign flips; degrees fed to radian trig.

Numeric anchor. $h = 500$ km, $\eta = 30^\circ$: $\theta_o = 32.6319^\circ$, $\Lambda = 2.6319^\circ$ (mean-R).

2.2 Slant range

Equation. Direct form in the look angle:

$$R_s = (R + h) \cos \eta - \sqrt{R^2 - (R + h)^2 \sin^2 \eta}.$$

The **minus** root is the near (visible-surface) intersection; the plus root exits the far side of the Earth. A negative discriminant means the LOS misses the Earth — $\sin \eta > R/(R+h)$, beyond the limb; the boundary is $\eta_{max} = \arcsin \frac{R}{R+h}$ with tangent-ray range $\sqrt{2Rh + h^2}$.

Pitfalls. Plus-root selection (returns ~11 Mm at nadir instead of h); silent NaN on a beyond-horizon geometry instead of an actionable error; near η_{max} , $dR_s/d\eta \rightarrow \infty$ — solvers should re-parameterize in Λ or ε there.

Numeric anchors. $h = 500$ km, $\eta = 30^\circ$: $R_s = 585,110.5$ m (mean-R). Nadir: $R_s = h$ exactly. $\eta_{max}(500 \text{ km}) = 68.01^\circ$ (mean-R).

2.3 Circular-orbit kinematics

Equations.

$$v = \sqrt{\frac{\mu}{R+h}}, \quad T = 2\pi \sqrt{\frac{(R+h)^3}{\mu}}, \quad v_g = v \frac{R}{R+h}.$$

The $R/(R+h)$ factor: satellite and nadir point share one angular rate ω ; linear speed is ω times each circle's radius. It is a projection of angular motion, not a velocity-vector projection.

Assumptions & validity. Two-body, circular, no J_2 (~0.1 % LEO period effect); non-rotating Earth for v_g — Earth rotation adds up to $\pm 465 \cos(\text{lat})$ m/s (~6.6 % equatorial) to the true relative ground speed, unmodeled.

Pitfalls. $\sqrt{\mu/R}$ instead of $\sqrt{\mu/(R+h)}$ (3.8 % at 500 km); **orbital** v **where ground** v_g **belongs** in smear/line-rate math — a +7.8 % error at 500 km that looks plausible; μ in km³/s² mixed with meters.

Numeric anchors (mean-R): $v = 7616.6$ m/s, $T = 94.469$ min at $h = 500$ km.

2.4 GSD — nadir and off-nadir

Equations. IFOV = p/f ; nadir GSD = $(p/f) h$. Off-nadir, the beam width IFOV $\cdot R_s$ (measured $\perp$ to the LOS) projects onto the local tangent plane; the LOS makes angle θ_o (**incidence at the target, not the look angle** η) with the surface normal:

$$\text{GSD}_{\text{in-plane}} = \frac{\text{IFOV} \cdot R_s}{\cos \theta_o}, \quad \text{GSD}_{\perp} = \text{IFOV} \cdot R_s.$$

The $1/\cos\theta_o$ applies **only** in the tilt plane (the plane containing the LOS and the local vertical); the orthogonal horizontal direction is itself $\perp$ to the LOS and projects length-preserving. In RADIANT's convention the LOS tilt is carried by `geometry.path_zenith_rad` in the **along-track** plane, so `along_track_m` receives the $1/\cos\theta_o$ elongation and `cross_track_m` the range-only factor. The formula is exact on the sphere: differentiating the viewing triangle gives $R d\Lambda/d\eta = R_s/\cos\theta_o$ identically.

Pitfalls. $\cos\eta$ for $\cos\theta_o$ (2.75 % at the anchor geometry, unbounded toward the limb); $\cos$ vs $1/\cos$ (GSD must *grow* off-nadir); applying the obliquity factor to both directions (a further ~19 % area error at the anchor); quoting one off-nadir “GSD” without naming the direction.

Numeric anchors (mean-R, $p = 10\ \mu\text{m}$, $f = 2\ \text{m}$, $h = 500\ \text{km}$): nadir 2.500000 m; $\eta = 30^\circ$: in-plane 3.473901 m — the wrong-angle value 3.378 m is explicitly rejected by the anchor test.

2.5 Swath and access

Equations. Swath from the cross-track GSD and detector format: $W = \text{GSD}_{\text{cross}} \cdot N_{\text{pix}}$; area access rate $\dot{A} = W v_g$; ground range from the central angle, $R\Lambda$. Wide-FOV swath uses the full spherical mapping $\Lambda(\eta)$ — the flat-Earth $2h \tan \eta_{\text{half}}$ is 0.3 % low at $\pm 15^\circ$ and diverges beyond $\sim 30^\circ$.

2.6 Smear kinematics and TDI line-rate matching

Equations. Uncompensated nadir pushbroom: ground motion $v_g t_{\text{int}}$ maps to the focal plane through the magnification f/R_s :

$$d_{\text{img}} = v_g t_{\text{int}} \frac{f}{R_s} \quad (= v_g t_{\text{int}} f/h \text{ at nadir}).$$

TDI requires the line clock to match the image velocity: $p f_{\text{line}} = v_{\text{img}} \iff t_{\text{line}} = \text{GSD}/v_g$; residual per-stage mismatch smear multiplies by N stages. The MTF consequence of d_{img} is the smear sinc of the spatial chapter's *Platform kernels: jitter and smear*.

Pitfalls. Orbital v for v_g (+7.8 %); f/R or $f/(R+h)$ for the magnification; ground meters compared to focal-plane microns without f/R_s ; conflating t_{int} with t_{line} in TDI ($t_{\text{int}} = N t_{\text{line}}$; the matching condition constrains t_{line}).

Numeric anchor. $h = 500\ \text{km}$, $p = 10\ \mu\text{m}$, $f = 2\ \text{m}$, $t_{\text{int}} = 1\ \text{ms}$: $d_{\text{img}} \approx 2.8$ pixels — uncompensated millisecond integration is not viable; the matched line time is $\sim 354\ \mu\text{s}$.

2.7 Solar geometry

Equation. Spherical law of cosines on the astronomical triangle:

$$\cos \theta_z = \sin \phi \sin \delta + \cos \phi \cos \delta \cos H$$

(ϕ latitude, δ declination, H hour angle from solar noon). Declination comes from the day of year, and the hour angle from either the local solar time or the local time of the ascending node.

Pitfalls. Clock time vs apparent solar time (equation of time, $\pm 4^\circ$ in H); elevation returned where zenith is expected (downstream $\cos \theta_z$ irradiance factors get sin/cos swapped); sign conventions on δ across hemispheres.

Numeric anchor. $\phi = 35^\circ\text{N}$, $\delta = 23.44^\circ$, $H = 15^\circ$: $\theta_z = 17.4256^\circ$. Noon collapses to $\theta_z = \phi - \delta$ exactly.

2.8 Attitude: Euler ZYX

Equation. Intrinsic z-y'-x'' (yaw $\psi \rightarrow$ pitch $\theta \rightarrow$ roll φ), active, right-handed, column vectors:

$$\mathbf{R} = \mathbf{R}_z(\psi) \mathbf{R}_y(\theta) \mathbf{R}_x(\varphi),$$

gimbal lock at $\theta = \pm 90^\circ$. This matches `scipy Rotation.from_euler('ZYX', ...)` exactly; extraction via $\theta = -\arcsin R_{31}$, $\psi = \text{atan2}(R_{21}, R_{11})$, $\varphi = \text{atan2}(R_{32}, R_{33})$.

Pitfalls. Intrinsic z-y'-x'' equals *extrinsic* x-y-z (same product) but not intrinsic x-y'-z''; active/passive transposes; the $+\sin \theta$ corner of $\mathbf{R}_y$ sits opposite to $\mathbf{R}_z/\mathbf{R}_x$.

2.9 Sampling on the ground

Focal-plane Nyquist $1/(2p)$ projects to ground Nyquist $1/(2\text{GSD})$ — direction-dependent off-nadir (use the direction's GSD from *GSD — nadir and off-nadir* above). The optics-vs-sampling budget ($Q = \lambda F_\# / p$, aliasing, folded MTF) lives in the spatial chapter's *Sampling: Nyquist, Q, and folded MTF*; the geometry chapter's contribution is the GSD that scales it to the ground.

Radiometric Signal Chain

The radiometry foundations of RADIANT: the governing equations that carry a photon from source emission to a photoelectron count, each stated with its derivation, validity limits, classic implementation pitfalls, and a pinned numeric anchor. This chapter covers the radiometric quantities themselves; the spatial (PSF/MTF) treatment is the Spatial Model chapter and the noise taxonomy is the Noise Model chapter.

Every numeric anchor below was independently re-derived from the physics literature, without reference to the implementation, and then verified against the implementation.

3.1 Notation

Symbols and canonical units are defined once for the whole manual in the front matter, Notation and Symbols; this chapter uses them without redefinition. RADIANT canonical units throughout: wavelength in μm , angles in radians, lengths in meters, spectral radiance in $\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$.

Where a formula needs λ in meters (anything with h , c , k_B in it), we write λ_m explicitly. The μm -to-m bookkeeping is the single most common bug source in this entire subject; every entry below calls out where the 10^{-6} lives.

3.2 The chain at a glance

RADIANT models the end-to-end signal chain as ten sequential stages, the first of which resolves the scene geometry. Each stage is a pure function that transforms an immutable `ChainState`, adding radiometric frames, noise terms, and MTF contributions:

```
GeometryStage → SourceStage → AtmosphereStage → OpticsStage → PlatformStage  
→ SpectralIntegrationStage → DetectorStage → ReadoutStage → CalibrationStage  
→ PerformanceStage
```

The final `ChainState` contains everything needed to compute performance metrics (SNR, NEDT, NIIRS, MTF). The foundations in this chapter are the radiometric content of that chain: Geometry establishes ranges and angles; Source builds $L(\lambda)$ from the equations below; Atmosphere applies τ_{atm} and adds path radiance; Optics applies throughput and finalizes the radiometric regime; SpectralIntegration collapses spectra to photoelectrons using the photon-conversion entry. A stage-by-stage map is given in [How the Foundations Feed the Chain](#) at the end.

3.3 Foundations

3.3.1 Planck spectral radiance

Equation.

$$B(\lambda, T) = \frac{2hc^2}{\lambda_m^5} \cdot \frac{1}{e^x - 1} \cdot 10^{-6}, \quad x = \frac{hc}{\lambda_m k_B T}$$

The first factor evaluates in W/m²/sr/m; the trailing 10⁻⁶ is the Jacobian $d\lambda_m/d\lambda_{\mu m}$ that converts to the canonical W/m²/sr/μm.

Symbols. λ_m [m] wavelength in meters; T [K]; h [J·s]; c [m/s]; k_B [J/K]; x dimensionless; output B [W/m²/sr/μm].

Derivation. Planck's law follows from Bose–Einstein statistics of photon modes in a cavity: the mode density per unit volume per unit frequency is $8\pi\nu^2/c^3$, each mode carries mean energy $h\nu/(e^{h\nu/k_B T} - 1)$, and converting energy density to radiance (multiply by $c/4\pi$) and changing variables from ν to λ (Jacobian c/λ^2) gives the per-meter form $2hc^2/\lambda_m^5 \cdot (e^x - 1)^{-1}$. A per-μm spectral interval is 10⁻⁶ of a per-m interval, hence the final factor. The dimensionless argument can equivalently be written $x = c_2/(\lambda_m T)$ with the second radiation constant $c_2 = hc/k_B$.

Assumptions & validity. Thermal-equilibrium emission into vacuum (refractive index $n = 1$; in a medium B scales as n^2). Unpolarized, Lambertian (direction-independent) radiance. Valid for all $\lambda, T > 0$. Numerically, $x \gtrsim 700$ overflows $\exp()$ in float64; RADIANT truncates the Wien tail to exactly zero there (the true value is below float64 underflow, so no physics is lost). $T = 0$ returns exactly zero radiance.

Pitfalls.

- The per-μm 10⁻⁶ applied zero times or twice — results off by 10⁶ or 10¹².
- Using λ in μm inside λ^5 without conversion — a 10³⁰ scale error partially masked by the 10⁻⁶ Jacobian.
- $\exp(x) - 1$ instead of $\expm1(x)$: catastrophic cancellation for $x \ll 1$ (Rayleigh–Jeans regime). RADIANT uses `numpy.expm1`.
- Confusing the radiance constant $c_{1L} = 2hc^2$ with the exitance constant $c_1 = 2\pi hc^2$ — a stray factor of π . Radiance B and exitance $M = \pi B$ are different quantities.

Numeric anchor. $B(10 \mu\text{m}, 300 \text{ K}) = 9.92403333 \text{ W/m}^2/\text{sr}/\mu\text{m}$ ($x = 4.796$, near the LWIR peak of a 300 K blackbody at $\lambda_{\text{max}} = 2898 \mu\text{m}\cdot\text{K}/300 \text{ K} \approx 9.66 \mu\text{m}$). Also $B(4 \mu\text{m}, 300 \text{ K}) = 0.721976423 \text{ W/m}^2/\text{sr}/\mu\text{m}$.

References. [Planck 1901], [Siegel & Howell §1.6], [CODATA 2018].

3.3.2 Temperature derivative of Planck radiance — the NEDT kernel

Equation.

$$\frac{\partial B}{\partial T} = B(\lambda, T) \cdot \frac{x}{T} \cdot \frac{e^x}{e^x - 1} \quad \left[\text{W/m}^2/\text{sr}/\mu\text{m/K} \right]$$

Symbols. As in the Planck entry; output in $\text{W/m}^2/\text{sr}/\mu\text{m}/\text{K}$. The fully expanded form is $\partial B/\partial T = (2h^2c^3 \cdot 10^{-6})/(\lambda_m^6 k_B T^2) \cdot e^x/(e^x - 1)^2$.

Derivation. Write $B = A/(e^x - 1)$ with $A = 2hc^2/\lambda_m^5$ fixed in T . Since $x = hc/(\lambda_m k_B T)$, $dx/dT = -x/T$. Chain rule: $dB/dT = A e^x (e^x - 1)^{-2} \cdot (x/T) = B \cdot (x/T) \cdot e^x/(e^x - 1)$. This is the kernel of every NEDT computation: a scene temperature change ΔT produces a radiance change $\Delta L \approx (\partial B/\partial T) \Delta T$.

Assumptions & validity. The derivative itself is exact — no linearization. Using it as the NEDT kernel assumes ΔT small enough that B is locally linear: excellent for $\Delta T \lesssim$ a few K at terrestrial temperatures, degrading in the Wien regime where the fractional error of the linearization grows like $(x/2)(\Delta T/T)$. Undefined at $T = 0$ (RADIANT raises).

Pitfalls.

- Dropping the $e^x/(e^x - 1)$ factor (writing $\partial B/\partial T = B \cdot x/T$) — a Wien-limit shortcut that is $\sim 0.8\%$ low at $10\ \mu\text{m}/300\ \text{K}$ and silently regime-dependent.
- Carrying $(e^x - 1)^1$ where the expanded form needs $(e^x - 1)^2$.
- Sign check: $\partial B/\partial T > 0$ for every λ, T . A negative value anywhere is a bug.
- All the per- μm 10^{-6} hazards of the Planck entry, now propagating straight into NEDT.

A physical sanity check worth internalizing: at $300\ \text{K}$ the MWIR derivative is $\sim 5.5\times$ smaller than LWIR in absolute per- μm terms, but its *fractional* contrast $(\partial B/\partial T)/B$ is larger ($4.00\ \%/K$ at $4\ \mu\text{m}$ vs $1.61\ \%/K$ at $10\ \mu\text{m}$) — the classic MWIR-vs-LWIR contrast trade.

Numeric anchor. $\partial B/\partial T (10\ \mu\text{m}, 300\ \text{K}) = 0.159971567\ \text{W/m}^2/\text{sr}/\mu\text{m}/\text{K}$ (audit value, verified against central finite differences to $\leq 3 \times 10^{-10}$ relative).

References. [Planck 1901], [Holst 2008], [CODATA 2018].

3.3.3 Band-integrated and band-averaged radiance

Equation.

$$L_{\text{band}} = \int_{\lambda_1}^{\lambda_2} L(\lambda) d\lambda \approx \sum_{i=1}^{N-1} \frac{1}{2} (L_i + L_{i+1}) (\lambda_{i+1} - \lambda_i) \quad [\text{W/m}^2/\text{sr}]$$

$$\bar{L} = \frac{L_{\text{band}}}{\lambda_2 - \lambda_1} \quad [\text{W/m}^2/\text{sr}/\mu\text{m}]$$

Symbols. λ_1, λ_2 [μm] band edges; L_i [$\text{W/m}^2/\text{sr}/\mu\text{m}$] samples on the spectral grid; grid spacing in μm .

Derivation. Integrating a per- μm spectral density over a wavelength interval expressed in μm yields $\text{W/m}^2/\text{sr}$ with *no* extra conversion factor — the units cancel by construction. The trapezoid rule is the RADIANT quadrature (via `numpy.trapezoid`); its general form handles non-uniform grids correctly because the spacing enters per interval. The band average divides by the full span $\lambda_2 - \lambda_1$; if a relative spectral response $R(\lambda)$ is involved, the correct band average is instead $\int L R d\lambda / \int R d\lambda$.

Assumptions & validity. Trapezoid accuracy requires the grid to resolve the integrand: trivially satisfied for a smooth blackbody (201 points across a band gives $< 10^{-6}$ relative error in LWIR/MWIR), but for

atmospherically filtered radiance the grid must resolve the $\tau(\lambda)$ line structure or the result is quadrature-limited, not physics-limited.

Pitfalls.

- $\Delta\lambda$ in the wrong units (m or nm instead of μm) — 10^{-6} or 10^3 scale errors on the integral.
- Dividing by the sample count N instead of the wavelength span when band-averaging.
- Applying the band average *and* multiplying by $\Delta\lambda$ somewhere downstream — integrating twice.
- Treating the band as a top-hat when a real spectral response weighting exists.
- Coarse-grid bias: even trapezoid on a 5-point grid across 8–12 μm is 0.4 % low; a left-Riemann sum is worse.

Numeric anchor. $\int_8^{12} B(\lambda, 300 \text{ K}) d\lambda = 38.5004239 \text{ W/m}^2/\text{sr}$ (adaptive-quadrature truth; band average 9.62510598 W/m²/sr/ μm). Sanity: that is 26.3 % of the total $\sigma T^4/\pi = 146.15 \text{ W/m}^2/\text{sr}$, consistent with blackbody band-fraction tables.

References. [Siegel & Howell §1.6], [Holst 2008].

3.3.4 Brightness temperature — spectral (monochromatic) inversion

Equation. Given spectral radiance L [W/m²/sr/ μm] at wavelength λ , the brightness temperature is the closed-form inverse of Planck:

$$T_B = \frac{hc}{\lambda_m k_B} \cdot \frac{1}{\ln\left(1 + \frac{2hc^2 \cdot 10^{-6}}{\lambda_m^5 L}\right)} \quad [\text{K}]$$

Symbols. λ_m [m]; L [W/m²/sr/ μm]; the prefactor $hc/(\lambda_m k_B) = c_2/\lambda_m$ [K]; the log argument is dimensionless (both terms in W/m²/sr/ μm at that λ).

Derivation. Solve $L = (2hc^2 \cdot 10^{-6}/\lambda_m^5)/(e^x - 1)$ for x : $e^x - 1 = 2hc^2 \cdot 10^{-6}/(\lambda_m^5 L)$, so $x = \ln(1 + r)$ with r the radiance ratio, and $T_B = hc/(\lambda_m k_B x)$. Because $B(\lambda, T)$ is strictly monotonic in T at fixed λ , the inverse is unique. For a blackbody, T_B at every wavelength equals the physical temperature; λ -dependence of T_B diagnoses a graybody or selective emitter.

Assumptions & validity. Blackbody spectral shape at that wavelength ($\varepsilon = 1$); a graybody's brightness temperature sits *below* its thermodynamic temperature. Monochromatic: L must be the spectral radiance at λ , or a band narrow enough that L is flat across it. Defined for any $L > 0$.

Pitfalls.

- The 10^{-6} inside the log: L is per- μm , so the $2hc^2/\lambda_m^5$ term must be converted to per- μm before forming the ratio — omitting it is a 10^6 error in the log argument.
- Dropping the “1+” (Wien approximation) — fine for $x \gg 1$, but a hidden ~1 K bias near the band peak (at 10 μm /300 K, $x \approx 4.8$).
- $\log(1 + r)$ evaluated naively for tiny r (very hot / Rayleigh–Jeans cases) — use $\log1p$.
- λ in μm inside λ^5 (same as the Planck entry).

Numeric anchor. $T_B(10\ \mu\text{m}, L = 9.0\ \text{W/m}^2/\text{sr}/\mu\text{m}) = 294.054730\ \text{K}$; re-forwarding gives $B(10\ \mu\text{m}, 294.054730\ \text{K}) = 9.00000000\ \text{W/m}^2/\text{sr}/\mu\text{m}$ (round-trip exact to 10 digits).

References. [Planck 1901], [NIST ITS-90], [Holst 2008].

3.3.5 Band radiance temperature — band-integrated inversion

Equation. Given a measured band radiance L_{meas} [W/m²/sr] over $[\lambda_1, \lambda_2]$, the band radiance temperature T_R solves

$$F(T_R) \equiv \int_{\lambda_1}^{\lambda_2} B(\lambda, T_R) d\lambda - L_{\text{meas}} = 0$$

No closed form exists; solve numerically (RADIANT uses Brent’s method on a bracketing interval).

Symbols. L_{meas} [W/m²/sr]; T_R [K]; band edges in μm .

Derivation. Existence and uniqueness follow from monotonicity: $\partial B/\partial T > 0$ for every λ, T (previous entry, strictly positive), so $F'(T) = \int \partial B/\partial T d\lambda > 0$ — F is strictly increasing. $F(T \rightarrow 0^+) = -L_{\text{meas}} < 0$ and $F(T \rightarrow \infty) \rightarrow +\infty$, so a strictly increasing continuous function with a sign change has exactly one root. This licenses bracketing solvers with a guaranteed single solution — which is exactly why RADIANT validates the bracket before calling Brent.

Assumptions & validity. Blackbody spectral *shape* within the band; a graybody or atmospherically filtered scene yields an *effective* T_R , not the physical temperature. With a spectral response $R(\lambda) \geq 0$ the same monotonicity argument holds. The quadrature grid used inside the solver must match the one used to produce L_{meas} , or a quadrature mismatch masquerades as a temperature bias (a 0.15 % radiance error is ~0.1 K at LWIR).

Pitfalls.

- Inverting the *band-averaged* radiance through the *monochromatic* Planck at band center — close for narrow bands, a systematic bias for wide bands (Planck is nonlinear across the band).
- Solver bracket too narrow (e.g. [200, 400] K fails on hot targets). RADIANT brackets [1, 10,000] K and raises an actionable error if L_{meas} falls outside the image of the bracket.
- Comparing per- μm to band-integrated units (W/m²/sr/ μm vs W/m²/sr) — a $(\lambda_2 - \lambda_1)$ factor error.
- Confusing tolerance on radiance with tolerance on temperature: at 3–5 μm , $dL_{\text{band}}/dT \approx 4\ \%/K$ of L_{band} , so a 1 % radiance tolerance is ~0.25 K.

Numeric anchor. For 8–12 μm and $L_{\text{meas}} = 30.0\ \text{W/m}^2/\text{sr}$: $T_R = 285.450727\ \text{K}$ (re-forward reproduces 30.0000000 W/m²/sr). The self-consistency case $L_{\text{meas}} = 38.5004239\ \text{W/m}^2/\text{sr}$ recovers exactly 300.000000 K.

References. [Siegel & Howell §1.6], [Press et al. §9.3] (Brent’s method).

3.3.6 Graybody emission (scene targets)

Equation.

$$L_{\text{emit}}(\lambda, T) = \varepsilon(\lambda) \cdot B(\lambda, T) \quad \left[\text{W/m}^2/\text{sr}/\mu\text{m} \right], \quad 0 \leq \varepsilon \leq 1$$

The complete at-surface leaving radiance in the thermal IR adds the reflected downwelling term:

$$L_{\text{leave}}(\lambda) = \varepsilon(\lambda) B(\lambda, T_{\text{surf}}) + [1 - \varepsilon(\lambda)] L_{\text{down}}(\lambda)$$

Symbols. $\varepsilon(\lambda)$ dimensionless emissivity (scalar graybody or spectral); T_{surf} [K]; L_{down} [W/m²/sr/μm] downwelling sky radiance.

Derivation. Kirchhoff's law in local thermodynamic equilibrium equates spectral directional absorptance and emissivity, $\alpha(\lambda) = \varepsilon(\lambda)$. Energy balance for an opaque surface ($\tau = 0$) gives $\rho + \alpha = 1$, hence $\rho = 1 - \varepsilon$: the emission and sky-reflection terms are complementary. High- ε surfaces emit; low- ε surfaces mirror the sky.

For **scene targets and backgrounds**, ε is a legitimate *independent* material input — it describes what the material is, and RADIANT accepts it directly (contrast with optical elements, next entry).

Assumptions & validity. Local thermodynamic equilibrium (fails only for luminescent/lasing media — not passive-EO scenes). Strictly, Kirchhoff holds per direction: $\varepsilon(\lambda, \theta, \phi) = \alpha(\lambda, \theta, \phi)$; using a hemispherical ε for directional radiance assumes a diffuse emitter (water at grazing angles is the classic violator). ε is assumed independent of T over the range of interest. RADIANT validates $\varepsilon \in [0, 1]$ and refuses to extrapolate a spectral emissivity table outside its wavelength range.

Pitfalls.

- Applying $\rho + \varepsilon = 1$ across bands — pairing a *solar-band* ρ with a *thermal-band* ε . Kirchhoff is spectral; visible albedo says nothing about 10 μm emissivity.
- Forgetting the reflected-downwelling term: for $\varepsilon = 0.95$ at LWIR it is a ~5 % correction, and omitting it biases retrieved brightness temperatures.
- Allowing $\varepsilon > 1$ or $\varepsilon + \rho > 1$ through unvalidated inputs.

Numeric anchor. $\varepsilon = 0.95$ at $\lambda = 10 \mu\text{m}$, $T = 300 \text{ K}$: $L = 0.95 \times 9.92403333 = 9.42783166 \text{ W/m}^2/\text{sr}/\mu\text{m}$. Reflected-sky term for $\rho = 0.05$, $L_{\text{down}} = 5.0 \text{ W/m}^2/\text{sr}/\mu\text{m}$: $0.250000000 \text{ W/m}^2/\text{sr}/\mu\text{m}$ (2.6 % of the emitted term).

References. [Kirchhoff 1860], [Siegel & Howell §3], [Holst 2008].

3.3.7 Kirchhoff-derived emissivity of optical elements

Equation. For any optical element *inside the sensor*, emissivity is never an input — it is derived:

$$\text{mirrors: } \varepsilon(\lambda) = 1 - R(\lambda) \quad \text{transmissive elements: } \varepsilon(\lambda) = 1 - T(\lambda) - R(\lambda)$$

Symbols. $R(\lambda)$ reflectance, $T(\lambda)$ transmittance, $\varepsilon(\lambda)$ emissivity — all dimensionless, all per wavelength, constrained by $T + R \leq 1$.

Derivation. Same physics as the graybody entry — Kirchhoff plus energy balance — but the *architectural* consequence differs. A scene target's ε is a material property the user legitimately knows. An optical element's ε , R , and T are locked together by energy conservation: a user who specifies both R and ε for a mirror has over-specified the energy balance, and the two claims can contradict. RADIANT therefore accepts R (and T where applicable) and *derives* ε ; the derived value then drives the warm-optics self-emission term $L_{\text{nf}}(\lambda) = \varepsilon_{\text{opt}}(\lambda) B(\lambda, T_{\text{opt}})$ that the optics stage adds to the background.

For a refractive element modeled with the cavity model, the generalized form $\varepsilon_{\text{eff}} = T_2 n^2 (1 - \beta) / D$ applies (n^2 enhancement for emission inside a dielectric medium); for a simple refractive element with no bulk absorption model, $\varepsilon = 0$ — the residual $1 - T$ is predominantly reflection, and assigning it to emission would fabricate thermal background.

Assumptions & validity. Same LTE assumptions as Kirchhoff generally. Per-wavelength: the constraint is enforced on the full spectral grid, not on band averages.

Pitfalls.

- Accepting ε as an independent parameter for an optical surface — an over-specification of the energy balance. RADIANT raises KirchhoffViolationError instead.
- Applying $\varepsilon = 1 - R$ to a *transparent* element (missing the T term).
- Letting $T + R > 1$ slip through unvalidated — RADIANT enforces $T + R \leq 1$ within tolerance and raises otherwise.
- Assigning the whole $1 - T$ of a lens to emission when most of it is reflection — over-predicts warm-optics background.

Numeric anchor. Gold mirror with $R = 0.98$ (flat): $\varepsilon = 1 - 0.98 = 0.02$ (dimensionless), and its self-emission at $T_{\text{opt}} = 300$ K, $\lambda = 10$ μm is $0.02 \times 9.92403333 = 0.198480667$ W/m²/sr/ μm .

References. [Kirchhoff 1860], [Wolfe & Zissis §5].

3.3.8 Top-of-atmosphere solar irradiance

Equation.

$$E_{\text{sun}}(\lambda) = \pi B(\lambda, T_{\text{sun}}) \left(\frac{R_{\text{sun}}}{d} \right)^2 k_{\text{scale}} \quad \left[\text{W/m}^2/\mu\text{m} \right]$$

with distance scaling $E(d) = E(1 \text{ AU})/d_{\text{AU}}^2$.

Symbols. $T_{\text{sun}} = 5778$ K effective photospheric temperature; $R_{\text{sun}} = 6.957 \times 10^8$ m; d Sun–target distance [m] (d_{AU} in AU); $(R_{\text{sun}}/d)^2 \approx 2.18 \times 10^{-5}$ dimensionless geometric dilution; k_{scale} the flat calibration factor that makes $\int E_{\text{sun}} d\lambda = S_0 = 1361$ W/m² exactly.

Derivation. The Sun subtends $\Omega_{\text{sun}} = \pi(R_{\text{sun}}/d)^2$ at the observer (small-angle; the π here is the projected-solid-angle integral of a uniform disk). Irradiance on a normal plane is radiance $\times$ solid angle: $E = \Omega_{\text{sun}} B(\lambda, T_{\text{sun}})$, which regroups into the form above. Because a 5778 K blackbody with IAU nominal

R_{sun}/AU integrates to $\sim 1365 \text{ W/m}^2$ rather than the observed $S_0 = 1361 \text{ W/m}^2$, RADIANT applies the flat k_{scale} correction: the spectral *shape* stays exactly Planck while the integral recovers S_0 to machine precision. Distance scaling is conservation of intensity: $E \propto 1/d^2$; Earth's seasonal range $d = 0.9833\text{--}1.0167 \text{ AU}$ gives a $\pm 3.4 \%$ annual swing.

Assumptions & validity. Blackbody spectral shape — real solar spectra deviate at the few-percent level (Fraunhofer lines, UV/VIS deficit). A tabulated `astm_e490` model is planned to replace the shape without breaking the API. Small-angle solid angle is excellent ($\Omega_{\text{sun}} \approx 6.8 \times 10^{-5} \text{ sr}$).

Pitfalls.

- Confusing irradiance with radiance: E_{sun} is $\text{W/m}^2/\mu\text{m}$; the equivalent-Lambertian radiance used in single-scatter path-radiance formulas is E_{sun}/π [$\text{W/m}^2/\text{sr}/\mu\text{m}$].
- Applying $1/d^2$ twice (once in the irradiance table, once in code) or not at all.
- Using the solar radius/distance ratio unsquared.

Numeric anchor. An uncalibrated 5778 K blackbody Sun gives $E(0.55 \mu\text{m}) = 1756.84 \text{ W/m}^2/\mu\text{m}$ at 1 AU ($\Omega_{\text{sun}} = 6.79427 \times 10^{-5} \text{ sr}$). The flat calibration multiplies that by $k_{\text{scale}} = 0.99574$, so RADIANT's shipped model returns $1749.35 \text{ W/m}^2/\mu\text{m}$ — within 7 % of the measured $\approx 1870 \text{ W/m}^2/\mu\text{m}$ [ASTM E490], the real Sun exceeding the blackbody near $0.55 \mu\text{m}$. The calibrated model integrates to $S_0 = 1361.0 \text{ W/m}^2$ over $0.05\text{--}50 \mu\text{m}$ to $\sim 10^{-6}$ relative.

A note on the temperature, because two values circulate: the IAU nominal solar effective temperature is 5772 K, while 5778 K is the older value that pairs with the nominal radius and 1 AU to reproduce S_0 most closely. RADIANT uses **5778 K** (`blackbody_5778`). The choice moves $E(0.55 \mu\text{m})$ by 0.5 % before calibration and *nothing at all* after it in integrated terms, since k_{scale} absorbs the difference — it is a shape parameter, not an energy one.

References. [Kopp & Lean 2011] ($S_0 = 1361 \text{ W/m}^2$), [ASTM E490], [Wehrli 1985].

3.3.9 Reflected solar radiance

Equation.

$$L_{\text{refl}}(\lambda) = f_r(\lambda; \theta_{\text{sun}}, \theta_{\text{obs}}) \cdot E_{\text{sun}}(\lambda, d) \cdot \cos \theta_{\text{sun}} \quad \left[\text{W/m}^2/\text{sr}/\mu\text{m} \right]$$

For the Lambertian special case $f_r = \rho/\pi$:

$$L_{\text{refl}}(\lambda) = \frac{\rho(\lambda) E_{\text{sun}}(\lambda) \cos \theta_{\text{sun}}}{\pi}$$

Symbols. f_r [sr^{-1}] BRDF; E_{sun} [$\text{W/m}^2/\mu\text{m}$] solar spectral irradiance on a plane normal to the rays at the target (exo-atmospheric if τ_{atm} is applied separately downstream — which is RADIANT's convention); θ_{sun} [rad] solar zenith angle; $\cos \theta_{\text{sun}}$ projects normal-incidence irradiance onto the tilted surface; d [AU] Sun–target distance.

Derivation. By definition of BRDF, $dL_o = f_r dE_i$: outgoing radiance is BRDF times incident irradiance *on the surface*, which is the normal-plane irradiance times $\cos \theta_{\text{sun}}$. For a Lambertian surface, reflected exitance

is $M = \rho E \cos \theta_{\text{sun}}$ and a perfectly diffuse surface has $M = \pi L$, giving the $1/\pi$ [sr⁻¹]. The Sun is treated as a single collimated directional source (no hemisphere integration) — v1 convention.

Assumptions & validity. BRDF model fidelity is the limit — Lambertian is the baseline, Phong adds a glint lobe (next entry). Collimated-sun is fine except in near-terminator penumbra. $\cos \theta_{\text{sun}} \geq 0$ required: RADIANT clamps sun-below-horizon to zero radiance *with intent* (returns zeros, never silent negative radiance) and rejects zenith angles outside $[0, \pi/2]$ at construction.

Pitfalls.

- Missing $1/\pi$ (returns an exitance-like value, high by $\pi \approx 3.14$) or dividing by 2π (confusing hemisphere solid angle 2π with the projected-solid-angle integral $\int \cos \theta d\Omega = \pi$).
- Omitting $\cos \theta_{\text{sun}}$, or using *elevation* where *zenith* is expected (a $\cos \leftrightarrow \sin$ swap).
- Degrees passed to $\cos()$ expecting radians.
- Applying a $\cos \theta_{\text{view}}$ as well — for a Lambertian surface the radiance has no view-angle cosine (the $\cos \theta_v$ in received power exactly cancels the $1/\cos \theta_v$ in projected source area).
- Double-counting the atmosphere: E_{sun} must be exo-atmospheric if τ_{atm} is applied separately (RADIANT's case), at-surface if not.

Numeric anchor. With the audit's adopted $E_0(0.55 \mu\text{m}) = 1870 \text{ W/m}^2/\mu\text{m}$ (exo-atmospheric, $\tau = 1$), $\rho = 0.3$, $\theta_{\text{sun}} = 30^\circ$: $L = 154.647755 \text{ W/m}^2/\text{sr}/\mu\text{m}$. This anchor scales linearly with the adopted solar table — RADIANT's `blackbody_5778` model gives a few-percent-lower E_0 at $0.55 \mu\text{m}$ (previous entry), so rescale by $E_{0,\text{impl}}/1870$ before comparing implementation output against it.

References. [Nicodemus 1977], [Holst 2008], [ASTM E490].

3.3.10 BRDF normalization — Lambertian and Phong

Equation. BRDF definition: $f_r(\omega_i, \omega_o) = dL_o/dE_i$ [sr⁻¹]. Energy conservation requires $\int_{\text{hemi}} f_r \cos \theta_o d\Omega_o \leq 1$ for every incidence direction. RADIANT implements both of:

$$\text{Lambertian: } f_r = \frac{\rho}{\pi} \quad \text{Phong: } f_r = \frac{\rho_d}{\pi} + \rho_s \cdot \frac{n+2}{2\pi} \cdot \cos^n \alpha$$

Symbols. ρ [—] total hemispherical reflectance; ρ_d, ρ_s [—] diffuse and specular shares with $\rho_d + \rho_s = \rho$; $n \geq 0$ [—] Phong exponent (higher = narrower lobe); α [rad] angle between the observer direction and the mirror-reflection direction.

Derivation. *Lambertian:* the projected-solid-angle integral over the hemisphere is $\int \cos \theta d\Omega = 2\pi \int_0^{\pi/2} \cos \theta \sin \theta d\theta = \pi$ — so $f_r = \rho/\pi$ integrates to exactly ρ . That is why the divisor is π , not 2π . *Phong:* with the lobe centered on the normal ($\alpha = \theta$), $\int_{\text{hemi}} \cos^n \alpha \cos \theta d\Omega = 2\pi \int_0^{\pi/2} \cos^{n+1} \theta \sin \theta d\theta = 2\pi/(n+2)$; the factor $(n+2)/(2\pi)$ therefore makes the specular lobe integrate to exactly ρ_s at normal incidence, so $\rho_d + \rho_s \leq 1$ suffices for energy conservation there. The weaker $(n+1)/(2\pi)$ variant seen in graphics literature normalizes $\int \cos^n \alpha d\Omega$ *without* the $\cos \theta$ throughput factor — it conserves lobe solid-angle weight, not reflected energy, and can over-reflect by up to $(n+2)/(n+1)$. RADIANT uses the energy-normalizing $(n+2)/(2\pi)$.

Assumptions & validity. Lambertian: exact energy conservation for any $\rho \leq 1$; view-independent radiance. Phong: phenomenological (no reciprocity in the classic form, no Fresnel dependence); the normalization is exact only when the mirror direction coincides with the normal — at grazing incidence part of the lobe falls below the horizon and the actual reflected fraction is $< \rho_s$ (energy lost, never gained: conservative but biased). Adequate for glint order-of-magnitude budgets; use microfacet models for rendering-grade physics. RADIANT evaluates $\alpha = |\theta_{\text{obs}} - \theta_{\text{sun}}|$ in the plane of incidence (azimuthally symmetric v1 geometry).

Pitfalls.

- ρ instead of ρ/π — off by π in every reflected radiance; arguably the single most common radiometry bug in existence. Or $\rho/(2\pi)$ from the hemisphere-vs-projected-solid-angle confusion.
- $(n+1)/(2\pi)$ where energy normalization $(n+2)/(2\pi)$ is intended, or vice versa without documenting the convention.
- Measuring α from the surface normal instead of from the mirror-reflection direction.
- $\cos^n \alpha$ not clamped at $\alpha > \pi/2$ — odd n on a negative cosine produces negative radiance. RADIANT clamps $\cos \alpha$ at zero.
- Adding a Lambertian *radiance* term to a Phong *BRDF* term (unit mismatch: one already has the $1/\pi$ folded into a radiance, the other is sr^{-1}).

Numeric anchor. Lambertian, $\rho = 0.3$: $\int f_r \cos \theta d\Omega = 0.300000000$ (dimensionless, $= \rho$ exactly). Phong with $(n+2)/(2\pi)$ at normal incidence: the hemispherical integral over ρ_s equals 1.000000000 (dimensionless) for $n = 1, 10, 100$ (audit adaptive quadrature).

References. [Nicodemus 1977], [Phong 1975], [Lewis 1994].

3.3.11 Point-source irradiance and pixel solid angle

Equation. A point source of spectral intensity $I(\lambda)$ [W/sr/μm] at range R [m] through path transmittance $\tau_{\text{atm}}(\lambda)$ delivers at-aperture spectral irradiance

$$E_{\text{ap}}(\lambda) = \frac{I(\lambda) \tau_{\text{atm}}(\lambda)}{R^2} \quad \left[\text{W/m}^2/\mu\text{m} \right]$$

The pixel-subtended solid angle (small-angle) is

$$\Omega_{\text{pix}} = \frac{p_x p_y}{f^2} = \text{IFOV}_x \cdot \text{IFOV}_y \quad [\text{sr}]$$

and the target solid angle is $\Omega_t = A_t/R^2$ [sr].

Symbols. $I(\lambda)$ [W/sr/μm] — for a small Lambertian-emitting facet $I = L \cdot A_t$; R [m] slant range; p_x, p_y [m] pixel pitches; f [m] focal length; A_t [m²] projected target area.

Derivation. Inverse square is conservation of energy through expanding spheres: intensity is fixed, irradiance $= I/R^2$. Aperture spectral power is then $\Phi(\lambda) = E_{\text{ap}} A_{\text{ap}}$ [W/μm]. RADIANT's spectral-integration stage implements the equivalent radiance route: $\Phi = L_{\text{target}} \cdot A_{\text{ap}} \cdot \Omega_t$ with $\Omega_t = A_t/R^2$ — identical to $I/R^2 \cdot A_{\text{ap}}$ because $I = L A_t$. The blur spreads this energy per the PSF; the ensquared-energy fraction EE_{box} multiplies the in-pixel signal, applied exactly once, downstream in spectral integration.

Assumptions & validity. Point-source regime valid when the source's angular extent $\ll$ IFOV and $\ll$ PSF width. $\Omega = A/R^2$ is the small-angle approximation — exact solid-angle formulas matter only for IFOV $\gtrsim 0.1$ rad, never for imaging sensors. Flat normal-incidence footprint assumed; off-nadir the ground footprint grows by $1/\cos(\text{incidence})$. τ_{atm} is the *slant-path* transmittance, not the vertical column.

Pitfalls.

- Applying both $1/R^2$ and Ω_{pix} to the same term — double-counting the geometry. Radiance-based extended-scene math ($L \cdot \Omega_{\text{pix}} \cdot A_{\text{ap}}$) and irradiance-based point-source math ($I/R^2 \cdot A_{\text{ap}}$) are *alternative routes*, never multiplied together.
- τ_{atm}^2 (two-way path) sneaking in from radar heritage — passive EO is one-way.
- Range in km fed to a formula expecting m — a 10^6 error in E_{ap} .
- Applying EE_{box} anywhere but once, in spectral integration.

Numeric anchor. $I = 100 \text{ W/sr/}\mu\text{m}$, $\tau = 0.7$, $R = 500 \text{ km}$: $E_{\text{ap}} = 2.80000000 \times 10^{-10} \text{ W/m}^2/\mu\text{m}$. Pixel solid angle for GSD = 3 m at $R = 500 \text{ km}$: $\Omega_{\text{pix}} = 3.60000000 \times 10^{-11} \text{ sr}$ (IFOV = 6 μrad).

References. [Holst 2008], [Wolfe & Zissis §1].

3.3.12 Sub-pixel fill fraction and radiance mixing

Equation. When the target's solid angle is smaller than the pixel's:

$$ff = \frac{\Omega_t}{\Omega_{\text{pix}}} = \frac{A_t}{R^2 \Omega_{\text{pix}}} \quad (0 < ff \leq 1, \text{ clamped at } 1)$$

The pixel-averaged apparent radiance is the area-weighted mix (with path radiance filling the whole pixel uniformly):

$$L_{\text{mixed}}(\lambda) = ff \cdot L_{\text{target}}(\lambda) \cdot EE_{\text{box}} + (1 - ff) L_{\text{bg}}(\lambda) + L_{\text{path}}(\lambda)$$

Symbols. ff [—] fill fraction; L_{target} , L_{bg} [$\text{W/m}^2/\text{sr}/\mu\text{m}$] pure target and background contributions at the aperture (path radiance separated out); L_{path} [$\text{W/m}^2/\text{sr}/\mu\text{m}$] path radiance; EE_{box} [—] applied to the compact target term only.

Derivation. The pixel sees Ω_{pix} ; the target fills Ω_t of it and background fills the rest — hence the $ff / (1 - ff)$ split. The **regime-consistency identity** ties the sub-pixel and point-source routes together: the target's contribution to aperture irradiance via the point-source route is $L_t A_t \tau / R^2$, and via the sub-pixel route it is $L_t \cdot ff \cdot \Omega_{\text{pix}} \cdot \tau$. Since $ff \cdot \Omega_{\text{pix}} = A_t / R^2 = \Omega_t$ by construction, the two routes agree identically — that identity is what the fill-fraction anchor test pins. Path radiance is generated along the line of sight, not at the target, so it is added once, unweighted by ff , and never multiplied by EE_{box} .

Assumptions & validity. Areal mixing assumes the target sits fully within one IFOV footprint (no straddling-pixel split modeling). Clamps to $ff = 1$ when the target overfills the pixel (the regime then belongs to extended-scene handling). Fill fraction from geometry requires positive area, range, pitches, and focal length; otherwise RADIANT falls back to the explicit `fill_fraction` parameter rather than guessing.

Pitfalls.

- Area vs side-length confusion: $ff = A_t/\text{GSD}^2$, not l_t/GSD .
- Forgetting the $(1 - ff)$ background complement — the unfilled pixel area still sees background radiance.
- Applying EE_{box} to the background term — background is extended; EE applies to the compact target only, and RADIANT guards this explicitly.
- Splitting path radiance by ff — it fills the pixel uniformly.
- Deriving ff from nadir GSD^2 off-nadir — the footprint grows with incidence.

Numeric anchor. $A_t = 1 \text{ m}^2$, $\text{GSD} = 3 \text{ m}$ (i.e. Ω_{pix} matched to a 3 m footprint): $ff = 0.11111111$ (dimensionless).

References. [Holst 2008], RADIANT Source/Target System doc §3.

3.3.13 Photon conversion — spectral power to photoelectron rate

Equation. Photon energy $E_{\text{ph}} = hc/\lambda_m$ [J]. The photoelectron rate from spectral power $\Phi(\lambda)$ [W/ μm] at the detector is

$$\dot{n}_e = \int \text{QE}(\lambda) \Phi(\lambda) \frac{\lambda_m}{hc} d\lambda \quad [\text{e}^-/\text{s}], \quad \lambda_m = \lambda \cdot 10^{-6}$$

integrated over λ in μm ; total signal $N_e = \dot{n}_e \cdot t_{\text{int}}$ [e-]. RADIANT packages the same physics as an aperture-referred spectral responsivity

$$R(\lambda) = A_{\text{ap}} \Omega_{\text{pix}} \tau_{\text{opt}}(\lambda) \text{QE}(\lambda) \frac{\lambda_m}{hc} \quad [\text{e}^-/\text{s per W/m}^2/\text{sr}/\mu\text{m}]$$

with $R_{\text{band}} = \int R(\lambda) d\lambda$ used for backward propagation ($L_{\text{aperture}} = N_e/(R_{\text{band}} t_{\text{int}})$).

Symbols. $\Phi(\lambda)$ [W/ μm]; $\text{QE}(\lambda)$ [e-/photon, dimensionless]; $\lambda_m/(hc)$ [photons/J] converts watts to photons/s; t_{int} [s]; A_{ap} [m^2]; Ω_{pix} [sr]; τ_{opt} [—].

Derivation. A watt of monochromatic light at wavelength λ carries $1/E_{\text{ph}} = \lambda_m/hc$ photons per second. Weight by QE per photon and integrate across the band. The 10^{-6} lives in λ_m : λ in μm must be converted to meters so $\lambda_m/(hc)$ has units 1/J.

Assumptions & validity. One photoelectron per detected photon, weighted by QE (no avalanche gain — gain is a separate downstream stage). λ/hc must stay *inside* the integral: photon energy varies across the band, and pulling out a band-center value biases wide bands (for 3–5 μm the edge-to-edge photon-energy ratio is 5/3). QE convention (per incident vs per absorbed photon) must be consistent, with fill factor and window transmission counted once — inside QE or as separate factors, never both. Shot noise follows as $\sqrt{N_e}$ only if N_e is a true photon count: any premature gain scaling corrupts the Poisson statistics.

Pitfalls.

- The 10^{-6} : using λ in μm directly in $\lambda/(hc)$ — rate high by 10^6 .
- Energy-weighted vs photon-weighted band averages confused: QE tables are per-photon; responsivity in A/W is per-energy ($R_{A/W} = \text{QE} \cdot q\lambda_m/hc$) — mixing them double-counts λ .

- Applying t_{int} inside a “rate” function so downstream multiplies by it again.
- The electron charge $q = 1.602 \times 10^{-19}$ C appearing where it does not belong — electron *counts* are dimensionless; charge enters only when converting to amperes.

Numeric anchor. Monochromatic $\Phi = 10^{-12}$ W at $\lambda = 4 \mu\text{m}$ with $\text{QE} = 0.7$: $\dot{n}_e = 1.40955264 \times 10^7$ e-/s (check: $E_{\text{ph}}(4 \mu\text{m}) = 4.96611464 \times 10^{-20}$ J, and 0.7×10^{-12} W / $4.96611464 \times 10^{-20}$ J reproduces the rate).

References. [Holst 2008], [Janesick 2001], [CODATA 2018].

3.4 How the foundations feed the chain

The ten stages consume the foundations above in a fixed order. Each stage is a pure function of the chain state and the parameters; all inter-stage data flows through one immutable state object, which no stage mutates.

#	Stage	Radiometric role	Foundations used	Detailed in
0	Geometry	Slant range R , incidence, solar geometry	— (feeds R , θ_{sun} to everything)	Ch. 2
1	Source	Build $L_{\text{target}}(\lambda)$, $L_{\text{bg}}(\lambda)$; tentative regime	Planck, graybody, reflected solar, BRDF, brightness/radiance temperature converters, point-source intensity	this chapter
2	Atmosphere	$L_{\text{ap}} = L \tau_{\text{atm}} + L_{\text{path}}$	band integration on the $\tau(\lambda)$ grid	Ch. 4
3	Optics	Throughput τ_{opt} , A_{ap} , Ω_{pix} ; warm-optics self-emission; final regime	Kirchhoff emissivity, Planck (self-emission), pixel solid angle	Ch. 5
4	Platform	Smear/jitter degradation; EE_{box} from the fully degraded PSF	—	Ch. 5
5	Spectral Integration	Spectral $\rightarrow$ scalar, exactly once; EE_{box} applied exactly once	photon conversion, point-source/sub-pixel regime radiometry, fill-fraction mixing	this chapter
6	Detector	Noise budget (16 terms)	$\sqrt{N_e}$ shot statistics	Ch. 6
7	Readout	TDI, binning, coadd, gain, ADC scaling	—	Ch. 6
8	Calibration	Post-NUC residuals, drift, bias budget — calibration noise terms and accuracy inputs; the stage adds terms and never a frame	Planck at the calibration points, band responsivity	Ch. 7
9	Performance	SNR, NEDT, NIIRS, system MTF	$\partial B / \partial T$ (NEDT kernel), band responsivity (backward propagation)	Ch. 8

Three chain-glue equations worth stating here because they are pure radiometry:

At-aperture radiance (Atmosphere stage):

$$L_{\text{ap}}(\lambda) = L_{\text{target}}(\lambda) \tau_{\text{atm}}(\lambda) + L_{\text{path}}(\lambda)$$

where τ_{atm} follows Beer–Lambert along the slant path and L_{path} is atmospheric self-emission and scatter into the line of sight.

Extended-scene signal electrons (Spectral Integration stage — the only spectral-to-scalar collapse in the chain):

$$N_e = \int_{\lambda_1}^{\lambda_2} L_{\text{ap}}(\lambda) \tau_{\text{opt}}(\lambda) A_{\text{ap}} \Omega_{\text{pix}} \text{QE}(\lambda) \frac{\lambda_m}{hc} t_{\text{int}} d\lambda$$

with EE_{box} multiplying the target term only in point-source and sub-pixel regimes, never in extended scenes and never on the background term.

SNR (Performance stage): $\text{SNR} = N_e / \sigma_{\text{total}}$ with σ_{total} the RSS of the noise budget; contrast SNR uses $N_{e,\text{target}} - N_{e,\text{bg}}$ in the numerator. NEDT divides the noise-equivalent radiance by the band-integrated $\partial B / \partial T$ — which is why the NEDT-kernel entry above carries per-K units end to end.

3.4.1 Dimensional Trace

Stage	Input units	Output units	Conversion
Geometry	orbit/attitude params	m, rad	trigonometry
Source	K, dimensionless ϵ/ρ	W/m ² /sr/μm	Planck + ϵ ; BRDF · E _{sun} · cos θ
Atmosphere	W/m ² /sr/μm	W/m ² /sr/μm	× $\tau_{\text{atm}} + L_{\text{path}}$
Optics	W/m ² /sr/μm	W/m ² /sr/μm (+ A _{ap} [m ²], Ω _{pix} [sr] published)	× τ_{opt}
Platform	— (radiometric pass-through)	EE _{box} [—]	PSF ensquared energy
Spectral Integration	W/m ² /sr/μm, m ² , sr, s	e-	× A · Ω · QE · λ _m /hc · t _{int} , ∫ dλ
Detector	e-	e- RMS (noise terms)	noise model
Readout	e-	e- (scaled), DN	× N _{TDI} · M _{bin} · N _{coadd} ; ÷ gain
Performance	e-	dimensionless (SNR), K (NEDT), NIIRS	S/σ; σ _L / ∫ (∂B/∂T) dλ

Every row must check; integrating a per-μm density over μm is the only “conversion-free” collapse, and the two 10^{−6} factors in the chain (Planck Jacobian, photon-energy λ_m) each appear exactly once.

3.5 Assumptions and Limitations

- Lambertian target emission; BRDF structure available only through the Lambertian and normalized-Phong models (no microfacet, no measured BRDF tables)
- Sun as a collimated directional source; blackbody-5778 solar spectral shape (few-percent deviations from the real spectrum; tabulated ASTM E490 model planned)
- No molecular spectroscopy or fluorescence
- No atmospheric refraction or scintillation
- No ghost images, BSDF scatter, or chromatic aberration
- No optical crosstalk between pixels
- No temporal variability in scene
- The project’s scope-decision register records the full list of deferred effects

Atmosphere Models

The physics of RADIANT’s two atmosphere model families — the closed-form **simple parametric** model and the **library-backed** interpolated / tabulated / MODTRAN-import models — organized by model rather than by change history. What each model computes, from what first principles, and which measurement or test holds each claim in place.

This document is the *physics*. The measured accuracy of every model against the MODTRAN run set is recorded in the MODTRAN-parity validation record; the architecture contract (AtmosphericQuantities, backend dispatch, guard structure) is specified in the atmosphere architecture document; which model an operator should reach for is the User’s Guide’s model-selection guidance. Atmospheric **turbulence** is a spatial effect and is documented with the rest of the MTF cascade in the spatial chapter’s *Atmospheric turbulence*; it is not repeated here.

Symbols. λ wavelength [μm]; ζ zenith angle at a path segment’s **lower** endpoint [rad]; θ_s solar zenith [rad]; Θ scattering angle [rad]; h altitude above mean sea level [m]; H_i species scale height [m]; τ transmittance [–]; OD optical depth [–]; m air mass [–]; σ volume extinction coefficient [km^{-1}]; w precipitable water [cm]; ω_0 single-scattering albedo [–]; $B(\lambda, T)$ Planck spectral radiance [$\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$]; L radiance [$\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$]; E irradiance [$\text{W}/\text{m}^2/\mu\text{m}$]; $R_E = 6.371 \times 10^6$ m.

4.1 Two families, one contract

Every model in RADIANT delivers the same product bundle — a transmittance, an upwelling path radiance, a downwelling emission term, and (for the composed topologies) the segment products a path is assembled from. What differs is where the numbers come from:

Family	Members	Physics origin	Geometry response
Closed-form parametric	simple	Beer-Lambert on four species, calibrated band-by-band against MODTRAN	Recomputed for every geometry
Library-backed	interpolated, tabulated, modtran (tape7 import)	Real MODTRAN 6 radiative transfer, stored and served	interpolated interpolates over declared axes; the other two are geometry-agnostic files
Identity	exo	$\tau \equiv 1, L \equiv 0$ — the cosmic vacuum	None (exact)

The simple model is the only backend that can serve an arbitrary path topology (down-looking, up-looking, level, grazing, twilight), because the segment evaluators are built on its species model. The library families serve the topologies their runs actually measured, and refuse the rest rather than approximating them.

4.2 The simple parametric model

4.2.1 Beer-Lambert structure

Transmittance along a path is the exponential of a slant optical depth built from four species, each carried on its own exponential density profile:

$$\tau_{atm}(\lambda) = \exp[-OD_{tot}(\lambda)], \quad OD_{tot}(\lambda) = \sum_i m_i OD_{i,vert}(\lambda)$$

with $i \in \{\text{mol}, \text{aer}, \text{h}_2\text{O}, \text{gas}\}$ and m_i the species air mass of *Air mass* below. The **vertical** column each species presents between two altitudes is the analytic exponential integral

$$\text{col}_i(h_{lo}, h_{hi}) = \int_{h_{lo}}^{h_{hi}} e^{-h/H_i} dh = H_i \left(e^{-h_{lo}/H_i} - e^{-h_{hi}/H_i} \right) \quad [\text{km}]$$

which is endpoint-symmetric, so one segment read in either direction presents one column — the structural reason transmittance is single-valued (*Path topologies*).

Scale heights: $H_{\text{mol}} = 8000$ m, $H_{\text{aer}} = 1200$ m, $H_{\text{h}_2\text{O}} = 2000$ m. The well-mixed-gas floor rides the molecular profile by construction (*The well-mixed-gas absorption floor*).

4.2.2 Molecular (Rayleigh) scattering

The published Rayleigh constant is a **total vertical optical depth**, dimensionless — not a per-km coefficient:

$$\tau_{R,vert}(\lambda) = 0.0088 \lambda_{\mu\text{m}}^{-4.09}$$

At $\lambda = 0.55 \mu\text{m}$ this evaluates to 0.101484, against the published whole-atmosphere value 0.0973–0.10 ([Hansen & Travis 1974], [Bucholtz 1995]). The sea-level volume extinction the slant integral needs is **derived** from it through the exponential profile's own identity $\tau_{vert} = \sigma_0 H$:

$$\sigma_{\text{mol}}(\lambda) = \frac{\tau_{R,vert}(\lambda)}{H_{\text{mol}}} = 0.012686 \text{ km}^{-1} \text{ at } 0.55 \mu\text{m} \quad (\text{published } \approx 0.0116 \text{ km}^{-1})$$

and is scaled along the path by $e^{-h/H_{\text{mol}}}$.

Deriving σ_0 rather than storing it independently is what prevents the classic error of consuming the published optical depth *as* the km^{-1} coefficient, which inflates every VIS/NIR molecular optical depth by exactly the column depth, $\approx 8\times$.

4.2.3 Aerosol (Mie) extinction

Aerosol extinction at the reference wavelength is fit from the Koschmieder visibility relation, then carried spectrally by an Ångström power law:

$$\sigma_{\text{aer}}(0.55 \mu\text{m}) = \frac{3.912}{V_{km}} [\text{km}^{-1}], \quad \sigma_{\text{aer}}(\lambda) = \sigma_{\text{aer}}(0.55) \left(\frac{\lambda}{0.55} \right)^{-\alpha}$$

The constant $3.912 = -\ln(0.02)$ is the 2 % contrast threshold defining meteorological visibility. Three canonical aerosol regimes:

Regime	Ångström α	Aerosol SSA ω_{aer}
rural	1.3	0.95
urban	1.5	0.85
maritime	0.7	0.99

The power law is a **scattering** model: good in VIS/SWIR, weak but usable through the MWIR, and wrong in the LWIR where real IR aerosol extinction is absorption-dominated and roughly flat. It is therefore clamped at the MWIR–LWIR boundary $\lambda_{\text{clamp}} = 5.0 \mu\text{m}$ — beyond it the extinction is frozen at its $5 \mu\text{m}$ value rather than decaying unphysically toward zero, and the model warns once per run when the clamp engages.

4.2.4 Water vapor — the curve of growth

Water vapor is **not** linear in Beer’s law over a real band, because the strong lines saturate long before the windows between them do. The calibrated model is a 17-region curve of growth:

$$\text{OD}_{\text{h}_2\text{o}}(\lambda) = k(\lambda) w_{\text{eff}}^{b(\lambda)}$$

where w_{eff} [cm] is the path water amount — total precipitable water times the traversed fraction of the $H_{\text{h}_2\text{o}} = 2 \text{ km}$ exponential column. The exponent is fit region by region against the real MODTRAN 6 water ladder (runs D4 / A1 / D5, $\text{H}_2\text{O} \times 0.5 / \times 1 / \times 2$): **sub-linear** $b \approx 0.2\text{--}0.8$ in the saturated absorption bands, **super-linear** $b \approx 1.3\text{--}1.75$ in the LWIR where the e-type continuum dominates and the absorber’s own pressure broadening compounds with amount.

This replaced a five-Lorentzian line fit whose far wings made the MWIR water response $\approx 5\times$ too steep.

Standard-profile water columns ([McClatchey 1972], carried into MODTRAN MODELS 1–6) supply the default when the operator selects a climate profile and leaves the water column at its schema default:

Profile	PWV [cm]	Profile	PWV [cm]
tropical	4.11	midlat_winter	0.85
midlat_summer	2.92	subarctic_summer	2.08
us_standard	1.40	subarctic_winter	0.42

4.2.5 The well-mixed-gas absorption floor

Each of the same 17 regions carries a **water-independent** vertical optical depth — the CO_2 $4.3 \mu\text{m}$ and $15 \mu\text{m}$ bands, N_2O , the O_3 $9.6 \mu\text{m}$ band, and O_2 / CH_4 overtones — in excess of what Rayleigh and aerosol already supply:

$$\text{OD}_{\text{region}}(w) = \text{floor_od} + k w_{\text{eff}}^b$$

The floor rides the molecular scale height — it is defined as a fraction of the molecular column — and it enters the single-scattering albedo denominator as a **pure absorber** (*Single-scatter solar path radiance*). Without it a model attributes the MWIR CO₂ floor to water and evaluates $\omega_0 \approx 1$ for space columns.

The calibrated table, exactly as shipped (`_CALIBRATED_GAS_REGIONS`):

Region [μm]	floor_od	k	b	Region [μm]	floor_od	k	b
0.30–0.45	0.1262	0.0000	1.000	3.10–3.50	0.1370	0.5824	0.457
0.45–0.70	0.1375	0.0025	0.874	3.50–5.00	0.4494	0.0944	0.808
0.70–1.30	0.0402	0.1245	0.434	5.00–7.50	1.3543	1.7850	0.530
1.30–1.50	0.0000	1.0933	0.327	7.50–8.00	0.9424	0.9210	0.673
1.50–1.75	0.0217	0.0282	0.645	8.00–9.40	0.1494	0.0992	1.204
1.75–2.05	0.0000	1.1186	0.216	9.40–9.90	0.8877	0.0409	1.701
2.05–2.40	0.0747	0.0320	0.843	9.90–10.00	0.3013	0.0379	1.805
2.40–3.10	0.7440	0.9666	0.560	10.00–12.00	0.0471	0.0602	1.750
				12.00–14.29	0.5956	0.1398	1.583

Spectral shape *within* a region is flat: the model’s contract is band-integrated fidelity, not line structure. Wavelengths outside 0.30–14.29 μm clamp to the edge regions’ calibration.

The eight bold floor_od values are a re-fit of the visible, near-infrared and short-wave rows. floor_add is the measured opacity *in excess of* Rayleigh + aerosol and is clamped at zero, so an over-large Rayleigh reference pins the visible rows to that clamp and leaves the model too transmissive. Only floor_od moves in such a re-fit: k and b are solved in closed form from the MODTRAN ladder alone and never see the RADIANT model, so their bit-identity across a re-fit is the check that the calibration *reference* changed and the fit did not.

The fit is evaluated on the ladder’s own grid. floor_add is a *difference* of two band optical depths, and the two have to be weighted alike. MODTRAN’s native grid is uniform in wavenumber (1 cm^{-1} , so $\Delta\lambda \propto \lambda^2$); a uniform- λ reference disagrees with it wherever the spectrum is steep, always toward an over-large floor — a few hundredths of an optical depth in the visible, ≤ 0.0004 beyond 1.3 μm . The generator evaluates the reference on the ladder’s grid, so the two weightings match.

The same convention removes a *coverage* mismatch in the first row. The delivered tape7 grid starts at 0.374953 μm , so the measured 0.30–0.45 μm optical depth is really the 0.375–0.45 μm mean, while an unrestricted reference spans the whole region — including 0.30–0.375 μm , where Rayleigh alone is enormous. Measured over the same interval the row reads 0.642 against the ladder’s 0.768 and carries **0.1262**, within 0.014 OD of the 0.45–0.70 μm floor: a continuous short-wavelength deficit rather than an artificial 0.16 OD step at the 0.45 μm edge, which is physically what one expects, since whatever supplies that deficit does not know where a table boundary is. **Its residual limitation: the row is fitted from 0.375–0.45 μm and applied across 0.30–0.45 μm ,** because no anchor data exists below 0.374953 μm .

At the fit’s own anchor (us_standard, rural 23 km, nadir full column) the model’s 0.45–0.70 μm band optical depth reads 0.4566 against MODTRAN’s 0.4561 — **0.1 %**. RMS band-mean τ parity across thirteen full-column anchors is 0.0111 on 0.45–0.70 μm and 0.0292 on 0.40–0.90 μm .

One residual is recorded rather than tuned away: 0.14 optical depths is far more than real 0.45–0.70 μm gas chemistry supplies — the O_3 Chappuis band contributes ~ 0.03 — so part of the visible floor is standing in for an aerosol-model deficit. The band total is right; the attribution between gas and aerosol is not resolved by this fit, and the 0.30–0.45 μm row landing beside it is further evidence of that.

The 8.00–9.40 / 9.40–9.90 / 9.90–10.00 μm partition is an ozone split. One flat row across 8.00–10.00 μm is a 2 μm slab covering both the clean window and the O_3 ν_2 fundamental, so the model carries no identifiable ozone opacity at all: on a nadir full column $\tau(9.60)$ and $\tau(8.70)$ agree to six figures. Cutting the partition at the band edges the delivered ladder shows — water-free OD rises $0.24 \rightarrow 0.89$ across 9.372–9.416 μm and falls $0.52 \rightarrow 0.33$ across 9.901–9.911 μm — gives a band-core floor **5.9× the adjacent window’s**, which is the ozone. Two consequences follow directly from the table:

- The **ozone share of the in-feature floor is arithmetic**, not fitted: $(0.8877 - 0.1494)/0.8877 = 0.832$, the band-core floor above the continuum floor the neighbouring window measures.
- The floor still rides the **molecular** scale height, so the identified ozone is placed at ~ 8 km rather than ~ 25 km. Full columns barely notice; partial columns do. Placing it correctly is separate work that this split unblocks and does not itself perform.

The whole table is one vintage, regenerated from a single run of the generator, which reproduces its own output bit-for-bit on a re-run.

4.2.6 The region-edge smoothstep blend

The table is piecewise-constant, but it is **not** read as a step function. Across each of the sixteen interior region edges the three coefficients $c \in \{\text{floor_od}, k, b\}$ are joined by a C^1 smoothstep ramp of half-width $h_w = 0.02$ μm (full width 0.04 μm):

$$u(\lambda) = \text{clip}\left(\frac{1}{2} + \frac{\lambda - \lambda_{\text{edge}}}{2h_w}, 0, 1\right), \quad S(u) = u^2(3 - 2u), \quad c(\lambda) = c_{lo} + (c_{hi} - c_{lo})S(u)$$

with $S(0) = 0$, $S(1) = 1$, $S'(0) = S'(1) = 0$ — the ramp meets the flat calibrated regions with matching value *and* slope — and $S(\frac{1}{2}) = \frac{1}{2}$, so the edge itself carries the exact arithmetic mean of the two regions. At the 0.70 μm edge, for instance, k evaluates to $(0.0025 + 0.1245)/2 = 0.0635$ exactly, and floor_od to $(0.1375 + 0.0402)/2 = 0.08885$ — real work on both sides of the edge, where an uncalibrated pair of zeros would have contributed nothing.

Outside the ramps nothing changes: a λ at or beyond h_w from every edge keeps the bit-identical calibrated coefficient. Every region is wider than $2h_w$ — the narrowest is 0.10 μm at 9.90–10.00 μm , two and a half times the full ramp width — so no two ramps overlap, which is the invariant that stops a future refit from silently invalidating the blend.

Read literally, the step table made $\tau(\lambda)$ jump at every edge and made a band-mean τ that straddled an edge **sampling-grid-dependent**. Both effects, and the size of the adopted blend, are in the parity document §2.7.

h_w is a documented module constant, not a schema parameter (the KOSCHMIEDER precedent): it is a numerical property of the calibration table, not a tuneable physical quantity.

4.2.7 Air mass — plane-parallel inside 80°, per-species spherical past it

The slant path length over the *absorbing* vertical extent of the segment is

$$L_{\text{slant}} = \frac{\Delta h_{\text{absorbing}}}{\cos \zeta}, \quad \Delta h_{\text{absorbing}} = \min(|h_{\text{sen}} - h_{\text{tgt}}|, h_{\text{atm,top}} - \min(h_{\text{sen}}, h_{\text{tgt}}))$$

so a ground site viewing a 700 km target traverses 100 km of air, not 700 km. The stored air mass is therefore exactly $m = \sec \zeta$, and a wholly exo path returns $m = 1$ by the ExoAtmosphere convention.

$\sec \zeta$ is the honest plane-parallel primitive and it is used unchanged for $\zeta \leq \text{SPHERICAL_SWITCH_RAD} = 80^\circ$. Past 80° every column — the observer column, the full ground-to-sensor column, the solar column, and any ColumnSegmentSpec — hands over to the exact spherical slant integral, **per species**:

$$m_i = \frac{S_i(r_0; h_{lo} \rightarrow h_{hi}; H_i)}{\text{col}_i}, \quad r_0 = (R_E + h_{lo}) \sin \zeta_{lo}$$

where S_i is the density-weighted spherical column about the ray's perigee radius r_0 (grazing_column.grazing_slant_column_km, a graded-grid quadrature anchored analytically against Chapman's grazing limit). The well-mixed-gas floor rides m_{mol} , because it is defined as a fraction of the molecular column.

Three properties matter, and all three are measured rather than asserted:

- **It has to be per species.** Water vapor's 2 km profile hugs the tangent point far harder than the 8 km molecular one, so a single corrected scalar cannot serve both — at $\zeta = 89.4^\circ$ the two diverge by $2.27 \times$ in error.
- **The direction is toward more signal.** The spherical air mass is always the smaller one, so transmittance and SNR move **up** past 80° and never down.
- **It is a step, not a blend.** The size of that step, and why 80° rather than the 89.5° ceiling, are in the parity document §2.4.

The solar column's old 89.5° clamp retires with the hand-over: the spherical route has no ceiling, so a twilight scene at $\theta_s = 89.9^\circ$ now gets its own column. ZENITH_CEILING_RAD (89.5°) still bounds the *observer* zenith, because that is the column-air-mass validity ceiling for the plane-parallel form.

4.2.8 One linearization convention — the slant column

Two calibrated terms — the water curve of growth and the gas floor — are *column* optical depths, the integral of no local coefficient. Wherever the model needs a **local** extinction (the single-scatter weights of *Single-scatter solar path radiance*, the level arm of *Path topologies*), they must be linearized: divided by a reference column to produce an equivalent per-km coefficient.

Because the curve of growth is sub-linear, that choice is not neutral. Linearizing against the vertical column instead of the slant one scales the effective water weight by m_{h2o}^{b-1} , and ω_0 with it, wherever water absorbs.

All three evaluators now linearize against the slant column — the amount actually traversed, which is what the linearization is *of*. column_segment_optical_depth publishes slant_column_mol_km / _aer_km / _h2o_km provenance under the same key names the near-horizon branch already used, so the convention is inspectable.

Note the scope precisely: this is the convention for the **linearized local weights**. The optical depth itself is still built as *vertical column* $\times$ *species air mass* (*Air mass*); the two are consistent because m_i is defined as the ratio of the slant to the vertical column.

At $\zeta = 0$ the air mass is exactly 1, so slant $\equiv$ vertical and every vertical anchor is bit-identical across this change.

4.2.9 Single-scatter solar path radiance

The scattered term on any column is the classic single-scatter source weighted by the column's own emissivity-equivalent $1 - \tau$:

$$L_{\text{path,scat}}(\lambda) = \frac{E_{\text{sun}}(\lambda)}{4\pi} \cos \theta_s \omega_0(\lambda) P(\Theta, \lambda) [1 - \tau_{\text{atm}}(\lambda)]$$

with E_{sun} the top-of-atmosphere solar spectral irradiance and 4π the full-sphere phase-function normalization.

Single-scattering albedo. Extinction-weighted, with the pure absorbers in the denominator only:

$$\omega_0(\lambda) = \frac{\sigma_{\text{mol}} + \omega_{\text{aer}} \sigma_{\text{aer}}}{\sigma_{\text{mol}} + \sigma_{\text{aer}} + \sigma_{\text{h2o}} + \sigma_{\text{gas}}}$$

Phase function. Combined by *scattering* cross-section — a photon absorbed by water does not contribute scattered radiance at all — and normalized so an isotropic scatterer gives $P \equiv 1$:

$$P(\Theta, \lambda) = \frac{\sigma_{\text{mol}} P_R(\Theta) + \omega_{\text{aer}} \sigma_{\text{aer}} P_{HG}(\Theta)}{\sigma_{\text{mol}} + \omega_{\text{aer}} \sigma_{\text{aer}}},$$

$$P_R(\Theta) = \frac{3}{4} (1 + \cos^2 \Theta), \quad P_{HG}(\Theta) = \frac{1 - g^2}{(1 + g^2 - 2g \cos \Theta)^{3/2}}, \quad g = 0.7$$

Where the species proportions are evaluated — the lower endpoint. ω_0 and $P(\Theta)$ depend on the *relative* proportions of the four species, and those change with altitude. All evaluators take them at the segment's **lower endpoint**: the densest air in the path, the end the `L_toward_lower` product emerges from, and the choice the grazing and level evaluators always made.

The retired alternative — the segment's arithmetic-mean altitude — put the weights, for any column taller than ≈ 40 km, above the altitude where the aerosol and water coefficients underflow to zero in double precision. The consequence was not a small bias but a silent model substitution: ω_0 evaluated to exactly 1 (no absorption at all) and the Henyey-Greenstein forward peak collapsed onto the isotropic-Rayleigh 1.5, so a tall column scattered as if the atmosphere held no aerosol whatever `visibility_km` said. The adoption measurement is in the parity document §2.3.

This term is **provisional below 3 μm** . Single scattering under-predicts the daytime VIS/NIR sky, where multiple scattering dominates; a `UserWarning` fires when the evaluation grid extends below `SCATTERED_SKY_PROVISIONAL_MAX_UM = 3.0  $\mu\text{m}$` and a solar geometry with the sun above the local horizon is supplied. A pure-thermal MWIR/LWIR call warns about nothing, and neither does a night scene on a VIS grid.

4.2.10 Thermal path radiance — Kirchhoff emission at a height-resolved temperature

The thermal term on any column is a one-slab Kirchhoff graybody whose emissivity is derived from the column's own transmittance by Kirchhoff's law, never taken as an independent input:

$$L_{\text{path,therm}}(\lambda) = [1 - \tau_{\text{atm}}(\lambda)] B(\lambda, T_{\text{eff}}(\lambda))$$

It applies whether or not the sun is up — a night down-looking scene has scattered $\equiv 0$ and thermal > 0 — and one implementation serves both directions of travel.

What T_{eff} is. It is the single temperature that makes this one-slab form reproduce the **layered formal solution** of the segment's own non-isothermal air. Slice the segment into N sub-layers; layer i has slant optical depth $\delta_i(\lambda)$ and temperature T_i . Ordering the layers **from the end the radiation escapes**, the emergent radiance is the discrete formal solution of the non-scattering LTE transfer equation:

$$L(\lambda) = \sum_i B(\lambda, T_i) (1 - e^{-\delta_i(\lambda)}) e^{-c_i(\lambda)}, \quad c_i(\lambda) = \sum_{j < i} \delta_j(\lambda)$$

The weights telescope exactly:

$$\sum_i (1 - e^{-\delta_i}) e^{-c_i} = 1 - e^{-\sum_i \delta_i} = 1 - \tau(\lambda)$$

so $L = (1 - \tau) \langle B \rangle$ with $\langle B \rangle$ a convex combination of the layer Planck functions, and

$$T_{\text{eff}}(\lambda) = B^{-1}(\langle B \rangle(\lambda))$$

is guaranteed to lie between the coldest and warmest layer in the segment.

Three properties are structural, not fitted:

- **The total optical depth is untouched.** Only the altitude the emission is weighted at changes, so every τ in the model is bit-identical across this construction.
- **Isothermal is exact.** Every T_i equal returns that temperature exactly, for every τ , every direction and every layer count. A level arm is isothermal, so the level topologies keep a single exact graybody temperature.
- **Direction is geometry, not a fork.** `escape` names which endpoint the radiance leaves from. One model serves both directions: the down-looking `evaluate` term uses `escape="upper"` and the up-looking segment product `escape="lower"`. A direction-blind optical-depth-weighted mean temperature was measured against the same anchors and rejected — it degrades the MWIR everywhere, which is the measurement that forces `escape` into the model.

Where the opacity sits in altitude. The curve of growth fixes each species' *total* column optical depth and says nothing about its vertical distribution, which is what the weighting needs. That distribution is taken from first principles, with no fitted coefficient:

- **Scattering species** (Rayleigh, aerosol) — extinction is proportional to number density, so the weighting profile is the species' own density scale height, unchanged.
- **Pressure-broadened absorbers** (the well-mixed-gas floor, water vapor) — a Lorentz line's absorption coefficient goes as number density *times* the collisional half-width, and the half-width goes as total pressure, so $\alpha \propto \rho_a p_{\text{air}}$ and the emission weighting rides the harmonic combination

$$H_{\text{emit}} = \left(\frac{1}{H_a} + \frac{1}{H_{\text{air}}} \right)^{-1}$$

which takes the well-mixed floor from 8 km to **4 km** and water from 2 km to **1.6 km**.

- **Layered absorbers** (ozone) — a species produced photochemically in the mid stratosphere has no scale height at all, so it is placed on a Gaussian layer instead:

$$\rho_{\text{O}_3}(z) \propto \exp\left(-\frac{1}{2} \left(\frac{z - z_0}{w}\right)^2\right), \quad z_0 = 25 \text{ km}, w = 5 \text{ km},$$

the US Standard Atmosphere ozone profile's peak and half-width.

The ozone split. The calibrated `floor_od` is one number per spectral region covering CO₂, N₂O, CH₄ and O₃, and only the first three are well mixed. Inside the 9.40–9.90 μm O₃ ν_2 region the floor is therefore partitioned before it is placed:

$$\text{share}_{\text{O}_3}(\lambda) = 1 - \frac{\text{floor}_{\text{continuum}}(\lambda)}{\text{floor}(\lambda)} = \frac{0.8877 - 0.1494}{0.8877} = 0.8317 \quad \text{inside the band},$$

where `floorcontinuum` is the same region table evaluated with the band's row carrying its clean-window neighbor's floor. That share of the gas floor rides the ozone layer; the remainder keeps the 4 km pressure-broadened placement. Three properties follow:

- **the share is arithmetic, not fitted** — it is read off the two committed table rows (the well-mixed-gas floor's 8.00–9.40 and 9.40–9.90 μm regions), so a re-fit moves it automatically and there is no coefficient in the emission model to go stale. An absorption *band* stands on top of the continuum its neighbors carry, so the excess — not the total — is ozone;
- **placement is continuous in λ** — both floors pass through the same region-edge smoothstep, so the share inherits its C^1 ramp at 9.40 and 9.90 μm rather than needing a second ramp implementation that could drift from the first;
- **τ is untouched** — the two parts sum to the floor the model already had. This is a redistribution in altitude and nothing else, as every placement rule here is.

Outside the band the share is exactly zero, no layer is constructed, and T_{eff} is bit-identical to the four-species form. The 9.90–10.00 μm long-wave tail (floor 0.3013, 3.3 $\times$ its continuum) is a **documented exception**: part of it is ozone too, but the split is scoped to the band core the parity is measured in, so a narrow-band product centered on 9.95 μm still places its ozone at 4 km.

The sub-layer count is a convergence-tested quadrature parameter, not a tuning knob: `EMISSION_LAYERS_PER_SPECIES` = 32, whose discretization error against a 512-per-species reference is $\max |\Delta T_{\text{eff}}| = 0.016$

K over the whole anchor set — two orders of magnitude below the model’s own ≈ 4 K accuracy against MODTRAN. It is therefore fixed, not exposed as a user parameter.

Vacuum limit. A segment with no opacity has no emission to weight; the function falls back to the temperature of its densest air so the value stays finite. The radiance it multiplies is exactly zero there, so it is unobservable — it exists only so nothing returns NaN.

4.2.11 The hemispheric downwelling terms

Two products describe what falls on a *surface* — hemispheric irradiances, not directional radiances, and consumed by the reflected-diffuse terms of the target and ground-background arms.

Thermal. A target-anchored graybody on the **vertical** target $\rightarrow h_{\text{atm,top}}$ column — the sky the target actually sees, so the sensor’s altitude and viewing zenith deliberately do not enter a hemispheric flux at the target:

$$E_{\text{sky,thermal}}(\lambda) = \left[1 - \tau_{\text{sky,vert}}(\lambda)^D\right] \pi B(\lambda, T(h_{\text{tgt}} + z_{\text{em}})), \quad L_{\text{atm,down}}(\lambda) = \frac{E_{\text{sky,thermal}}(\lambda)}{\pi}$$

$T(\cdot)$ is the fixed-lapse ICAO standard-atmosphere lookup (6.5 K/km, floored at the 216.65 K tropopause above 11 km). The emission-height offset $z_{\text{em}} = 200$ m is fitted to the real up-looking MODTRAN H-runs (downwelling is dominated by near-surface air).

The flux-diffusivity exponent D is not fitted. It is a geometric identity,

$$D = \sec 48.2^\circ = 1.50030 \dots,$$

the secant of the diffusivity angle the entire downwelling reference set was run at: every reference deck is an up-looking MODTRAN run at 48.2° , and the quantity this model is scored against is $\pi L(48.2^\circ)$ — the hemispheric-flux proxy, itself validated to $\sim 15\%$ against a true hemispheric-flux table. Scoring against a 48.2° pencil and then weighting the emissivity by any other factor mixes two different hemispheric approximations; taking the exponent from the reference geometry removes a free parameter instead of re-tuning one. (The textbook Elsasser diffusivity factor 1.66 is $\sec 53.13^\circ$ — a different angle and a different approximation, not what this reference set uses.)

Because z_{em} is fitted *through this one closed form*, a directional path-radiance product cannot inherit it. That is why the height-resolved thermal path-radiance model and this graybody coexist deliberately: they are different products, not two versions of one.

The obvious refinement was measured and declined. Computing the downwelling directly — evaluating the sky column’s emergent radiance at 48.2° from the height-resolved machinery, and retiring both fitted constants — was scored against a nine-rung ladder of measured hemispheric-proxy downwelling from 0 to 80 km, as $\text{RMS} |\ln(\text{model}/\pi L_{\text{MODTRAN}})|$ over nine rungs $\times$ two bands. The shipped pairing of $\sec 48.2^\circ$ with the near-surface emission height scores **1.9233**; swapping in the layered temperature scores 1.9385, and restricted to the four tropospheric rungs — the only ones where the two temperatures differ, the standard profile being isothermal above the tropopause — it is worse still, 0.3087 against 0.4771.

The band split behind that is worth knowing, because it says what would have to change first. Isolating the temperature shows the layered form is a 3.8× improvement in the LWIR — the near-surface proxy runs 5–22 K too warm on a semi-transparent window that genuinely samples the whole column — and a 3.7× degradation in the MWIR, running 3.5–10.6 K too cold. With the region-flat 3–5 μm gas floor supplying too little opacity, the layered weighting reaches too high into cold air, and the fitted warm bias was compensating for a τ -shape error rather than a temperature error. The layered temperature therefore waits on a finer MWIR gas model, not on the emission-placement machinery.

Adopting sec 48.2° raises the downwelling effective emissivity everywhere — by $\sec 48.2^\circ / 1.1 = 1.364\times$ in the optically thin limit, and asymptotically not at all where the sky column already saturates — so every reflected-sky term rises.

Scattered. On the same vertical slab:

$$E_{\text{sky,scattered}}(\lambda) = E_{\text{TOA}}(\lambda) \cos \theta_s \omega_{0,\text{eff}}(\lambda, \text{aerosol}) [1 - \tau_{\text{down,vert}}(\lambda)]$$

$\omega_{0,\text{eff}}$ is **not** the internal column ω_0 of *Single-scatter solar path radiance*. It is a MODTRAN-derived effective single-scattering albedo: band-median values per aerosol regime over VIS 0.4–0.7 / NIR 0.7–1.4 / SWIR 1.4–2.5 μm, edge-extended outside, obtained by inverting *this closed form* against the real ground-level diffuse-flux tables — so it absorbs MODTRAN’s multiple-scatter contribution as an effective parameter of this formula. Examples: rural 0.791 / 0.698 / 0.187; urban 0.423 / 0.430 / 0.263.

The two ω_0 definitions coexist on purpose: the flux-fit table is not a valid substitute inside the phase-function-weighted L_{path} integral, and the internal column ω_0 over-predicted diffuse sky irradiance by $\approx 1.3\times$ (VIS rural) to $\approx 5\times$ (SWIR urban).

4.2.12 Path topologies — segments, arms, and whole paths

Every non-trivial path is a **composition of path segments**: one piece of path between two points, evaluated once, read from either end. Two spec types, because there are two topologies and they are not variations of one form:

Spec	Fields	Topology	Air mass
ColumnSegmentSpec	$h_{\text{low}}, h_{\text{high}}$ [m], ζ_{low} [rad]	endpoint-minimum — the path’s lowest point is an endpoint	<i>Air mass</i>
LevelArmSpec	altitude [m], length [m]	interior-tangent — the lowest point is in the middle	none at all

One evaluated product, `SegmentQuantities`, carries **one** τ and **two** directional radiances $L_{\text{toward upper}}$, $L_{\text{toward lower}}$. Transmittance is reciprocal, so a segment has one τ no matter which way it is read; path radiance is not, because the emitting and scattering layers are weighted by the transmittance of the material *between* them and the receiver, and that weighting reverses with direction.

The lower-endpoint convention. Every column segment’s zenith is keyed to its **lower** endpoint. Two structural reasons: it is the one endpoint the two travel directions share, so a single scalar describes the segment rather than the reading of it; and it is the angle MODTRAN’s Card 3 wants when $H_1 \leq H_2$, so no convention translation sits between RADIANT and its truth source.

The level arm. A level path cannot be served by the column machinery at all, and the reason is structural rather than accuracy: a column's optical depth is $\int e^{-h/H} dh$ between two equal altitudes, which is **exactly zero**, and its air mass is $\sec(\pi/2)$, which is undefined. The arm is instead

$$\tau(\lambda) = \exp[-\alpha(\lambda, h) L], \quad L_{\text{path}}(\lambda) = [1 - \tau(\lambda)] B(\lambda, T_{\text{eff}}) + \text{single-scatter source}$$

with α the **local** extinction at the arm's altitude [km^{-1}] and L the **true spherical chord** between the endpoints [km] — not a flat-Earth range. No new calibration is introduced: the water and gas terms are linearized exactly as *One linearization convention* prescribes, against a reference column independent of the arm's own length, which is what makes α a property of altitude alone and τ a pure exponential in L .

The arm's one approximation is constant density. On a spherical Earth the chord dips below its endpoints by the tangent-height depression

$$\Delta h \approx \frac{L^2}{8R_E}, \quad \text{mean sag over the chord} \approx \frac{2}{3}\Delta h$$

so the real path samples slightly *denser* air and the model **under-states** optical depth. The horizon guard bounds the error: $\Delta h < 100$ m computes clean, up to 2 km warns, beyond 2 km raises. On the 2 km water scale height that is a 3.4 % water-density error at the clean edge, 6.8 % on the longest grid arm ($L = 100$ km), and a factor ≈ 1.9 at the raise threshold — which is why 2 km raises rather than warns.

$\tau(2L) = \tau(L)^2$ holds exactly for this arm, and it is precisely where a correlated- k band model disagrees: strong lines saturate first and flux keeps leaking through the windows between them. That divergence is measured against the horizontal MODTRAN grid in the parity document §2.6.

The whole level path. The sky behind a level target is *not* an arm composed with a continuation joined at the target plane. A level ray is tangent at the chord **midpoint**, so the sensor sits on its descending half, and the traversed column is, per species,

$$S_i = 2 S(r_p; h_p \rightarrow h_{\text{arm}}) + S(r_p; h_{\text{arm}} \rightarrow h_{\text{top}}), \quad r_p = \sqrt{r_{\text{arm}}^2 - (L/2)^2}$$

evaluated once, in one place. Rooting a single ascending arc at the sensor — the up-looking branch's shape — would drop the arm entirely, recovering only 83.0 % of the true traversed molecular column for a 100 km arm at 3 km and 75.1 % for a 150 km arm at 10 km. A zero-length arm reduces the whole-path evaluator **exactly** to the grazing evaluator at $\zeta = \pi/2$, so the level and ascending sky topologies join without a step. The 8 km-at-sea-level case is degenerate — its perigee sits 1.3 m *below* the ellipsoid — so the integration floor is clamped at MSL and the model warns.

Sky radiance along the LOS. The radiance a receiver at h_{start} sees looking up along a ray of zenith ζ with nothing behind the atmosphere but cold space:

$$L_{\text{sky}}(h_{\text{start}}, \zeta) = \text{SegmentQuantities}(\text{ColumnSegmentSpec}(h_{\text{start}}, h_{\text{atm, top}}, \zeta)) \cdot L_{\text{toward lower}} + \tau_{\text{seg}} L_{\text{beyond}}$$

with $L_{\text{beyond}} \equiv 0$ — there is nothing behind the atmosphere but cold space.

The 2.7 K cosmic background contributes $< 10^{-9}$ W/m²/sr/μm anywhere in the 0.3–14 μm working range, so the composition is a no-op and the module is a thin, well-named wrapper rather than a second physics implementation. **The ray starts at the sensor, not the target:** a background behind a target cannot depend on where along the ray the target sits, and splitting one column at the target plane swaps part of a warm ground-anchored graybody for a cold target-anchored one — the segment model is not additive across a join.

Composition rules. For two adjoining segments A (nearer the source) and B :

$$\tau(A \cup B) = \tau_A \tau_B, \quad L(A \cup B) = L_A \tau_B + L_B$$

which, with the vacuum identities $\tau_V \equiv 1$, $L_V \equiv 0$, collapse an exo-altitude target's path to the published fields with no arithmetic performed at all.

4.2.13 Twilight — per-altitude illumination and the solar tangent transit

Whether a point is lit is decided per altitude rather than by a global $\theta_s < \pi/2$ bound, so solar zenith spans the closed $[0, \pi]$:

$$\text{sunlit}(h, \theta_s) \iff \theta_s \leq \frac{\pi}{2} \quad \text{or} \quad (R_E + h) \sin \theta_s \geq R_E \iff h \geq R_E (\sec \delta - 1), \quad \delta = \theta_s - \frac{\pi}{2}$$

so a 60 km booster is sunlit at 5° solar depression while the ground beneath it (shadow height ≈ 24 km) is not. The assumption is a **sharp terminator** — opaque sphere, point Sun, no refraction; the ≈ 200 m penumbral blur and the $\approx 0.5^\circ$ of unmodeled refractive lift are documented, not smoothed.

For a **sunlit** target with $\theta_s > \pi/2$ the direct beam is a tangent transit, not a descending column, so the solar transmittance is the two-arm decomposition

$$\tau_{\text{sun}} = \tau(\text{tangent} \rightarrow \text{target}) \cdot \tau(\text{tangent} \rightarrow \text{TOA})$$

about the tangent radius $r_0 = (R_E + h_{tgt}) \sin \theta_s$. For a **shadowed** target τ_{sun} is exactly 0 and the scattered-sky solar component is already identically zero; the thermal sky is untouched.

This branch is provisional. The twilight transit carries the largest optical depths anywhere in RADIANT (30–70 air masses), where both the exponential-in-column transmittance and the unmodeled refraction are at their worst. The twilight pair Q7/Q8 *was* delivered — with Card-3 ANGLE hand-set to 93° / 96° and LENN = 1, and the hand edit verified against the matrix by the Card-3 echo sweep — but both rows are dev_only: no library family and no radiometric parity test consumes them, so the transit's **transmittance remains unanchored** (parity document §1 and §3). Treat it as an order-of-magnitude bound.

Note also that RADIANT models the target as a horizontal Lambertian facet, so the direct solar term is multiplied by $\cos \theta_s$ clamped at zero: for any $\theta_s > \pi/2$ the direct term vanishes regardless of τ_{sun} , because the beam arrives from below the facet. τ_{sun} is still published correctly, because it is an inspectable physical quantity that a non-horizontal target model would consume.

4.3 The library-backed models

4.3.1 What a library family is

A family is a set of pre-computed MODTRAN states at discrete geometry points, plus a declaration of which geometry fields are **interpolation axes**. Each stored node carries a wavelength grid, a transmittance, a path radiance, a downwelling emission array, and the full five-field run geometry it was rendered at.

Shipped families, with their axes and direction:

Family	Direction	Axes	Coverage
us_standard_zenith_fan	down	path_zenith_rad	ground target, sensor 100 km, ζ 0–60°
midlat_summer_ladders	down	sensor_altitude_m, target_altitude_m	targets 0–29 km; sensor 35 km / 100 km / GEO; nadir
midlat_summer_sensor_ladder	down	sensor_altitude_m	ground target; sensor 3–100 km + GEO; nadir
midlat_summer_boost_offnadir	down	sensor_altitude_m, target_altitude_m, path_zenith_rad	targets 0–100 km; sensor 100 km / GEO; ζ 0–60°
midlat_summer_upwelling_offnadir	down	sensor_altitude_m, path_zenith_rad	ground target; sensor 10/100 km / GEO; ζ 0–60°
midlat_summer_boost_ladder	down	sensor_altitude_m, target_altitude_m	targets 0–100 km, 12 rungs; explicit-dir only
midlat_summer_uplooking_ladder	up	target_altitude_m	ground sensor, targets 0–20 km, vertical
midlat_summer_uplooking_zenith_fan	up	target_altitude_m, path_zenith_rad	ground sensor, targets 0–20 km, ζ 0–60°
midlat_summer_uplooking_sensor_ladder	up	sensor_altitude_m	observer 0–100 km, full column, fixed 48.2°
midlat_summer_sst_column_fan	up	path_zenith_rad	ground sensor, full column, ζ 0–78.5° (sec 1–5); explicit-dir only
midlat_summer_sst_column_fan_site900m	up	path_zenith_rad	900 m site, full column, ζ 0–78.5° (sec 1–5); explicit-dir only

Direction is a first-class property. An up-looking run measures one travel direction, so an up-looking family stores its radiance under a **different NPZ key** (`path_radiance_toward_lower`), carries a `los_direction` marker, and is served through a different query entry point; the two entry points refuse each other's families rather than reading the wrong product.

The two SST fans differ only in their rendered lower endpoint (0 m and 900 m) and are **siblings, not one family with a sensor axis**: two rungs cannot be interpolated between, the 900 m block has no 48.2° rung to make a rectangular grid with, and the 0 m fan had to stay byte-identical. A site at any other elevation is served by simple.

4.3.2 Geometry interpolation in $\log\tau$

Transmittance is interpolated in **optical-depth** space, not in τ :

$$\ln \tau_{\text{query}} = \text{interp}(\text{axes}; \ln \tau_{\text{nodes}}), \quad \tau_{\text{query}} = \exp(\ln \tau_{\text{query}})$$

Beer-Lambert gives $\tau = e^{-\text{OD}}$ with OD linear in path length and absorber amount, so linear interpolation of $\ln \tau$ preserves that linearity and linear interpolation of τ does not.

Path radiance and downwelling emission interpolate **linearly**. They are additive quantities with no Beer-Lambert exponential in path length, so log space would buy nothing and would misbehave at zeros.

No extrapolation. A query outside the convex hull of the nodes always raises. The one query served from outside the node span is not an exception but an application of the rule — see *Vacuum-equivalence identities*.

4.3.3 Wavelength resampling in log- τ , and why the two must commute

A chain runs on its own wavelength grid, which is normally not the stored one, so the geometry-interpolated spectra must be resampled. That resample runs in the **same** log- τ space:

$$\tau(\lambda_{\text{query}}) = \exp[\text{interp}_{\lambda}(\ln \tau_{\text{stored}})]$$

Both operations are then linear in $\ln \tau$, so they **commute** and the answer no longer depends on their order. Resampling linearly in τ instead returns the arithmetic mean of the bracketing samples where the physics gives the geometric mean — and by AM–GM the arithmetic mean is strictly larger, so the old convention was biased one way, upward.

The key equation the tests use as their truth anchor: for $\tau(\lambda) = e^{-a\lambda}$ the optical depth is *exactly* linear in λ , so a log-linear resample is exact at every query wavelength and a cell midpoint must return $\sqrt{\tau_i \tau_{i+1}}$.

The convention is universal across all three file-backed backends. `tabulated` and `modtran` each perform only one resample — no operation order to get wrong — but a different *convention* meant the same stored MODTRAN column returned different numbers depending on which backend served it. All three now carry every τ -like array (`transmittance`, τ_{up} , τ_{sun} , $\tau_{full,up}$) through one shared implementation. Scope, in both directions:

- **log- τ :** transmittance only, in every backend.
- **linear:** L_{path} , $L_{\text{atm,down}}$, and every irradiance.

A query on the file's own grid short-circuits the log round trip and is **bit-identical** to the stored array.

4.3.4 Zenith axes interpolate in $\sec \zeta$

A zenith-angle axis is mapped $\zeta \mapsto \sec \zeta$ before interpolation. Combined with the log- τ geometry interpolation above, this is Beer-Lambert-**exact** between nodes: the path length along a zenith axis is $\propto \sec \zeta$, so $\ln \tau$ linear in $\sec \zeta$ reproduces Beer at every query angle. Linear-in-angle carried a measured in-band bias at fan midpoints (parity document §2.10).

The transform is internal — user-facing coordinates, bounds and errors stay in radians. Angles $\geq \text{_MAX_ZENITH_RAD}$ ($\approx 88.8^\circ$) are **refused**, because $\sec \zeta$ diverges at the horizon.

Batch-2 gave the up-looking direction the same axis, sampled as a *uniform sec ladder* rather than a uniform angle ladder: $\sec \zeta = 1.0/1.4999/2.0$ for the target fan and $1/1.5/2/3/4/5$ for the full-column SST fan. Three sec rungs is the minimum that can **test** linearity in sec rather than assume it. The $85^\circ/88^\circ/89.5^\circ$ probes (M6–M8) were run but are deliberately **not shipped as nodes**: at and past the 88.8° ceiling the mapping is

unvalidated, and a shipped node there would sit inside a hull the interpolator may traverse. They serve as physics anchors instead (parity document §2.4).

4.3.5 TAU_FLOOR — a lower clamp only

Stored transmittance is floored at $\text{TAU_FLOOR} = 10^{-30}$ (equivalently $\text{OD} \approx 69$) before the logarithm, so $\ln \tau$ is finite everywhere and an opaque band resamples to that floor rather than to $-\infty$ or NaN. One definition, imported by every consumer.

The floor is deliberately **not** matched by a cap at 1.0. A $\tau > 1$ array is invalid data, and capping it would convert a mis-scaled tape7 into a plausible-looking column — exactly the kind of silent repair RADIANT forbids. Instead it survives the resample and fails loudly downstream. Negative τ raises inside the resample itself.

4.3.6 Vacuum-equivalence identities

Two identities let a family serve a query outside its node span *without* extrapolating, because the intervening path is provably vacuum. Both are gated on the shared constant $\text{_VACUUM_EQUIVALENT_ALTITUDE_M} = 100 \text{ km} = h_{\text{atm,top}}$.

Sensor axis. MODTRAN’s atmosphere ends at $h_{\text{atm,top}}$, so a sensor above a node at or above the column top sees an *exactly identical* column — the added path has zero extinction and zero emission. This is why the ladders duplicate their 100 km TOA state at a 40,000 km node: an orbital sensor then falls inside the hull, and the duplication is a physical identity, not an approximation. The same identity exempts such a query from the non-axis geometry-mismatch warning.

Target axis. The mirror. A family whose measured target ceiling reaches $h_{\text{atm,top}}$ integrated the *entire* column, so its top-of-column run **is** the observer leg for any exo-altitude target; the query is clamped to the ceiling node and the clamp is recorded in the segment provenance under `exo_target_vacuum_clamp`.

The guard asks the **family**, never a hard-coded name: `uplooking_target_ceiling_m` reads the highest target altitude the family’s own runs measure — from the `target_altitude_m` axis hull when it carries one, from the recorded fixed value when it does not. A **partial-column** family (ceiling below $h_{\text{atm,top}}$) is **refused**, with an actionable error naming its measured ceiling and the families that do serve the scene: composing its top rung with the vacuum identity would join a measured leg to an invented one. The clamp is gated at $h_{\text{atm,top}}$ and nowhere else — a 50 km target through a 20 km ladder is 40 km of real atmosphere and still fails the hull check, unchanged.

4.3.7 The hybrid up-looking composition

An up-looking run family is **one leg of data**, so an up-looking chain run on the interpolated backend is a declared hybrid: the family serves the leg it measured, and a `SimpleAtmosphere` companion — built pre-chain from the same `atmosphere.*` parameters a `model = "simple"` run would use — serves the rest.

Leg	Served by	Why
observer (sensor $\rightarrow$ target)	the up-looking run family	this <i>is</i> the rendered column, and it dominates a ground-to-air scene

Leg	Served by	Why
illumination (solar column + sky above the target)	the SimpleAtmosphere companion	no rung of a sensor→target ladder is the column <i>above</i> the target, and the down-looking proxy query an up-looking family would need is refused by construction
sky at aperture (sensor → $h_{\text{atm,top}}$)	the SimpleAtmosphere companion	reading a partial ladder's top rung as “the sky” would be extrapolation past the hull

Two independently-calibrated models in one answer is a real modeling compromise, so it is never silent: a `UserWarning` is raised, an INFO record is logged, and `stage_outputs["atmosphere"]["topology_provenance"]["backend_split"]` names which leg came from which model. Where the two models must agree — τ_{sun} , E_{TOA} , and both E_{sky} terms, all served by the companion alone — they are bit-identical.

The compromise is accepted only while it stays declared. The warning, the INFO record and the `backend_split` marker are part of what makes it acceptable and must never be softened into silence. The measured divergence between the two legs is in the parity document §2.11.

The one re-audit condition: the split exists because an up-looking family is one leg of data. A family that is self-contained — carrying its own solar column and its own sensor → $h_{\text{atm,top}}$ sky — or a scene whose target is a blackbody, where the illumination terms vanish, makes the companion unnecessary *for that scene*, and the split should be dropped there rather than declared.

A **level** path on an up-looking family is refused, not approximated: a level arm has zero vertical extent and a local zenith of $\pi/2$ everywhere, so no rung of a column ladder is that path and no interpolation between rungs produces it.

4.4 What the models do not represent

Recorded here because a physics document that omits its own boundaries is misleading. Each item is tracked in the repository, where the measured consequences are tabulated in full.

- **Line structure inside a calibrated region.** The simple model's spectral shape is flat within each of the 17 regions; the 0.04 μm edge ramps remove the discontinuity, not the underlying flatness. It is the dominant residual of the thermal path radiance.
- **Multiple scattering.** The single-scatter source under-predicts the daytime VIS/NIR sky; the sub-3 μm provisional warning is what says so to an operator.
- **Refraction.** The geometry is unrefracted. It is the dominant geometric error inside the horizon guard's warn band, and the decks that would calibrate it have not been run.
- **Stratospheric structure.** The fixed-lapse ICAO profile is isothermal above the tropopause, so real stratospheric warming is not represented. Measured on the downwelling ladder: MODTRAN's recovered emission temperature reaches 268 K at the 50 km stratopause and falls to 178 K at 80 km, while the profile answers 222.65 K at every one of those altitudes. This is why the five stratospheric rungs of the downwelling ladder are insensitive to *any* emission-placement choice — every candidate form returns the same temperature there — and why they carry no information about the hemispheric downwelling terms' fitted constants.

- **Grazing-arc opacity distribution.** A grazing arc's air lies along the arc, not along the vertical between its endpoints; the emission weighting is approximate there, though the *total* optical depth is exact. Measured on a ground-rooted column at 85/88/89.5°, the two placements differ by $\leq 1.2\%$ in band-mean thermal radiance — under the 3–6 % model/MODTRAN residual there, so those runs cannot settle it. The effect is only measurable on an arc rooted at an elevated tangent point *below* the tropopause (modeled separation 7.9 %/9.8 % at 5 km, 16.1 %/13.1 % at 8 km, MWIR/LWIR, and exactly 0 at 15 km for the isothermal reason above); the reference runs for that geometry are authored and not yet executed.
 - **O₃ emission altitude.** The well-mixed-gas floor lumps CO₂/N₂O/CH₄ with O₃, which peaks near 25 km, so 9.6 μm emission is placed too low. The τ side is settled: the 8–10 μm region is split at the band edges (*The well-mixed-gas absorption floor*), so the ozone share of the in-feature floor is arithmetic — 0.832 — rather than a free parameter. The *placement* is not: the identified opacity still rides the molecular scale height, so the model puts ozone near 8 km instead of 25 km. Measured on fourteen matched pairs with the split table in place, the 9.4–9.9 μm feature reads RMS $|\ln \text{ratio}|$ 0.3581 (it was 0.1519 when the flat slab under-supplied the in-band opacity), and moving the τ -determined share onto a 25 km layer recovers it to 0.1456 — within 7 % of the sweep's unconstrained best, 0.1365 at share 0.70. The remaining degeneracy in the layer's center and width is unchanged in kind (0.133–0.191 across 20–30 km centers and 3–8 km widths) because the profile is isothermal wherever the layer sits.
 - **Polarization, 3D/heterogeneous atmospheres, time dependence, adjacency, aurora/airglow, cloud microphysics.** Out of scope for this version of the model.
-

4.5 References

This chapter cites, by the keys of the manual's single References list: [Hansen & Travis 1974] and [Bucholtz 1995] for the Rayleigh whole-atmosphere optical depth, [McClatchey 1972] for the standard-profile water columns, [Henyey & Greenstein 1941] for the aerosol phase function, [Elsasser 1942] for the diffusivity-factor context of D , and [Chapman 1931] for the grazing-limit anchor of the spherical slant column.

Spatial Model

PSF construction, the MTF budget, ensquared energy, sampling, and the dual-path architecture that keeps them consistent — as implemented in RADIANT. Numeric anchors are literature values re-derived independently of the implementation, then checked against it.

Symbols used throughout: λ wavelength [μm] ($\lambda_m = \lambda \cdot 10^{-6}$ m where SI is needed); D aperture diameter [m]; f focal length [m]; $F_{\#} = f/D$; ϵ central-obscuration ratio; p pixel pitch [m]; ν spatial frequency at the focal plane [cy/m] (cy/mm in worked numbers); $\nu_c = 1/(\lambda F_{\#})$ optical cutoff; $\nu_{Nyq} = 1/(2p)$ Nyquist.

5.1 The two spatial paths

RADIANT maintains **two parallel spatial representations, both rooted in the same complex pupil**:

- the **PSF path** — every spatial degradation enters as a convolution kernel on one EffectivePSF; EE_box, RER, FWHM, Strehl, LSF and ERF are computed only from it;
- the **MTF product path** — the frequency-domain budget $\text{MTF}_{sys}(\nu) = \prod_i \text{MTF}_i(\nu)$ consumed by MTF-at-Nyquist, folded MTF, and GIQE/NIIRS.

The two are the same information: for incoherent imaging the OTF is simultaneously the Fourier transform of the PSF and the normalized autocorrelation of the complex pupil $P = A e^{i2\pi W/\lambda}$,

$$\text{OTF}(\nu) = \frac{\mathcal{F}\{\text{PSF}\}(\nu)}{\mathcal{F}\{\text{PSF}\}(0)} = \frac{(P \star P)(\lambda f \nu)}{(P \star P)(0)},$$

so the FFT of the fully convolved EffectivePSF must agree with the MTF product. RADIANT checks this on every chain run that computes the spatial path, to an absolute tolerance of 2×10^{-2} ($\approx 2\times$ the measured full-chain discretization floor).

Why a separate $\text{MTF}_{diffraction} \times \text{MTF}_{aberration}$ factoring is forbidden: aberrations are a phase modification *inside* the pupil, entering the autocorrelation as a phase *difference across the shear* under the integral — they cannot be factored into (clear-aperture MTF) $\times$ (aberration-only multiplier). True aberrated OTFs can go negative (defocus contrast reversal); no positive scalar multiplier can reproduce that. Only genuinely independent *image-plane convolutions* (detector aperture, diffusion, IPC, jitter, smear) and statistically independent ensemble averages (turbulence) multiply. The one deliberate MTF-only term is TDI mis-registration (*TDI mis-registration — the one MTF-only term*, below) — a readout-timing effect with no instantaneous kernel, excluded from the consistency comparison.

5.2 Diffraction-limited PSF

Equation. Unaberrated circular pupil (Airy):

$$I(r) = I_0 \left[\frac{2J_1(v)}{v} \right]^2, \quad v = \frac{\pi r}{\lambda F_{\#}},$$

first zero at $v_1 = 3.8317060$, i.e. $r_1 = 1.2196699 \lambda F_{\#}$. Encircled energy $EE(v) = 1 - J_0^2(v) - J_1^2(v)$ (unobscured only): 0.837785 at the first dark ring. With central obscuration ϵ :

$$I(v) = \frac{I_0}{(1 - \epsilon^2)^2} \left[\frac{2J_1(v)}{v} - \epsilon^2 \frac{2J_1(\epsilon v)}{\epsilon v} \right]^2$$

— core narrows, energy moves into the rings.

Implementation. RADIANT does not evaluate the Bessel form; it FFT-propagates the sampled complex pupil (aperture mask incl. obscuration and spider vanes, plus wavefront phase) and takes $|\cdot|^2$ — the Airy pattern is the *test oracle*, not the algorithm. Polychromatic PSFs are photon-flux-weighted sums of monochromatic PSFs, one per wavelength on a sub-sampled grid whose size is a parameter.

Assumptions & validity. Scalar Fraunhofer diffraction; $F_{\#} \gtrsim 2$ for the scalar approximation; grid effects controlled by the pupil-sampling and PSF-oversampling parameters, whose defaults are 128 pupil samples across and 8× oversampling.

Pitfalls. 1.22 vs the exact 1.2196699 (0.03 % — matters for tight baselines); amplitude vs irradiance PSF; forgetting the $(1 - \epsilon^2)^{-2}$ energy renormalization; applying the unobscured EE formula to an annular pupil.

Numeric anchor. $r_1 = 19.5147 \mu\text{m}$ at $\lambda = 4 \mu\text{m}$, $F_{\#} = 4$.

5.3 Optical MTF from the pupil autocorrelation

Equation. Clear circular pupil, incoherent:

$$\text{MTF}(\tilde{\nu}) = \frac{2}{\pi} \left[\arccos \tilde{\nu} - \tilde{\nu} \sqrt{1 - \tilde{\nu}^2} \right], \quad \tilde{\nu} = \nu / \nu_c \leq 1, \quad \nu_c = \frac{1}{\lambda F_{\#}}.$$

RADIANT computes the general case (obscuration, spiders, aberrations) by direct autocorrelation of the complex pupil — a single `MTF_optics` term, never a factored product (*The two spatial paths*, above).

Assumptions & validity. Incoherent illumination (the coherent cutoff is $\nu_c/2$ — a classic factor-2 trap); monochromatic per- λ , spectrally weighted for broadband.

Pitfalls. λ -unit slips in ν_c (λ must be in mm for cy/mm); NaN instead of 0 above cutoff; treating the annular-pupil MTF as clear-MTF $\times$ scalar (the true ratio crosses 1: ≈ 0.78 at $0.3\nu_c$ but ≈ 1.10 at $0.8\nu_c$ for $\epsilon = 0.3$).

Numeric anchor. $\text{MTF}(0.5\nu_c) = \frac{2}{\pi} \left(\frac{\pi}{3} - \frac{\sqrt{3}}{4} \right) = 0.391002$; $\nu_c = 62.5 \text{ cy/mm}$ at $\lambda = 4 \mu\text{m}$, $F_{\#} = 4$.

5.4 Wavefront error, Zernike modes, and Strehl

Equations. Wavefront OPD enters the pupil phase as $e^{i2\pi W/\lambda}$. Noll-indexed Zernike defocus $Z_4 = \sqrt{3}(2\rho^2 - 1)$: with $W = a_4 Z_4$, the RMS OPD is exactly a_4 and the peak-to-valley is $2\sqrt{3}\sigma$. Maréchal Strehl approximation from RMS OPD σ_{OPD} :

$$S \approx \exp\left[-\left(\frac{2\pi\sigma_{OPD}}{\lambda}\right)^2\right].$$

RADIANT reports two Strehls: the **PSF-derived** `strehl` — degraded-PSF peak over the diffraction-limited `reference_psf` peak, with the *same* detector kernels applied to both so detector effects cancel — and the analytic `strehl_marechal` diagnostic.

Assumptions & validity. Maréchal reliable for $\sigma \lesssim \lambda/10$; Noll normalization (coefficient = RMS) — the Wyant convention differs by $\sqrt{3}$ -type factors. Zernike modes are orthonormal on the unobscured unit disk; annular pupils strictly need annular polynomials.

Pitfalls. Waves vs radians in the Maréchal exponent (factor $(2\pi)^2$); P-V vs RMS (the $2\sqrt{3}$ is defocus-specific); using Maréchal to reconstruct a PSF or MTF (it is a scalar diagnostic only).

Numeric anchors. $S(\sigma = \lambda/14) = 0.817569$ (the classic ≈ 0.8 diffraction-limited threshold); quarter-wave P-V defocus $\rightarrow S = 0.814$.

5.5 Detector-plane kernels: pixel aperture, diffusion, IPC

Pixel aperture. A photosite of linear width w integrates the image — a rect convolution:

$$\text{MTF}_{det}(\nu) = \left| \frac{\sin(\pi w \nu)}{\pi w \nu} \right|,$$

first zero at $\nu = 1/w$. With areal fill factor FF, $w = p\sqrt{\text{FF}}$ — the same width the PSF path's pixel kernel uses, sampled by area overlap, so the two paths agree.

Charge diffusion. Gaussian carrier spread of RMS σ_d : $\text{MTF}_{diff}(\nu) = \exp(-2\pi^2\sigma_d^2\nu^2)$.

Inter-pixel capacitance. Pitch-spaced coupling kernel with fraction α : $\text{MTF}_{IPC}(\nu) = (1 - 4\alpha) + 2\alpha \cos(2\pi\nu p)$ per axis (nearest-neighbor form).

Pitfalls. The sinc convention: NumPy's `np.sinc(x)` already includes π — `np.sinc(w·v)` is correct, `np.sinc(π·w·v)` double-counts π and moves the first zero to $1/(\pi w)$. Pitch vs aperture width when $\text{FF} < 1$. IPC and diffusion both exist in kernel form for the PSF path and analytic form for the MTF path — adding one side only breaks the agreement between them.

Numeric anchor. 100 % fill at Nyquist: $\sin(\pi/2)/(\pi/2) = 2/\pi = 0.636620$.

5.6 Platform kernels: jitter and smear

Jitter (random LOS motion, many cycles per integration): Gaussian blur of RMS σ at the focal plane ($\sigma = \sigma_j f$ for the RMS angular jitter σ_j of the notation table):

$$\text{MTF}_{jit}(\nu) = \exp(-2\pi^2\sigma^2\nu^2).$$

Smear (deterministic image motion $d = v_{image} t_{int}$ during integration): rect blur,

$$\text{MTF}_{smear}(\nu) = \left| \frac{\sin(\pi d \nu)}{\pi d \nu} \right|,$$

applied along the motion direction only. v_{image} derives from the **ground-track** velocity ($v_g = v R/(R+h)$, see the geometry chapter's *GSD — nadir and off-nadir*) times the magnification f/R_s — using orbital v instead of v_g inflates smear by h/R (+7.8 % at 500 km).

Assumptions & validity. Jitter Gaussian form requires jitter frequency $\gg 1/t_{int}$; comparable-period motion is neither pure jitter nor pure smear. A $d/\sqrt{12}$ “equivalent Gaussian” for smear is a second-moment match, not the correct MTF.

Pitfalls. The jitter exponent is $2\pi^2\sigma^2\nu^2$ exactly (equivalently $e^{-(2\pi\sigma\nu)^2/2}$); σ vs FWHM ($\times 2.3548$); a jitter MTF with a zero crossing indicates a wrongly applied sinc model.

Numeric anchors. $\sigma = 0.25 p$ at Nyquist: $e^{-\pi^2/32} = 0.734603$; $d = 0.5 p$ at Nyquist: $\sin(\pi/4)/(\pi/4) = 0.900316$.

5.7 Atmospheric turbulence

Equations. Fried parameter $r_0 \propto \lambda^{6/5}$ (so $r_0(\lambda_2) = r_0(\lambda_1)(\lambda_2/\lambda_1)^{6/5}$); long-exposure Kolmogorov MTF in angular frequency f_a [cy/rad]:

$$\text{MTF}_{LE}(f_a) = \exp \left[-3.44 \left(\frac{\lambda f_a}{r_0} \right)^{5/3} \right].$$

The 3.44 is half the structure-function constant 6.88. Turbulence is the one contributor legitimately multiplied into the budget without pupil-level treatment: in the long-exposure ensemble average the atmosphere is statistically independent of the pupil. The PSF path receives the matching kernel.

Assumptions & validity. Kolmogorov spectrum (infinite outer scale), weak fluctuations, full tilt averaging (long exposure). Short-exposure (tilt-removed) imaging needs Fried's corrected form — not the shipped default.

Pitfalls. Quoting r_0 at $0.5 \mu\text{m}$ and using it unscaled in the IR (at $4 \mu\text{m}$, r_0 is $12\times$ larger); 3.44 vs 6.88; gating on platform type instead of $r_0 > 0$ (RADIANT gates on the parameter `atmosphere.r0_m`, and the MTF term is written by **PerformanceStage**, not **AtmosphereStage** — the atmosphere stage publishes `r0_m` only).

Numeric anchors. $\text{MTF}_{LE} = 0.338398$ at $\lambda f_a / r_0 = 0.5$; $r_0 = 0.10 \text{ m} @ 0.5 \text{ } \mu\text{m} \rightarrow 1.21257 \text{ m} @ 4 \text{ } \mu\text{m}$.

5.8 TDI mis-registration — the one MTF-only term

Equation. Cross-track drift angle θ over N TDI stages displaces the N summed samples by $d = p \tan \theta$ each; averaging N displaced copies gives the Dirichlet kernel, which for sub-pixel per-stage drift reduces to a sinc in the **total** drift $Np\theta$:

$$\text{MTF}_{TDI}(\nu) \approx \left| \frac{\sin(\pi \nu N p \theta)}{\pi \nu N p \theta} \right|.$$

This is a readout-timing/registration effect on the time-aggregated image — there is no instantaneous spatial kernel, so it enters **only** the MTF product and is excluded from the Rule-4 consistency comparison.

Pitfalls. Tolerancing θ without N (degradation scales with total drift, so the θ tolerance tightens as N grows); the discrete Dirichlet form exceeds the sinc form by up to $\sim 0.6 \%$ at small N .

Numeric anchor. Total drift $0.5 p$ at Nyquist, $N = 16$: 0.9007 (Dirichlet) vs 0.900316 (sinc).

5.9 Ensquared energy (EE_box)

Equation. Fraction of the fully degraded PSF's energy inside the centered $n \times n$ pixel box:

$$EE_{n \times n} = \int_{-np/2}^{np/2} \int_{-np/2}^{np/2} \text{PSF}(x, y) dx dy.$$

Computed in PlatformStage from the **fully degraded** EffectivePSF (jitter, smear, turbulence included) and applied exactly once, in SpectralIntegrationStage, to point-source and sub-pixel target signals only — never to the background term, never in the extended regime.

Discretization. Each PSF cell is weighted by the fraction of its area inside the box (cell-area overlap) — second-order accurate; the earlier full-weight edge-cell scheme carried an $O(dx)$ bias ($+24 \%$ at $Q=2$ at default sampling) that overstated point-source SNR.

Pitfalls. Ensquared $\neq$ encircled (a square of side p is not a circle of diameter p — quoting the 83.8 % first-ring figure for a pixel box is a category error); Airy ring tails decay as $1/u^2$, so truncated normalization biases EE.

Numeric anchor. Unaberrated Airy at critical sampling ($Q = 2$): $EE_{1 \times 1} = 0.177327$ — only $\sim 18 \%$ of a point source's energy lands in the center pixel.

5.10 Sampling: Nyquist, Q, and folded MTF

Equations. $\nu_{Nyq} = 1/(2p)$;

$$Q = \frac{\lambda F_{\#}}{p}, \quad \frac{\nu_c}{\nu_{Nyq}} = \frac{2}{Q}.$$

- $Q = 2$: **critically sampled** — the optics pass nothing above Nyquist; no aliasing.
- $Q < 2$: undersampled (typical EO design point, chosen for SNR/GSD); the band $[\nu_{Nyq}, \nu_c]$ folds.
- $Q > 2$: oversampled.

The folded MTF adds the aliased response back onto the baseband. Sampling on a pitch p replicates the pre-sampling spectrum at integer multiples of the **sampling** frequency $\nu_s = 1/p = 2\nu_{Nyq}$, so

$$\text{MTF}_{fold}(\nu) = \sum_{k=-N}^{+N} \text{MTF}_{opt}(|\nu + k \nu_s|).$$

At $\nu = \nu_{Nyq}$ the $k = -1$ replica lands back on ν_{Nyq} , so the folded value there is twice the pre-sampling MTF (alias fraction $\rightarrow 1/2$) whenever the higher orders are negligible — and it is *zero* for optics that cut off below Nyquist, which is the oversampled sanity check. The alias fraction $(\text{MTF}_{fold} - \text{MTF}_{opt})/\text{MTF}_{fold}$ is evaluated only where $\text{MTF}_{fold} > 10^{-9}$ of its DC value; below that floor both terms are round-off, and the reported fraction is exactly zero — an oversampled design has no aliased energy. GIQE consumes RER from the PSF path and the MTF budget, as the Performance Metrics chapter describes.

Pitfalls. “Q = 1 critical” is a different (sampling-frequency) convention — RADIANT’s statement is $\nu_c/\nu_{Nyq} = 2/Q$, critical at $Q = 2$; sampling frequency $1/p$ vs Nyquist $1/(2p)$ — replicating the spectrum at ν_{Nyq} instead of ν_s puts the $k = -1$ copy on DC and adds $\text{MTF}(0) = 1$ to every system; ground-projecting Nyquist twice.

Numeric anchor. $\lambda = 0.55 \mu\text{m}$, $F_{\#} = 4$, $p = 10 \mu\text{m}$: $Q = 0.22$ — $\nu_c/\nu_{Nyq} = 9.1$, heavily aliased, the classic high-resolution VNIR design point.

5.11 RER and the edge response

Equation. The edge response function is the cumulative integral of the line spread function (itself the 1-D projection of the PSF); the relative edge response per axis is

$$\text{RER}_x = \text{ERF}_x(+p/2) - \text{ERF}_x(-p/2),$$

and the reported scalar is the **geometric mean** $\sqrt{\text{RER}_x \text{RER}_y}$ (GIQE-5 usage). All of LSF/ERF/RER derive from that one degraded PSF.

5.12 Scope and deferred effects

No ghost images, no measured-BSDF scatter beyond the total-integrated-scatter and halo model, no chromatic aberration model, no short-exposure turbulence. The project's scope-decision register records the full list.

Noise Model

Complete noise taxonomy, equations, scaling rules, and acquisition timing as implemented in RADIANT. All noise in electrons RMS, the manual’s canonical noise unit. Numeric anchors were re-derived from the literature independently of the implementation, then checked against it.

6.1 Overview — RSS composition

RADIANT models 16 independent noise sources, each in e- RMS. With statistical independence, variances add:

$$\sigma_{\text{total}} = \sqrt{\sum_{i=1}^{16} \sigma_i^2}$$

The load-bearing assumption is independence — each mechanism arises in a physically distinct process, so cross-covariances vanish. It is deliberately violated by fixed-pattern terms across frames of the same pixel (they are a static spatial pattern: they belong in a single-frame spatial budget but do not average down over frames — see TDI/co-add scaling below). All terms are input-referred to electrons at the sense node before combining; mixing DN-domain and e-domain terms in one RSS is the classic error.

Terms are classed **temporal** (frame-independent, average down as $\sqrt{N}$) or **spatial** (fixed-pattern).

Numeric anchor. $N_{\text{target}} = 10000$, $N_{bg} = 5000$, $N_{\text{dark}} = 1000$ e-, $\sigma_{\text{read}} = 50$ e-, PRNU 0.1 % ($\sigma_{PRNU} = 10$ e-): $\sigma_{\text{tot}} = \sqrt{16000 + 2500 + 100} = 136.381817$ e- RMS.

6.2 Noise Taxonomy

6.2.1 Photon shot noise (4 terms)

Poisson statistics of photon arrival: for N_e photoelectrons, $\sigma = \sqrt{N_e}$. QE < 1 preserves Poisson statistics (binomial thinning), so the square root is taken of **electrons**, never photons or DN ($\sigma_{DN} = \sqrt{S_{DN}/g}$, not $\sqrt{S_{DN}}$).

- Signal shot: $\sigma = \sqrt{S}$ (target electrons)
- Background shot: $\sigma = \sqrt{S_{bg}}$
- Nearfield shot: $\sqrt{S_{nf}}$ (warm-optics self-emission)
- Straylight shot: $\sqrt{S_{stray}}$

Pitfalls. Adding σ 's instead of variances; dark-frame subtraction removes the dark *mean*, never its shot noise (single-shot subtraction doubles the dark shot variance).

6.2.2 Detector-material noise (4 terms)

Dark shot: $\sigma = \sqrt{J_{\text{dark}} t_{\text{int}}}$ with J_{dark} in e-/s. Temperature scaling is Arrhenius:

$$J(T) = J_{\text{ref}} \exp \left[\frac{E_a}{k_B} \left(\frac{1}{T_{\text{ref}}} - \frac{1}{T} \right) \right]$$

(`detector.dark_activation_energy_eV`, `dark_reference_temperature_K`). Physics note: diffusion-limited material has $E_a \approx E_g$, generation–recombination-limited $E_a \approx E_g/2$ — conflating them mispredicts the cooling benefit by squaring the true ratio (audit anchor: $E_g = 0.1$ eV diffusion, 77→84 K: dark current $\times 4.56$).

G-R noise: $\sigma_{gr} = \sqrt{2g_{gr}J_{\text{dark}}t_{\text{int}}}$ ($g_{gr} \approx 0$ PV, ~ 1 PC — the same generation-AND-recombination doubling that costs photoconductors $\sqrt{2}$ in D*).

Johnson noise (from R_0A) and **1/f noise** (flicker coefficient integrated over a frequency band) follow the standard material forms.

6.2.3 ROIC noise (3 terms)

Read noise: fixed floor `read_noise_e_rms`; includes CDS benefit when `read_noise_is_post_cds = True` (default).

kTC reset noise: $\sigma_{kTC} = \sqrt{k_B T C_{\text{node}}}/q$ — suppressed to zero when CDS is enabled (`readout.cds_enabled = 1`, default).

ADC quantization:

$$\sigma_{ADC} = \frac{g}{\sqrt{12}}, \quad g = \frac{N_{\text{fullwell}}}{2^n} [\text{e-/DN}].$$

Valid when input noise $\gtrsim 1$ LSB (dithered); the error is uniform on $\pm \text{LSB}/2$ and its variance is $\text{LSB}^2/12$.

Pitfalls. Dividing by 12 instead of $\sqrt{12}$ (3.46 $\times$ understatement); counting quantization twice when a vendor's measured “read noise in e-” already includes the ADC; leaving σ_{ADC} in DN when the RSS is in e-.

Numeric anchor. 14-bit, 100 ke- full well: $g = 6.10352$ e-/DN, $\sigma_{ADC} = 1.76193$ e- RMS.

6.2.4 Spatial (fixed-pattern) noise (3 terms)

PRNU — multiplicative gain dispersion, **linear in signal** (not $\sqrt{N}$): $\sigma_{PRNU} = k S$ with k the residual non-uniformity after NUC. Shot/PRNU crossover at $N = 1/k^2$ (0.1 % $\rightarrow 10^6$ e-): above it, more integration cannot improve single-frame spatial SNR — only better flat-fielding can.

DSNU — dark-signal non-uniformity, a fixed spatial floor (`dsnu_e_rms`).

Clutter — scene-induced spatial variance, $\sigma_c \cdot S_{bg}$ (`detector.clutter_sigma`).

Pitfalls. Applying $\sqrt{\cdot}$ to PRNU; claiming temporal co-adding reduces FPN (perfectly frame-correlated for a staring pixel); applying k to signal+dark (dark dispersion is DSNU's job).

6.2.5 Other (2 terms)

Persistence — residual charge from the prior exposure, on an exponential-decay model — and **ROIC glow shot**, $\sqrt{R_{glow} t_{int}}$.

6.3 Acquisition scaling

6.3.1 TDI and co-adding

Summing N stages/frames: signal $\times N$; uncorrelated temporal noise $\times \sqrt{N}$; SNR $\times \sqrt{N}$. The architecture determines the read-noise and FPN behavior:

Noise category	Charge-domain (analog) TDI	Digital co-add
Signal	$\times N$	$\times N$
Shot-like temporal	$\times \sqrt{N}$	$\times \sqrt{N}$
Read / kTC / quantization	$\times 1$ (one read)	$\times \sqrt{N}$ (read per frame)
PRNU/DSNU across distinct pixels (cross-scan TDI)	effective $k/\sqrt{N}$	—
FPN, same pixels (staring co-add)	—	$\times N$ (coherent; no SNR gain)
Clutter (scene-correlated)	$\times N$	$\times N$

Well-fill cap: charge-domain TDI accumulates into one well — $N(S_1 + N_{d,1}) \leq N_{fullwell}$ bounds the useful N (audit anchor: 10 ke- per stage against a 100 ke- well caps N at 10, SNR gain at $\sqrt{10}$).

6.3.2 Binning

On-chip binning merges charge before read (read noise once per binned super-pixel); off-chip binning averages after read (read noise per contributing pixel, $\times \sqrt{n}$ in the sum). Signal scales with the binned area in both.

6.3.3 Frame timing

Frame period is stored independently of integration time:

$$f_{frame} = \frac{1}{T_{frame}}, \quad \text{duty} = \frac{t_{int}}{T_{frame}} \leq 1.$$

Unset `readout.frame_period_s` (0.0) defaults to t_{int} (duty 1.0, continuous readout); duty > 1 is rejected with an actionable error. Outputs `frame_period_s/frame_rate_hz/duty_cycle/frame_period_` defaulted in `stage_outputs["readout"]`.

6.4 Noise regime selection

`detector.noise_regime` controls the budget composition: "imaging" (temporal terms only, default) vs "detection" (adds PRNU, DSNU, clutter). The regime distinction exists because single-frame detection against a scene sees the spatial pattern as noise, while a calibrated imaging chain removes the static component.

6.5 When each term dominates

Regime	Dominant noise	Typical scenario
BLIP (background-limited)	Background shot	IR, long t_{int} , low read noise — the Performance Metrics chapter gives the BLIP criterion and f_{BLIP}
Read-noise limited	Read	Short t_{int} , low flux
Dark-current limited	Dark shot	Long t_{int} , warm detector
FPN-limited	PRNU/clutter	$S > 1/k^2$, no scan averaging

A system can exit BLIP *from above* as flux grows: $\sigma_{PRNU} = kN_{bg}$ outpaces $\sqrt{N_{bg}}$ past the crossover.

6.6 Parameter cross-reference

Volume III's Parameter Reference lists every detector and readout parameter with its type, default, and bounds.

Calibration Error Model

What a radiometric calibration leaves behind. A non-uniformity correction removes the pixel-to-pixel gain and offset dispersion exactly at the signal levels where it was derived; everywhere else a residual survives, and that residual — not the temporal noise floor — frequently sets the achieved NEDT of a cooled infrared system. This chapter states the residual model, the drift terms, the accuracy (bias) budget, and how each enters the noise composition of Chapter 6.

Two facts drive the whole design. **Correlated errors do not average down:** TDI stages, co-added frames, and binned pixels reduce temporal noise as $\sqrt{N}$, but a calibration residual is the same error in every frame, so it is exempt. **Precision is not accuracy:** a cal-source temperature uncertainty moves the entire array’s radiometric scale and must never be root-sum-squared into σ_{tot} , where it would masquerade as noise that integration could beat down.

7.1 Overview — a terms-only stage after readout

The calibration stage sits between readout and the metric layer, and it transforms nothing: it collapses no spectrum, moves no frame, and — this matters — contributes no MTF term. Residual fixed-pattern noise is spatial *noise*, not a spatial *degradation*: it has no point spread function kernel, so neither spatial path of Chapter 5 gains a contributor and the dual-path consistency check is untouched.

What the stage contributes is three kinds of term:

- **Residual noise terms** [e- RMS], all classed spatial and all calibration-correlated: post-NUC residual, cal-source non-uniformity, gain drift, offset drift, and — on the internal-shutter cal path only — narcissus fixed-pattern noise.
- **Bias terms** [dimensionless fractions of radiance], the accuracy budget: cal-source temperature and emissivity, spectral calibration, internal-cal offset, and absolute gain.
- **Published totals:** the calibration RSS σ_{cal} , the post-calibration total σ_{tot} , the bias total, and the calibration-limited NEDT.

The ordering is the physics. Readout applies the $\sqrt{N}$ scaling first; calibration appends afterwards. No term needs to be tagged “do not average me” for the exemption to hold — the chain position enforces it. Equivalently: the residual’s size *relative to the signal* is invariant in N , because both the residual and the full-scale reference it is computed against live in the same summed domain.

The type separation is the guarantee. A noise term and a bias term are different types. SNR and NEDT read noise terms only; the radiometric-accuracy metric reads bias terms only. Nothing in the model can quietly RSS one into the other.

Noise-regime interplay. The imaging noise regime excludes the raw PRNU, DSNU, and clutter terms on the grounds that they are “calibrated out”. The calibration residuals are precisely what survives that calibration, so they enter the total in *both* the imaging and the detection regimes. With an active scheme the detector’s PRNU and DSNU are re-read as **pre-correction** dispersions feeding the residual model rather than as noise terms of their own, so nothing is counted twice.

7.2 Cal points — where the correction is anchored

A correction scheme is defined by the signal levels at which it is exact. RADIANT supports one-, two-, and three-point schemes, and two ways of declaring the points.

7.2.1 Temperature-declared points

The physical form: the cal source is a blackbody at a stated temperature. Because the stage runs after the single spectral collapse, only scalars are available, so the mapping from temperature to signal is a band photon-radiance ratio anchored at the scene:

$$S(T_{cal}) = S_{scene} \frac{B_q(T_{cal})}{B_q(T_{scene})}, \quad B_q(T) = \int_{\lambda_1}^{\lambda_2} \frac{B(\lambda, T)}{E_{ph}(\lambda)} d\lambda$$

with B_q in photon/s/m²/sr and $E_{ph} = hc/\lambda_m$ [J]. Photon weighting, not energy weighting, is what the band integral uses, because the collected signal counts photons. Cal-point *fluxes* are never independent inputs in this mode — declaring both a cal temperature and its flux would over-specify the source.

Assumption. The spectral response $QE(\lambda)\tau(\lambda)$ varies slowly enough across the band that a flat-weighted photon-radiance ratio stands in for the response-weighted signal ratio. Only the *ratio* enters, and numerator and denominator are anchored at nearby temperatures, so the first-order response slope largely cancels; the approximation degrades for a strongly sloped response curve.

Numeric anchor. MWIR band 3.0–5.0 μm , scene at 300.0 K: $S(280 \text{ K}) = 0.46475 S_{scene}$ and $S(320 \text{ K}) = 1.96525 S_{scene}$. For a scene signal of 5.000×10^4 e- this places the cal points at $S_1 = 2.3238 \times 10^4$ e- and $S_2 = 9.8262 \times 10^4$ e-.

7.2.2 Flux-declared points

The laboratory form: an integrating sphere or flat-field source produces a known level relative to the scene, with no thermal anchor at all. The cal points are then declared directly as fractions of the scene signal,

$$S_i = f_i S_{scene}, \quad f_i \text{ dimensionless,}$$

with $f_1 < f_2 < f_3$ required. In this mode every temperature-anchored input — the cal temperatures, the source temperature uncertainty, the source non-uniformity in K, the band-center uncertainty, and the source emissivity uncertainty — is **rejected** as over-specification rather than silently ignored. A set parameter that does nothing is exactly the failure class the model exists to prevent. Drift, the direct gain-uncertainty bias, and the internal-shutter path carry no temperature anchor and behave identically in both modes.

7.2.3 When the thermal anchor does not describe the scene

The temperature mapping anchors at the declared scene temperature. On a reflective scene that temperature describes none of the collected signal: a visible/near-infrared band looking at a 300 K surface collects Kirchhoff-reflected sunlight, while a 300 K blackbody puts essentially nothing into that band. The diagnostic quantity is the in-band share of the scene-temperature blackbody's photon exitance,

$$\Phi_{band}(T) = \frac{\int_{\lambda_1}^{\lambda_2} B_q(\lambda, T) d\lambda}{\int_{0.05 \mu m}^{1000 \mu m} B_q(\lambda, T) d\lambda} \quad [\text{dimensionless}]$$

Numeric anchor. $\Phi_{band}(300 \text{ K})$ is 4.30×10^{-22} for a 0.50–0.85 μm band, 3.20×10^{-3} for a 3.0–5.0 μm band; a 77 K laboratory scene in an 8–12 μm band gives 1.99×10^{-5} . The discrimination spans nearly twenty orders of magnitude, so the 10^{-9} floor the model warns below is a diagnostic bound, not tuned physics.

Below that floor — or for a target declared purely reflective — the cal points are deterministic stand-ins: the residual *structure* (the shape of the parabola, the $\sqrt{N}$ exemption, the plateau between cal points) survives, but the absolute level rides the stand-in. The run proceeds under a loud advisory rather than silently, and the correct remedy is to declare the cal points as flux fractions.

7.2.4 The full-scale reference

The nonlinearity coefficient of *Post-NUC residual* below is referenced to full scale, so the reference S_{ref} must live in the same summed domain as the signal it is compared against: the counting detector's effective well when that branch ran, otherwise the per-pixel full-well capacity scaled by the accumulation gain the signal actually received through TDI, binning, and co-adding. One domain for numerator and denominator is what makes the residual's signal-relative size invariant in N .

7.3 Post-NUC residual — the correctability parabola

7.3.1 Two-point correction

Model each pixel's response as linear plus a small quadratic term, $y = g(S + \beta S^2/S_{ref}) + o$, with the curvature coefficient β dispersed across the array with 1σ magnitude σ_β . A two-point correction solves for each pixel's gain and offset so that the corrected output is exact at S_1 and S_2 . Carrying the exact linear-correction algebra through in the small- β limit leaves the per-pixel error

$$\Delta S = \beta \frac{(S - S_1)(S - S_2)}{S_{ref}},$$

so the array RMS residual is

$$\sigma_{NUC}(S) = \beta \frac{|(S - S_1)(S - S_2)|}{S_{ref}} \quad [\text{e- RMS}]$$

— a parabola that vanishes at both cal points, peaks between them, and grows quadratically outside them. The peak sits at the midpoint of the cal span and has the value $\beta (S_2 - S_1)^2 / (4S_{ref})$, which is the one number worth memorizing: the residual is quadratic in the *span*, so halving the separation between cal points quarters the worst-case residual inside it. Gain and offset dispersion (PRNU and DSNU) are removed exactly by the correction itself and re-enter only through the drift terms of *Drift between calibration events*.

Symbols. S scene signal [e-]; S_1, S_2 cal-point signals [e-]; β 1σ nonlinearity fraction at full scale [dimensionless]; S_{ref} full-scale reference [e-].

Pitfalls. Quoting the residual at a cal point (it is zero there by construction, which is why *Cal-source spatial non-uniformity* exists); referencing β to the signal instead of to full scale, which makes the residual cubic; taking S_{ref} in the per-pixel domain while S is in the summed TDI domain, which makes the residual appear to shrink with N_{TDI} .

Numeric anchor. $\beta = 0.5\%$, $S_1 = 2.000 \times 10^4$ e-, $S_2 = 8.000 \times 10^4$ e-, $S_{ref} = 1.000 \times 10^5$ e-: at the midpoint $S = 5.000 \times 10^4$ e- the residual is 45.00 e- RMS, and it is 0.00 e- RMS at either cal point.

7.3.2 One-point correction

An offset-only correction leaves the gain dispersion intact, acting on the departure from the single cal point:

$$\sigma_{1pt}(S) = k |S - S_1| \quad [\text{e- RMS}]$$

with k the pre-correction PRNU fraction. The contrast with the two-point correction above is the point of running two points at all: the one-point residual is *linear* in the departure and carries the full PRNU coefficient, so it is typically an order of magnitude larger.

Numeric anchor. $k = 0.5\%$, $S_1 = 2.000 \times 10^4$ e-, $S = 5.000 \times 10^4$ e-: 150.0 e- RMS — 3.3× the two-point residual at the same signal in the anchor above.

7.3.3 Three-point correction

With a third point between the extremes, the correction is piecewise: inside each bracketing segment it is exactly that pair's two-point solve, so the residual is that segment's own parabola, vanishing at all three cal points and peaking at a quarter of the corresponding two-point value when the mid point bisects the span. Outside the calibrated range the nearest segment extrapolates, exactly as the two-point model extrapolates past its own span. The model stops at three points by design; beyond that a real instrument uses a response curve, not more anchors.

Numeric anchor. Same β and S_{ref} as the two-point correction, with a mid point at 5.000×10^4 e-: the residual at 3.500×10^4 e- (the midpoint of the lower segment) is 11.25 e- RMS, a quarter of the 45.00 e- RMS two-point peak.

7.4 Cal-source spatial non-uniformity

A real calibration blackbody holds only ± 0.01 – 0.05 K (1σ) across its aperture. At cal time each pixel therefore views a slightly different source temperature and the correction imprints that pattern into its own coefficients.

With δT the pixel's local departure and $D_j = dS/dT$ evaluated at cal point j [e-/K], the imprinted error at cal view j is $\delta S_j = \delta T D_j$, and propagating it through the correction solve gives

$$\sigma_{unif}(S) = \Delta T_{unif} D_1 \quad (\text{one-point}), \quad \sigma_{unif}(S) = \Delta T_{unif} \frac{|D_1(S_2 - S) + D_2(S - S_1)|}{S_2 - S_1} \quad (\text{two-point})$$

with the three-point form interpolating piecewise between the bracketing pair's constants.

The combination is linear, not RSS. The same plate is viewed at both cal points and a cavity gradient is temperature-independent in kelvin to first order, so the two imprints are fully correlated. Treating them as independent and root-sum-squaring the weights would under-predict the residual inside the cal span.

Why this term matters out of proportion to its size. Unlike the nonlinearity parabola, it does **not** vanish at the cal points: $\sigma_{unif}(S_1) = \Delta T_{unif} D_1$ exactly. It is what retires the comfortable but false picture of a calibration that is perfect where it was taken. When it dominates, it sets a calibration-limited NEDT floor of approximately ΔT_{unif} itself, because the term and the scene derivative it is divided by are both proportional to dS/dT .

Numeric anchor. MWIR 3.0–5.0 μm , scene 300.0 K, scene signal 5.000×10^4 e-, cal points 280.0 K and 320.0 K: $D_1 = 950.3$ e-/K and $D_2 = 3123.9$ e-/K. With $\Delta T_{unif} = 0.020$ K the residual at the scene signal is 34.51 e- RMS. Against the scene's own derivative $dS/dT = 1795$ e-/K that is 0.0192 K — within 4 % of ΔT_{unif} , as the floor argument predicts.

7.5 Drift between calibration events

Between the cal event and the observation, the corrected gain and offset both walk. The v1 model is time-linear in both, with 1σ rates:

$$\sigma_{gain}(S) = r_g t_{cal} S \quad [\text{e- RMS}], \quad \sigma_{offset} = r_o t_{cal} \quad [\text{e- RMS}]$$

with r_g in s^{-1} (entered as %/hour), r_o in e-/s (entered as e-/hour), and t_{cal} the elapsed time since calibration [s]. Gain drift is proportional to signal and therefore scales with scene brightness; offset drift is signal-independent and dominates at low flux. Both are spatial and calibration-correlated, so both are exempt from $\sqrt{N}$ averaging.

The practical use of these two terms is cadence: they are what turns “how often must we calibrate” into a number, since they are the only residual terms that grow without bound between events.

Numeric anchor. $r_g = 0.20\%$ /hour and $r_o = 10.0$ e-/hour, 1800 s (30 min) since the cal event, scene signal 5.000×10^4 e-: $\sigma_{gain} = 50.00$ e- RMS and $\sigma_{offset} = 5.000$ e- RMS.

Deferred. Drift driven by focal-plane temperature change is not modeled; the v1 rates are time-linear only.

7.6 The internal-shutter cal path

A calibration flag inside the optical train sees only itself and the elements *behind* it. The correction derived from that view absorbs the aft-element emission but never the emission of the elements in *front* of the flag, which returns in operation as an uncorrected offset. The split is the photon-weighted in-band fraction of the near-field focal-plane irradiance contributed by the fore elements:

$$f_{fore} = \frac{\int_{\lambda_1}^{\lambda_2} \sum_{i \in \text{fore}} E_i(\lambda) \lambda d\lambda}{\int_{\lambda_1}^{\lambda_2} \sum_{i \in \text{all}} E_i(\lambda) \lambda d\lambda} \quad [\text{dimensionless}]$$

with $E_i(\lambda)$ the per-element near-field irradiance at the focal plane [W/m²/μm]. Each watt of in-band power carries λ/hc photons, so hc cancels in the ratio and only the factor λ survives as the weighting. The electron magnitude then rides the detector's own near-field conversion, $S_{fore} = S_{nf} f_{fore}$ [e-]: no radiometry is re-derived here, which is what keeps this term consistent with the warm-optics treatment of Chapter 3 and the appendix.

That uncorrected offset produces two distinct effects. Its mean is an **offset bias**, S_{fore}/S as a fraction of the scene signal, which belongs in the accuracy budget of *The bias budget* below. The part of it that varies across the array — cold-stop reflections, vignetting of the warm fore-optics — is **narcissus fixed-pattern noise**, a spatial noise term equal to a declared fraction of S_{fore} .

Assumption. The split weights per-element irradiance by photon count rather than by the QE spectrum. It is exact for flat QE across the band and first-order otherwise, because the per-element spectra are same-family graybody curves and the QE weighting largely cancels in the ratio.

7.7 Composition — how calibration enters the budget

The five residual terms are mutually independent mechanisms, so they combine in quadrature, and the calibration RSS combines the same way with the post-readout total:

$$\sigma_{cal} = \sqrt{\sigma_{NUC}^2 + \sigma_{unif}^2 + \sigma_{gain}^2 + \sigma_{offset}^2 + \sigma_{narc}^2}, \quad \sigma_{tot} = \sigma_{readout} \oplus \sigma_{cal}$$

Every term is appended in the electron domain, input-referred at the sense node, exactly as Chapter 6 requires of any noise term. All five are classed **spatial**, so a single-frame spatial budget contains them; none is temporal, so none averages down over frames. The downstream metrics prefer this post-calibration total wherever it is published: SNR and contrast SNR use σ_{tot} above, and the signal-to-clutter-plus-noise ratio adds σ_{cal} into its spatial RSS.

Worked example (MWIR staring, two-point cal). Band 3.0–5.0 μm, scene 300.0 K, scene signal 5.000×10^4 e-, full-scale reference 1.000×10^5 e-, cal points 280.0 K and 320.0 K (so $S_1 = 2.3238 \times 10^4$ e-, $S_2 = 9.8262 \times 10^4$ e-), nonlinearity $\beta = 0.5\%$, source uniformity $\Delta T_{unif} = 0.020$ K, drift rates $r_g = 0.20\%$ /hour and $r_o = 10.0$ e-/hour at $t_{cal} = 1800$ s, full-aperture cal path:

Term	Value	Class
Post-NUC residual	64.58 e- RMS	spatial, cal-correlated
Cal-source uniformity	34.51 e- RMS	spatial, cal-correlated
Gain drift	50.00 e- RMS	spatial, cal-correlated
Offset drift	5.000 e- RMS	spatial, cal-correlated
Calibration RSS σ_{cal}	88.81 e- RMS	—

Against a temporal (post-readout) total of 300.0 e- RMS this gives $\sigma_{tot} = 312.9$ e- RMS: the SNR falls from 166.7 to 159.8, a 4.1 % loss that no amount of integration time recovers.

Rule-4 non-interaction. None of these terms has a PSF kernel or an MTF factor. Adding an “FPN MTF” would double-count a spatial *noise* as a spatial *degradation*.

7.8 Calibration-limited NEDT

The reason the stage exists. Dividing the calibration RSS by the scene’s own thermal derivative gives the temperature resolution the calibration alone permits:

$$\text{NEDT}_{cal} = \frac{\sigma_{cal}}{dS/dT} \quad [\text{K}]$$

with dS/dT [e-/K] the band-integrated derivative the chain already computes for NEDT. This is a *floor*, not a contribution to be traded: because σ_{cal} is exempt from $\sqrt{N}$ averaging, lengthening the integration or adding TDI stages raises the signal and the residual together, leaving NEDT_{cal} where it was. A system whose temporal NEDT has been pushed below its calibration floor has bought nothing.

Numeric anchor. Continuing the *Composition* example, $dS/dT = 1795$ e-/K at the 300.0 K scene:

Quantity	Value
Temporal NEDT (readout total alone)	167.1 mK
Calibration-limited NEDT	49.5 mK
Achieved NEDT (post-calibration total)	174.3 mK

The calibration floor costs only 7.2 mK on this 167.1 mK system. Halve the temporal noise and the *same* floor costs 13.5 mK on an 83.6 mK system; halve it again and it costs 23.0 mK on a 41.8 mK system, which now achieves 64.8 mK. The floor does not move, so every improvement bought with integration time, cooling, or TDI stages is worth progressively less until the calibration — not the detector — is the thing to fix. That crossover, not the absolute numbers, is the design content of this chapter.

7.9 The bias budget — accuracy, not precision

A calibration-source error does not disperse pixel to pixel; it moves the whole array's radiometric scale. Such an error is reported as a fractional radiance bias and never appears in σ_{tot} .

Cal-source temperature. A 1σ uncertainty ΔT_{src} in the cal source's own temperature propagates through the photon-weighted band Planck log-derivative at the cal temperature:

$$b_{src} = \frac{\int_{\lambda_1}^{\lambda_2} (\partial B_q / \partial T) d\lambda}{\int_{\lambda_1}^{\lambda_2} B_q d\lambda} \Delta T_{src} \quad [\text{dimensionless}]$$

Cal-source emissivity. A pure scale error, $b_\varepsilon = \Delta\varepsilon / \varepsilon_{src}$.

Spectral calibration. A band-center error $\delta\lambda$ [μm] rigidly shifts the band the radiometry is computed over. The calibration absorbs the resulting scale error *at its own source temperature*, because the gain is derived there, so what survives on a scene is the difference of the band-shift log-derivatives:

$$g(T) = \frac{B_q(\lambda_2, T) - B_q(\lambda_1, T)}{\int_{\lambda_1}^{\lambda_2} B_q(\lambda, T) d\lambda} [\mu\text{m}^{-1}], \quad b_{spec} = |g(T_{scene}) - g(T_{cal})| \delta\lambda$$

The expression follows from the Leibniz rule with both band edges moving together. It is exactly zero when the scene sits at the cal temperature and grows with the scene-versus-cal separation; short-wave (Wien-side) bands are the sensitive ones. Band-*width* drift is a distinct and typically smaller mechanism and is not modeled.

Internal-cal offset. The uncorrected fore-optics emission of *The internal-shutter cal path* as a fraction of the scene signal, S_{fore}/S .

Absolute gain. A directly declared fractional uncertainty in the radiometric gain.

Composition. Independent sources RSS *within* the budget, and the total converts to a temperature at the scene through the chain's own derivative:

$$b_{tot} = \sqrt{\sum_i b_i^2}, \quad \Delta T = \frac{b_{tot} S}{dS/dT} \quad [\text{K}]$$

A scene with no thermal derivative reports the fractional value with a named failure for the kelvin conversion rather than a silent NaN.

Numeric anchors. $\Delta T_{src} = 0.10$ K on a 300.0 K cal source gives $b_{src} = 3.590 \times 10^{-3}$ (0.359 %) in a 3.0–5.0 μm band and 1.615×10^{-3} (0.161 %) in an 8.0–12.0 μm band — the MWIR band's steeper Planck slope makes it 2.2 $\times$ as unforgiving of source-temperature error. An emissivity uncertainty of 0.002 on an $\varepsilon_{src} = 0.99$ cavity gives $b_\varepsilon = 2.020 \times 10^{-3}$. A band-center uncertainty $\delta\lambda = 0.005$ μm with the scene at 300.0 K and the cal at 320.0 K in the 3.0–5.0 μm band gives $b_{spec} = 6.121 \times 10^{-4}$ (0.061 %), from $g(300 \text{ K}) = 1.5459 \mu\text{m}^{-1}$ and $g(320 \text{ K}) = 1.4235 \mu\text{m}^{-1}$.

Pitfalls. RSS-ing a bias into the noise total, which makes an accuracy error look like something $\sqrt{N}$ can fix; quoting the spectral-calibration bias at the cal temperature (it is zero there); reporting a bias in kelvin without stating the scene temperature it was converted at.

7.10 Assumptions, validity, and what is deferred

Assumption	Why it holds	What breaks it
Quadratic response over the cal span	The departure from linearity of a well-behaved photodiode is smooth and small over the calibrated range	Approach to saturation, where the response is not polynomial and the residual is a lower bound
Band photon-radiance ratio for cal points	Only the ratio enters, anchored at nearby temperatures, so the response slope largely cancels	Strongly sloped $QE\tau$ across the band; a scene whose signal is not thermal at the declared temperature (see <i>When the thermal anchor does not describe the scene</i>)
Fully correlated source-uniformity imprints	The same plate is viewed at every cal point and a cavity gradient is temperature-independent in K to first order	A source whose gradient changes shape with temperature
Time-linear drift	Adequate over a cal interval short compared with the thermal time constants driving it	Long intervals, or drift driven by focal-plane temperature excursions (not modeled)
Full PRNU correlation along a TDI column	Takes the correlated limit, an upper bound on the residual floor	Long TDI columns with partially decorrelated pixel gains — the true residual is then lower than modeled
Photon-weighted fore-optics split	Exact for flat QE; the per-element spectra are same-family graybody curves	Sharply structured QE across the band

Deferred by ruling. Focal-plane-temperature-driven drift; operations-level calibration (cal cadence trades, scene-based non-uniformity correction); partial along-column decorrelation of PRNU under long TDI. Each has a home in this stage when it arrives.

Boundary conditions. With the scheme off, the stage emits nothing and results are bit-identical to a chain without it — enabling the model is a deliberate act, never a default that quietly moves numbers. Cal points that coincide, are unordered, or are left unset with a scheme active are rejected with an actionable error rather than ill-conditioned arithmetic: as $S_2 \rightarrow S_1$ the two-point solve degenerates, and the model refuses instead of dividing by a vanishing span.

7.11 Parameter cross-reference

Volume III's Parameter Reference lists the `calibration.*` namespace with types, defaults, bounds, and entry units. Every cal-related default is the model-off limit.

Performance Metrics

The metric layer: how RADIANT turns the signal/noise/spatial chain state into SNR, NEDT, NIIRS, detection range, and their relatives. Numeric anchors were re-derived from the literature independently of the implementation. This is the one layer permitted to report a *named failure* instead of a number — a metric that cannot be computed for a scene says so, with a reason the caller can read; silently returning NaN remains forbidden anywhere in RADIANT.

Symbols: S target signal [e-]; S_{bg} background signal [e-]; σ_{tot} total noise [e- RMS] (see the Noise Model chapter); t_{int} [s]; p pixel pitch [m]; GSD [m].

8.1 SNR family

Equations.

$$\text{SNR} = \frac{S}{\sigma_{tot}}, \quad \text{CSNR} = \frac{S - S_{bg}}{\sigma_{tot}}, \quad \text{SCNR} = \frac{S - S_{bg}}{\sigma_{tot} \oplus \sigma_{clutter}}$$

($\oplus$ = RSS). **The denominator always contains the background and dark shot noise:** subtracting the background estimate in the numerator does not remove its Poisson fluctuations. In the thermal IR, where $\Delta S \ll S_{bg}$, using $\sqrt{S - S_{bg}}$ in the denominator overstates contrast SNR by large factors (1.9× at the audit anchor).

Numeric anchor. $S = 10000$, $S_{bg} = 5000$, $N_{dark} = 1000$ e-, $\sigma_{read} = 50$ e-, PRNU 0.1 %: $\sigma_{tot} = 136.38$ e-, SNR = 73.32, CSNR = 36.66.

8.2 NEDT

Equation.

$$\text{NEDT} = \frac{\sigma_{tot}}{dS/dT} \quad [\text{K}],$$

with the thermal responsivity band-integrated in the photon domain:

$$\frac{dS}{dT} = t_{int} G \int \frac{\lambda_m}{hc} QE(\lambda) \tau(\lambda) \frac{\partial L_\lambda}{\partial T} d\lambda \quad [\text{e}^-/\text{K}],$$

$\partial L_\lambda / \partial T$ the analytic Planck derivative (the Radiometric Chain chapter), G the étendue. RADIANT computes dS/dT in `SpectralIntegrationStage` (`ds_dt_e_per_K` stage output) and divides in the metric layer; a fallback `compute_nedt_from_snr` uses the SNR route when the derivative output is absent.

Pitfalls. Paraxial étendue $A_d \pi / (4F_\#^2)$ vs exact $A_d \pi / (4F_\#^2 + 1)$ — 6.25 % at $f/2$, 25 % at $f/1$ (understates NEDT); finite-difference $\Delta T = 1$ K instead of the analytic derivative; energy-domain radiance without the λ_m/hc photon factor; quoting mK NEDT that implies an impossible well fill (the audit's LWIR example: an $f/2$, 1 ms configuration implies $800\times$ a 100 ke- well — always sanity-check S_{bg} against `total_well_e`).

Numeric anchors. $dS/dT = 5000$ e-/K, $\sigma = 60$ e- $\rightarrow$ NEDT = 12.0 mK; $\int_8^{12} (\partial L_\lambda / \partial T) d\lambda = 0.6297$ W/m²/sr/K at 300 K.

8.3 NEI, NEP, and D*

Equations. Noise-equivalent irradiance at the aperture (photon domain):

$$\text{NEI} = \frac{\sigma_{tot}}{\eta_{sys} EE A_{ap} t_{int}} \quad [\text{ph/s/cm}^2 \text{ or /m}^2],$$

power form via hc/λ_m . Specific detectivity:

$$D^* = \frac{\sqrt{A_d \Delta f}}{\text{NEP}} \quad [\text{cm}\sqrt{\text{Hz/W}}], \quad \Delta f = \frac{1}{2 t_{int}},$$

with A_d in cm² (Jones convention). BLIP-limit PV detectivity $(\lambda_m/hc)\sqrt{\eta/(2Q_b)}$; photoconductors are $\sqrt{2}$ lower (generation and recombination noise).

Symbols in this section. η is the detector **quantum efficiency** [dimensionless] and η_{sys} the end-to-end **photon-to-electron efficiency** (optical throughput $\times$ QE), following the detector literature these expressions come from. This is the one place in the manual where η is not the look angle of the geometry chapter (notation §9); the two never appear in one equation. Q_b is the **background photon flux at the detector** [photon/s/cm²] that sets the BLIP limit — the same in-band background the noise chapter integrates, expressed as a photon irradiance on A_d rather than as electrons per frame.

Pitfalls. $\Delta f = 1/t_{int}$ vs $1/(2t_{int})$ — a $\sqrt{2}$ error of exactly the PV/PC magnitude, easily conflated; A_d in m² (inflates $D^* \times 100$); omitting the ensquared-energy factor in NEI (optimistic by $1/EE$); photon/energy domain mixing (the λ_m/hc lost or doubled — the NEP $\leftrightarrow$ electrons converter carries an absolute anchor for exactly this).

Numeric anchors. $\text{NEP} = 10^{-14}$ W, $A_d = (20 \text{ }\mu\text{m})^2$, $\Delta f = 1000$ Hz $\rightarrow D^* = 6.3246 \times 10^{12}$ cm $\sqrt{\text{Hz/W}}$; NEP anchor $\sigma = 100$ e-, $\eta = 0.7$, $\lambda = 10 \text{ }\mu\text{m}$, $t = 10$ ms $\rightarrow 2.8378 \times 10^{-16}$ W.

8.4 NIIRS via GIQE-5 (and the IIRS placeholder)

Equation.

$$\text{NIIRS} = c_0 + c_1 \log_{10} \text{GSD}_{\text{in}} + c_2 \log_{10} \text{RER} + c_3 \log_{10} \text{SNR} + c_4 H + c_5 G$$

with the literature coefficients (9.57, −3.32, 3.32, 1.559, −0.334, −0.01) pinned exactly by test; GSD in **inches** (geometric mean of the two ground-sample directions the geometry chapter derives), RER the geometric mean from the PSF path, H edge overshoot, G noise gain.

Fit-envelope gating. GIQE-5 was fit over GSD 1.18–31.5 in, RER 0.2–0.95, SNR 2–130. Outside the envelope RADIANT refuses by default (result-typed failure with `failure_reason`), computes only when `performance.niirs.allow_extrapolated = true`, and flags `niirs_extrapolated = 1`.

IIRS status. The MWIR/LWIR interpretability metric currently **reuses GIQE-5 verbatim** — same formula, envelope, and coefficients. Treat IR “IIRS” outputs as GIQE-5-on-IR until a real IIRS lands.

Numeric anchor. GSD = 1 m, RER = 0.9, SNR = 50 → NIIRS = 6.4268 (pure-literal test value).

8.5 Johnson criteria and minimum resolvable

Equations. Cycles resolved across a target’s critical dimension d_c at range R_s :

$$N_{cyc} = \frac{d_c}{2 \text{ IFOV } R_s}$$

(the factor 2 converting pixels to cycles), compared against the Johnson thresholds (detect ≈ 1 , recognize ≈ 4 , identify ≈ 8 cycles, 50 % probability). `johnson_range_m` inverts for range at a given task. Minimum-resolvable temperature/contrast couple the MTF budget to NEDT/SNR thresholds.

Pitfalls. Cycles vs pixels (factor 2); using nadir IFOV footprint at slant geometry; 50%-probability thresholds quoted as certainties.

8.6 Detection range

Equation. For a point source with atmospheric extinction, the signal-vs-range equation mixes $1/R_s^2$ geometric falloff with transmission along the path, $S(R) = S_{ref}(R_{ref}/R)^2 \tau(R)/\tau(R_{ref})$. The noise is **not** frozen at the reference range: the target’s own shot variance in electrons is the signal itself, so

$$\sigma^2(R) = S(R) + N_0^2, \quad N_0^2 = \sigma_{ref}^2 - S_{ref}$$

with N_0 the target-free floor (background shot, dark, read, kTC, quantization). RADIANT solves

$$\frac{S(R)}{\sqrt{S(R) + N_0^2}} = \text{SNR}_{threshold}$$

by bisection (point-source regime only; threshold parameter default 5.0). The criterion inverts in closed form — the signal a threshold T demands is $S^* = \frac{1}{2} \left(T^2 + \sqrt{T^4 + 4T^2 N_0^2} \right)$, so in vacuum the answer is $R = R_{ref} \sqrt{S_{ref}/S^*}$ with no root finding at all. Both forms agree at R_{ref} by construction, so the answer no longer depends on the range the chain was evaluated at. The superseded frozen-noise form solved $S(R)/\sigma_{ref} = T$ instead, and gave 123.4 km referenced at 25 km against 182.5 km referenced at 100 km for one unchanged configuration — a range that depended on where you asked from.

Pitfalls. Applying extended-scene SNR to the point-source range equation; forgetting that EE_box and jitter enter through the signal chain, not as post-hoc factors; α from a band-average τ over a very different path length; passing a signal and a total noise from different operating points, which leaves $\sigma^2 - S$ negative and is refused.

8.7 Saturation, dynamic range, BLIP

Equations. Well margin $20 \log_{10}(\text{FWC}/S_{tot})$ dB and ADC margin $20 \log_{10}(DN_{max}/S_{DN})$ dB; dynamic range

$$\text{DR} = 20 \log_{10} \frac{\text{FWC}}{\sigma_{floor}} \text{ dB}$$

with σ_{floor} the **dark-scene** floor (read + dark shot + ADC), not the shot-inflated full-well noise. BLIP: background shot exceeds all other noise RSS'd; $f_{BLIP} = \sigma_{bg,shot}/\sigma_{tot}$ ($\rightarrow 1$ fully BLIP; threshold $1/\sqrt{2}$); `blip_rate_e_per_s` reports the photon rate scale.

Pitfalls. 10 log vs 20 log (halves the dB); ADC full scale vs true full well (take the min); counting target shot as “background” in the BLIP test; quoting BLIP at one t_{int} as if general.

Numeric anchors. 100 ke-/50 e- $\rightarrow$ 66.02 dB; the audit’s marginal-BLIP case: $\sigma_{bg} = 70.7$ vs other-RSS 60.0 e-, $f_{BLIP} = 0.762$.

8.8 Radiometric inversions and sensitivity relatives

NE Δ L / NE Δ p — noise-equivalent radiance and reflectance differences, computed as radiance/SNR and reflectance/SNR. They are available as helpers but are **not part of the chain’s metric surface**, so a run does not report them. **Temperature retrieval** — band-radiance inversion with emissivity and temperature Jacobians; the forward band-integrated Planck radiance it inverts is anchored at 38.5 W/m²/sr for a 300 K scene over 8–12 μm .

8.9 Metric selection

Metrics are computed in five groups — radiometric, sampling/geometry, spatial-MTF, saturation, interpretability — each of which can be switched off. All five are on by default, and switching one off stops the computation rather than hiding the result: a metric another enabled group depends on is still computed, because the selection closes over dependencies.

Deselecting the whole Spatial-MTF group is the one case that saves real time. When nothing else needs a spatial input, the PSF convolution stack and the MTF product are skipped entirely, and with them the consistency check that normally compares the two spatial paths on every run — there is no spatial computation left to check.

Radiometric Model — Mixed Refractive and Reflective Optical Train

This appendix is an implementation recipe, and it is written like one. The chapters before it state each equation with its derivation, its validity limits and a pinned numeric anchor; this appendix instead gives the quantities in the order a program must compute them, in the ASCII notation the implementation uses, so that a reader working through the code can follow it line for line. There are no derivations or pitfalls sections here, and the equation blocks are deliberately not typeset as display mathematics — they are the recipe, not the physics. The physics they assemble is in the radiometric-chain and spatial chapters.

Two separate radiometric paths are modeled:

1. **Signal path** — transmission/reflection of source photons through the full optical train
2. **Thermal background path** — self-emission of each optical element propagated to the focal plane

Elements are typed as either REFRACTIVE or REFLECTIVE. The per-element equations differ fundamentally between the two types.

A.1 Element types and their radiometric character

A.1.1 Refractive element (lens, window, filter, beamsplitter substrate)

- Signal passes **through** the element
- Has bulk absorption (Beer-Lambert), entry and exit surface coatings
- System quantity is **transmittance** $T_{sys,i}(\lambda)$
- Thermal emission sourced from **bulk substrate + both surfaces**
- Cavity etalon effect between R1 and R2

A.1.2 Reflective element (mirror, cold stop, baffle)

- Signal **reflects off** the element
- No bulk propagation — surface interaction only
- System quantity is **reflectance** $Rho_{sys,i}(\lambda)$
- Thermal emission sourced from **surface coating only**
- No etalon — single surface interaction
- Emissivity governed by Kirchhoff: $eps_i(\lambda) = 1 - Rho_{sys,i}(\lambda) - A_{coat,i}(\lambda)$ where $A_{coat,i}$ is the coating absorption (if known); if not: $eps_i = 1 - Rho_{sys,i}$

A.2 Per-element radiometric quantities

A.2.1 Refractive element i

$$\begin{aligned}\text{beer}_i(\lambda) &= \exp(-\alpha_i(\lambda) \cdot d_i / \cos(\theta_{r,i})) \\ \text{denom}_i(\lambda) &= 1 - R1_i(\lambda) \cdot R2_i(\lambda) \cdot \text{beer}_i(\lambda)^2\end{aligned}$$

$$T_{\text{sys},i}(\lambda) = \frac{T1_i(\lambda) \cdot \text{beer}_i(\lambda) \cdot T2_i(\lambda)}{1 - R1_i(\lambda) \cdot R2_i(\lambda) \cdot \text{beer}_i(\lambda)^2}$$

$$\text{eps}_{\text{eff},i}(\lambda) = \frac{T2_i(\lambda) \cdot n_i(\lambda)^2 \cdot (1 - \text{beer}_i(\lambda))}{1 - R1_i(\lambda) \cdot R2_i(\lambda) \cdot \text{beer}_i(\lambda)^2}$$

$$L_{\text{thermal},i}(\lambda) = \text{eps}_{\text{eff},i}(\lambda) \cdot B(\lambda, T_i)$$

The element **transfer factor** applied to upstream signals passing through it:

$$C_i(\lambda) = T_{\text{sys},i}(\lambda) \quad [\text{refractive} - \text{transmission}]$$

A.2.2 Reflective element i

No bulk absorption. Single surface interaction. The coating has reflectance $Rho_i(\lambda)$ and absorptance $A_{\text{coat},i}(\lambda)$. Transmittance is zero (opaque mirror).

$$Rho_{\text{sys},i}(\lambda) = Rho_i(\lambda) \quad [\text{mirror reflectance}]$$

$$\text{eps}_i(\lambda) = 1 - Rho_i(\lambda) - A_{\text{coat},i}(\lambda) \quad [\text{surface emissivity via Kirchhoff}]$$

If coating absorptance is not separately characterized, conservative form:

$$\text{eps}_i(\lambda) = 1 - Rho_i(\lambda) \quad [\text{upper bound on emissivity}]$$

Thermal radiance from the mirror surface:

$$L_{\text{thermal},i}(\lambda) = \text{eps}_i(\lambda) \cdot B(\lambda, T_i)$$

The element **transfer factor** applied to upstream signals reflecting off it:

$$C_i(\lambda) = Rho_{\text{sys},i}(\lambda) \quad [\text{reflective} - \text{reflectance}]$$

A.2.3 Planck blackbody radiance (common to both types)

$$B(\lambda, T_i) = \frac{2 h c^2 / \lambda^5}{\exp(h c / \lambda k_B T_i) - 1}$$

Constants are the CODATA 2018 values of the Notation chapter's physical-constants table; they are defined once in the code and never re-entered here.

B is a **per-wavelength** spectral radiance in $\text{W}/\text{m}^2/\text{sr}/\mu\text{m}$. Evaluate the expression above with λ in meters and it returns $\text{W}/\text{m}^2/\text{sr}/\text{m}$; the per- μm form used everywhere in this manual carries the Jacobian factor 10^{-6} that converts the spectral density from per-meter to per-micrometer. Mixing the two is the single most common radiometry error.

A.3 Signal path through the full optical train

A.3.1 Cascaded system transfer

Each element applies its transfer factor $C_i(\lambda)$ — either $T_{sys,i}$ or $Rho_{sys,i}$ depending on type. The total signal transfer from source to focal plane is:

$$C_{total}(\lambda) = \prod C_i(\lambda) \quad \text{for } i = 1 \text{ to } N$$

Explicitly for a mixed train of N elements:

$$C_{total}(\lambda) = C_1(\lambda) \cdot C_2(\lambda) \cdot C_3(\lambda) \cdot \dots \cdot C_N(\lambda)$$

Where each C_i is either $T_{sys,i}$ (refractive) or $Rho_{sys,i}$ (reflective).

A.3.2 Signal irradiance at the focal plane

$$E_{signal,FP}(\lambda) = \pi \cdot L_{source}(\lambda) \cdot \sin^2(\theta_{FP}) \cdot C_{total}(\lambda)$$

Or equivalently using f-number:

$$E_{signal,FP}(\lambda) = \frac{\pi \cdot L_{source}(\lambda) \cdot C_{total}(\lambda)}{4 \cdot (f/\#)^2 \cdot (1 + M)^2}$$

A.3.3 Total in-band signal irradiance

$$E_{signal,total} = \int [\lambda_1 \text{ to } \lambda_2] E_{signal,FP}(\lambda) d\lambda$$

A.4 Thermal background path from each element

Each element i emits thermal radiation that propagates through all downstream elements $j > i$ before reaching the focal plane.

A.4.1 Downstream cumulative transfer from element i to focal plane

The thermal emission from element i is transferred by every downstream element using that element's transfer factor $C_j(\lambda)$:

$$\tau_{i,j}(\lambda) = \prod C_j(\lambda) \quad \text{for } j = i+1 \text{ to } N$$

Note: C_j is $T_{sys,j}$ if element j is refractive, $Rho_{sys,j}$ if reflective. Downstream reflective elements both redirect **and** attenuate the upstream thermal flux.

For the last element: $\tau_{i,N}(\lambda) = 1.0$

Explicitly for a 4-element train:

$$\begin{aligned}\tau_1(\lambda) &= C_2 \cdot C_3 \cdot C_4 \\ \tau_2(\lambda) &= C_3 \cdot C_4 \\ \tau_3(\lambda) &= C_4 \\ \tau_4(\lambda) &= 1.0\end{aligned}$$

A.4.2 The acceptance cone — one geometry for every element

The Lagrange invariant caps how much solid angle the focal plane can accept from an in-beam element. Whatever the internal layout, each element is seen through the reimaging optics and can fill at most the cone the working f-number sets:

$$\Omega_{\text{cone}} = 2 \cdot \pi \cdot (1 - \cos(\theta)), \quad \theta = \arctan(1 / (2 \cdot N_{\text{eff}}))$$

Where:

- N_{eff} = effective (post-cold-stop) f-number, f / D_{eff}
- θ = half-angle of the marginal ray

The exact form is used rather than the paraxial $\pi / (4 \cdot N_{\text{eff}}^2)$: the two agree to 0.5 % at $f/6$ and 5 % at $f/2$, and the paraxial form over-states the cone, exceeding 2π as $N_{\text{eff}} \rightarrow 0$, which no solid angle may do.

There is **no per-element geometry**. An element does not gain acceptance by sitting close to the focal plane: a private $\Omega_i = \pi (D_i/2)^2 / z_i^2$ is not bounded by the invariant — a 0.3 m mirror 1.0 m from the focal plane would claim 0.0707 sr against an $f/6$ cone of 0.0217 sr, 3.2× more than physics permits. Nor is there a cold-stop attenuation factor: in-cone emission arrives through the imaging path itself and cannot be blocked, while out-of-cone warm structure is taken to be blocked completely. What a cold stop does control is the size of the pupil, hence N_{eff} , hence Ω_{cone} .

A.4.3 Focal plane irradiance from each element

Every element, last or upstream, emits as a graybody into that one cone, attenuated by everything downstream of it:

$$E_{\text{FP},i}(\lambda) = \Omega_{\text{cone}} \cdot L_{\text{thermal},i}(\lambda) \cdot \tau_i(\lambda)$$

Fully expanded:

Refractive element i:

$$E_{\text{FP},i}(\lambda) = \Omega_{\text{cone}} \cdot \epsilon_{\text{eff},i}(\lambda) \cdot B(\lambda, T_i) \cdot \prod_{j=i+1 \text{ to } N} C_j(\lambda)$$

Reflective element i:

$$E_{\text{FP},i}(\lambda) = \Omega_{\text{cone}} \cdot \epsilon_i(\lambda) \cdot B(\lambda, T_i) \cdot \prod_{j=i+1 \text{ to } N} C_j(\lambda)$$

The structure is identical — only the emissivity model differs between the two types. The last element is not a special case: $\tau_N = 1$, so it contributes $\Omega_{\text{cone}} \cdot L_{\text{thermal},N}(\lambda)$.

Dimensionally, $\text{sr} \times [-] \times \text{W/m}^2/\text{sr}/\mu\text{m} \times [-] = \text{W/m}^2/\mu\text{m}$, the spectral irradiance the focal plane sees.

A.4.4 Total thermal background irradiance at the focal plane

$$E_{\text{background},\text{FP}}(\lambda) = \sum_{i=1 \text{ to } N} E_{\text{FP},i}(\lambda)$$

Total in-band thermal background:

$$E_{\text{background},\text{total}} = \int_{[\lambda_1 \text{ to } \lambda_2]} E_{\text{background},\text{FP}}(\lambda) d\lambda$$

A.5 Combined focal plane irradiance

$$E_{\text{FP},\text{total}}(\lambda) = E_{\text{signal},\text{FP}}(\lambda) + E_{\text{background},\text{FP}}(\lambda)$$

Signal-to-background ratio:

$$\text{SBR}(\lambda) = E_{\text{signal},\text{FP}}(\lambda) / E_{\text{background},\text{FP}}(\lambda)$$

A.6 Kirchhoff self-consistency check

A.6.1 Refractive element i

System reflectance (accounting for cavity):

$$R_{\text{sys},i}(\lambda) = \frac{R1_i + T1_i^2 \cdot R2_i \cdot \text{beer}_i^2}{1 - R1_i \cdot R2_i \cdot \text{beer}_i^2}$$

Check:

$$\begin{aligned} A_{\text{total},i}(\lambda) &= 1 - R_{\text{sys},i}(\lambda) - T_{\text{sys},i}(\lambda) \\ \text{eps}_{\text{eff},i}(\lambda) &\approx A_{\text{total},i}(\lambda) \end{aligned} \quad \leftarrow \text{must hold at each } \lambda$$

A.6.2 Reflective element i

$$\begin{aligned} A_{\text{total},i}(\lambda) &= 1 - \text{Rho}_{\text{sys},i}(\lambda) - A_{\text{coat},i}(\lambda) && [\text{if } A_{\text{coat}} \text{ known}] \\ &= 1 - \text{Rho}_{\text{sys},i}(\lambda) && [\text{conservative bound}] \\ \text{eps}_i(\lambda) &\approx A_{\text{total},i}(\lambda) && \leftarrow \text{must hold at each } \lambda \end{aligned}$$

A.7 Implementation recipe

Inputs per element i:

type_i = REFRACTIVE or REFLECTIVE

If REFRACTIVE:

R1_i(λ), T1_i(λ)	entry surface coating
R2_i(λ), T2_i(λ)	exit surface coating
alpha_i(λ)	bulk absorption coefficient

$n_i(\lambda)$ refractive index
 d_i substrate thickness
 T_i temperature
 $\theta_{r,i}$ refracted angle inside substrate

If REFLECTIVE:

$\rho_i(\lambda)$ surface reflectance
 $A_{coat,i}(\lambda)$ coating absorptance (optional; 0 if unknown)
 T_i temperature

System geometry (all elements – one cone, not one per element):

θ_{FP} focal plane convergence half-angle
 N_{eff} effective (post-cold-stop) f-number
 Ω_{cone} $2 \cdot \pi \cdot (1 - \cos(\arctan(1 / (2 \cdot N_{eff}))))$ [sr]

Per element, compute transfer factor C_i and emissivity:

If REFRACTIVE:

$beer_i = \exp(-\alpha_i \cdot d_i / \cos(\theta_{r,i}))$
 $denom_i = 1 - R1_i \cdot R2_i \cdot beer_i^2$
 $C_i = T1_i \cdot beer_i \cdot T2_i / denom_i$ [transmittance]
 $eps_i = T2_i \cdot n_i^2 \cdot (1 - beer_i) / denom_i$ [cavity emissivity]

If REFLECTIVE:

$C_i = \rho_i$ [reflectance]
 $eps_i = 1 - \rho_i - A_{coat,i}$ [surface emissivity]

Both types:

$L_i(\lambda) = eps_i(\lambda) \cdot B(\lambda, T_i)$

Signal path:

$C_{total}(\lambda) = \prod C_i(\lambda)$ for $i = 1$ to N
 $E_{signal}(\lambda) = \pi \cdot L_{source}(\lambda) \cdot \sin^2(\theta_{FP}) \cdot C_{total}(\lambda)$

Thermal background path (build tau as reverse cumulative product):

$\tau_N = 1.0$
 $\tau_i = \prod C_j(\lambda)$ for $j = i+1$ to N [reverse cumulative product]

For $i = 1$ to N :

$E_{FP,i}(\lambda) = \Omega_{cone} \cdot L_i(\lambda) \cdot \tau_i(\lambda)$

$E_{background}(\lambda) = \sum_{i=1}^N E_{FP,i}(\lambda)$

Total:

$E_{FP,total}(\lambda) = E_{signal}(\lambda) + E_{background}(\lambda)$

$$\text{SBR}(\lambda) = E_{\text{signal}}(\lambda) / E_{\text{background}}(\lambda)$$

Validation:

Refractive: $\text{eps_eff}, i \approx 1 - R_{\text{sys}, i} - T_{\text{sys}, i}$ for each optic

Reflective: $\text{eps}_i \approx 1 - \text{Rho}_i - A_{\text{coat}, i}$ for each mirror

A.8 Key physical distinctions summary

Property	Refractive element	Reflective element
Signal transfer factor	$T_{\text{sys}, i}$ (transmittance)	$\text{Rho}_{\text{sys}, i}$ (reflectance)
Bulk emission	Yes — Beer-Lambert path	No
Surface emission	Both surfaces (via T_2, n^2)	Single surface only
Cavity / etalon effect	Yes — $R_1 \cdot R_2$ denominator	No
Emissivity model	Full cavity eps_eff	Kirchhoff: $1 - \text{Rho} - A_{\text{coat}}$
Downstream attenuation	$T_{\text{sys}, j}$ per element	$\text{Rho}_{\text{sys}, j}$ per element
Sensitive to temperature	Bulk + surface	Surface only
Dominant in cold systems	Windows, filters	Warm fold mirrors

References

Full citations for the short keys used across the theory chapters (e.g. [Holst 2008], [Fried 1966]). One list, alphabetical by key; chapters cite by key only, adding a section locator after the key where one helps — [Holst 2008 §3.2].

Key	Full citation
[ASTM E490]	ASTM E490-00a(2019), <i>Standard Solar Constant and Zero Air Mass Solar Spectral Irradiance Tables</i> , ASTM International, West Conshohocken, PA.
[Bucholtz 1995]	Bucholtz, A., “Rayleigh-scattering calculations for the terrestrial atmosphere,” <i>Appl. Opt.</i> 34(15), 2765–2773 (1995).
[Boreman 2001]	Boreman, G. D., <i>Modulation Transfer Function in Optical and Electro-Optical Systems</i> , SPIE Press, 2001.
[Chapman 1931]	Chapman, S., “The absorption and dissociative or ionizing effect of monochromatic radiation in an atmosphere on a rotating earth,” <i>Proc. Phys. Soc.</i> 43, 26–45 (1931).
[CODATA 2018]	Tiesinga, E., Mohr, P. J., Newell, D. B., Taylor, B. N., “CODATA recommended values of the fundamental physical constants: 2018,” <i>Rev. Mod. Phys.</i> 93, 025010 (2021).
[Dereniak & Boreman 1996]	Dereniak, E. L., Boreman, G. D., <i>Infrared Detectors and Systems</i> , Wiley, 1996.
[Elsasser 1942]	Elsasser, W. M., <i>Heat Transfer by Infrared Radiation in the Atmosphere</i> , Harvard Meteorological Studies 6, 1942.
[Fried 1966]	Fried, D. L., “Optical Resolution Through a Randomly Inhomogeneous Medium for Very Long and Very Short Exposures,” <i>J. Opt. Soc. Am.</i> 56(10), 1372–1379 (1966).
[Gaskill 1978]	Gaskill, J. D., <i>Linear Systems, Fourier Transforms, and Optics</i> , Wiley, 1978.
[Hansen & Travis 1974]	Hansen, J. E., Travis, L. D., “Light scattering in planetary atmospheres,” <i>Space Sci. Rev.</i> 16, 527–610 (1974).
[Henyey & Greenstein 1941]	Henyey, L. G., Greenstein, J. L., “Diffuse radiation in the galaxy,” <i>Astrophys. J.</i> 93, 70–83 (1941).
[Goodman 2005]	Goodman, J. W., <i>Introduction to Fourier Optics</i> , 3rd ed., Roberts & Company, 2005.
[Harrington 2015]	Harrington, L., et al., “General Image Quality Equation, version 5 (GIQE-5),” NGA Technical Documentation, 2015.
[Holst 2008]	Holst, G. C., <i>Electro-Optical Imaging System Performance</i> , 5th ed., JCD Publishing / SPIE Press, 2008.
[Janesick 2001]	Janesick, J. R., <i>Scientific Charge-Coupled Devices</i> , SPIE Press, 2001.
[Johnson 1958]	Johnson, J., “Analysis of Image Forming Systems,” <i>Proc. Image Intensifier Symposium</i> , U.S. Army ERDL, 1958.
[Kirchhoff 1860]	Kirchhoff, G., “Ueber das Verhaeltniss zwischen dem Emissionsvermoegen und dem Absorptionsvermoegen der Koerper fuer Waerme und Licht,” <i>Ann. Phys.</i> 185(2), 275–301 (1860).
[Kopp & Lean 2011]	Kopp, G., Lean, J. L., “A new, lower value of total solar irradiance: evidence and climate significance,” <i>Geophys. Res. Lett.</i> 38, L01706 (2011).
[Lewis 1994]	Lewis, R. R., “Making shaders more physically plausible,” <i>Computer Graphics Forum</i> 13(2), 109–120 (1994).

Key	Full citation
[McClatchey 1972]	McClatchey, R. A., et al., <i>Optical Properties of the Atmosphere</i> , 3rd ed., AFCRL-72-0497, 1972.
[Nicodemus 1977]	Nicodemus, F. E., Richmond, J. C., Hsia, J. J., Ginsberg, I. W., Limperis, T., <i>Geometrical Considerations and Nomenclature for Reflectance</i> , NBS Monograph 160, National Bureau of Standards, 1977.
[NIST ITS-90]	Preston-Thomas, H., “The International Temperature Scale of 1990 (ITS-90),” <i>Metrologia</i> 27(1), 3–10 (1990).
[Noll 1976]	Noll, R. J., “Zernike polynomials and atmospheric turbulence,” <i>J. Opt. Soc. Am.</i> 66(3), 207–211 (1976).
[O’Neill 1956]	O’Neill, E. L., “Transfer Function for an Annular Aperture,” <i>J. Opt. Soc. Am.</i> 46(4), 285–288 (1956).
[Phong 1975]	Phong, B. T., “Illumination for computer generated pictures,” <i>Communications of the ACM</i> 18(6), 311–317 (1975).
[Planck 1901]	Planck, M., “Ueber das Gesetz der Energieverteilung im Normalspectrum,” <i>Ann. Phys.</i> 309(3), 553–563 (1901).
[Press et al.]	Press, W. H., Teukolsky, S. A., Vetterling, W. T., Flannery, B. P., <i>Numerical Recipes: The Art of Scientific Computing</i> , 3rd ed., Cambridge University Press, 2007.
[Schulz & Caldwell 1995]	Schulz, M., Caldwell, L., “Nonuniformity correction and correctability of infrared focal plane arrays,” <i>Infrared Physics & Technology</i> 36(4), 763–777 (1995).
[Siegel & Howell]	Siegel, R., Howell, J. R., <i>Thermal Radiation Heat Transfer</i> , 4th ed., Taylor & Francis, 2002.
[Vallado 2013]	Vallado, D. A., <i>Fundamentals of Astrodynamics and Applications</i> , 4th ed., Microcosm Press, 2013.
[Vincent 1990]	Vincent, J. D., <i>Fundamentals of Infrared Detector Operation and Testing</i> , Wiley, 1990.
[Wehrli 1985]	Wehrli, C., <i>Extraterrestrial Solar Spectrum</i> , Publication No. 615, Physikalisch-Meteorologisches Observatorium Davos / World Radiation Center (PMOD/WRC), Davos Dorf, Switzerland, 1985.
[Wertz & Larson 1999]	Wertz, J. R., Larson, W. J. (eds.), <i>Space Mission Analysis and Design</i> , 3rd ed., Microcosm/Kluwer, 1999.
[Wolfe & Zissis]	Wolfe, W. L., Zissis, G. J. (eds.), <i>The Infrared Handbook</i> , rev. ed., Environmental Research Institute of Michigan (for the Office of Naval Research), 1985.

Chapters may add keys; keep this table alphabetical and one-key-per-row. Citations here are to editions actually consulted — do not bump edition numbers without checking the cited section survives.